

# Course Overview

CS-2011: Introduction to Computer Systems (Fall 2022)  
Lecture 1

# Outline

## ■ Big Picture

- Course theme
- Five realities
- How the course helps with other CS/ECE courses

## ■ Academic integrity

## ■ Logistics and Policies

# The Big Picture

# Course Theme:

## (Systems) Knowledge is Power!

### ■ Systems Knowledge

- How hardware (processors, memories, disk drives, network infrastructure) plus software (operating systems, compilers, libraries, network protocols) combine to support the execution of application programs
- How you as a programmer can best use these resources

### ■ Useful outcomes from taking CS-2011

- Become more effective programmers
  - Able to find and eliminate bugs efficiently
  - Able to understand and tune for program performance
- Prepare for later “systems” classes in CS, ECE, ...
  - Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems, Computer Security, Malware Analysis, etc.

# It's Important to Understand How Things Work

## ■ Why do I need to know this stuff?

- Abstraction is good, but don't forget reality

## ■ Most CS courses emphasize abstraction

- (CE courses less so)
- **Abstract** data types
- Asymptotic analysis
  - e.g., Big-O notation (best case, average case, and worst case scenario of an algorithm)

## ■ These abstractions have limits

- Especially in the presence of bugs
- Need to understand details of underlying **implementations**
- Sometimes the abstract interfaces don't provide the level of control or performance you need

# Great Reality #1:


## Ints are not Integers, Floats are not Reals

### ■ Example 1: Is $x^2 \geq 0$ ?

■ Float's: Yes!

■ Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ??$



### ■ Example 2: Is $(x + y) + z = x + (y + z)$ ? (Associative?)

■ Unsigned & Signed Int's: Yes!

■ Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

# Computer Arithmetic

## ■ Does not generate random values

- Arithmetic operations have important mathematical properties

## ■ Cannot assume all “usual” mathematical properties

- Due to finiteness of representations
- Integer operations satisfy “ring” properties
  - Commutativity, associativity, distributivity
- Floating point operations satisfy “ordering” properties
  - Monotonicity, values of signs

## ■ Observation

- Need to understand which abstractions apply in which contexts
- Important issues for compiler writers and serious application programmers

# Great Reality #2:

## You've Got to Know Assembly

### ■ Chances are, you'll never write programs in assembly

- Compilers are much better & more patient than you are

### ■ But: Understanding assembly is key to machine-level execution model

- Behavior of programs in presence of bugs
  - High-level language models break down
- Tuning program performance
  - Understand optimizations done / not done by the compiler
  - Understanding sources of program inefficiency
- Implementing system software
  - Compiler has machine code as target
  - Operating systems must manage process state
- Creating / fighting malware
  - x86 assembly is the language of choice!



# Great Reality #3: Memory Matters

## Random Access Memory Is an Unphysical Abstraction

### ■ Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated

### ■ Memory referencing bugs especially pernicious

- Effects are distant in both time and space

### ■ Memory performance is not uniform

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

# Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.1399998664856
fun(3)	→	2.00000061035156
fun(4)	→	Segmentation fault

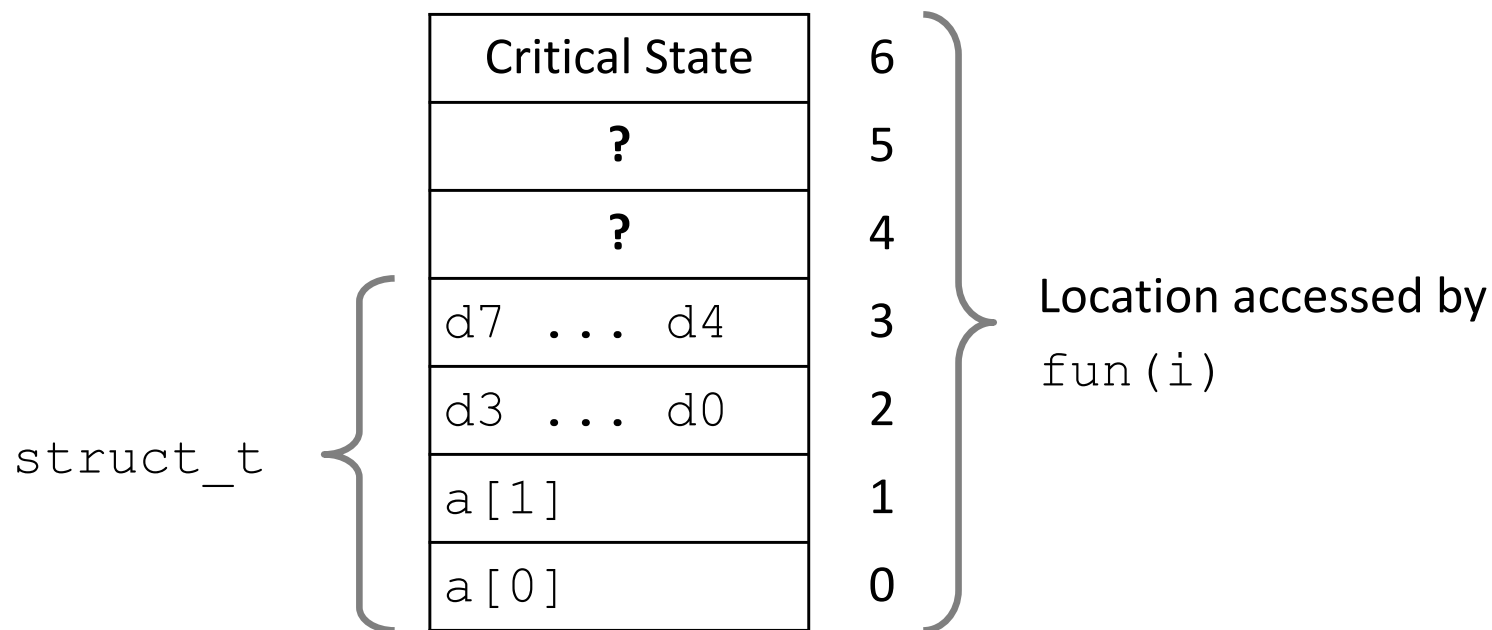
- Result is system specific

# Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.1399998664856
fun(3)	→	2.00000061035156
fun(4)	→	Segmentation fault

## Explanation:



# Memory Referencing Errors

## ■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

## ■ Can lead to nasty bugs

- Whether or not bug has any effect **depends on system and compiler**
- Action at a distance
  - Corrupted object logically unrelated to one being accessed
  - Effect of bug may be first observed long after it is generated

## ■ How can I deal with this?

- Program in Java, Ruby, Python, ML, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind)

## Great Reality #4: There's more to performance than asymptotic complexity

### ■ Constant factors matter too!

### ■ And even exact op count does not predict performance

- Easily see 10:1 performance range depending on how code written
- Must optimize at multiple levels: algorithm, data representations, procedures, and loops

### ■ Must understand system to optimize performance

- How programs compiled and executed
- How to measure program performance and identify bottlenecks
- How to improve performance without destroying code modularity and generality

# Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

**4.3ms**

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

**81.8ms**


2.0 GHz Intel Core i7 Haswell

■ Hierarchical memory organization

■ Performance depends on access patterns

- Including how step through multi-dimensional array

# Why The Performance Differs



# Great Reality #5:

## Computers do more than execute programs

### ■ They need to get data in and out

- I/O system critical to program reliability and performance

### ■ They communicate with each other over networks

- Many system-level issues arise in presence of network
  - Concurrent operations by autonomous processes
  - Coping with unreliable media
  - Cross platform compatibility
  - Complex performance issues



# Course Perspective

## ■ Most Systems Courses are Builder-Centric

- Computer Architecture
  - Design pipelined processor in Verilog
- Operating Systems
  - Implement sample portions of operating system
- Compilers
  - Write compiler for simple language
- Networking
  - Implement and simulate network protocols

# Course Perspective (Cont.)

## ■ This Course is Programmer-Centric

- Purpose is to show that by knowing more about the underlying system, one can be more effective as a programmer
- Enable you to
  - Write programs that are more reliable and efficient
  - Incorporate features that require hooks into OS
    - E.g., concurrency, *signal handlers*
- Cover material in this course that you won't see elsewhere
- Not just a course for dedicated hackers
  - **This course brings out the hidden hacker in everyone!**

# Role within CS/ECE Curriculum

Imperative  
Programming

**CS2011**

*Foundation of Computer Systems*  
*Underlying principles for hardware,  
software, and networking*

## CS Systems

- Cloud Computing
- Computer Security
- Operating Systems
- Compiler Design
- Database Applications
- Parallel Computing
- Distributed Systems
- Computer Networks
- Database Systems

## ECE Systems

- Computer Security
- Malware Analysis
- Intro to Embedded Systems
- Computer Networks
- Computer Architecture
- Wireless Networking
- Cyberphysical Systems

## CS Graphics

- Computer Graphics
- Comp. Photography

# Academic Integrity

## ■ DON'Ts

- Sharing code: by copying, retyping, **looking at**, or supplying a file
- Describing: verbal description of code from one person to another.
- Coaching: helping your friend to write a lab, line by line
- **Searching the Web** for solutions, discussions, tutorials, blogs, other universities' systems courses,... in English or any other language
- Copying code from a previous course or online solution
  - You are only allowed to use code I supply, or from the book author's website

## ■ OK

- Explaining how to use systems or tools to others
- Helping others with high-level design issues
  - Code / pseudo-code is **NOT** high level

# Textbooks

## ■ Randal E. Bryant and David R. O'Hallaron,

- [\*Computer Systems: A Programmer's Perspective\*](#), **Third Edition** (CS:APP3e), Pearson, 2016, ISBN 0-13-409266-X
  - **Hardcopy** also available on [Amazon](#) (rent, used, or new “more expensive than Pearson”)
  - Electronic edition available [here](#) (Rent 180 days \$44.99, Lifetime \$74.99)
  - Do not buy an *earlier edition* (1<sup>st</sup> / 2<sup>nd</sup>)
  - **DO NOT BUY A PAPERBACK/INTERNATIONAL EDITION (HAS MANY ISSUES)**
  - **DO NOT BUY KINDLE (BASED ON INTERNATIONAL VERSION)**
- Book Website: <https://csapp.cs.cmu.edu>
- This book really matters for the course! (**REQUIRED**)
  - How to solve labs
  - Practice problems typical of exam problems (or any other potential homework assignments)

## ■ Brian Kernighan and Dennis Ritchie,

- *The C Programming Language*, Second Edition, Prentice Hall, 1988
- Still the best book about C, from the originators
- Even though it does not cover more recent extensions of C

# Note on Textbook Global Edition

## ■ Quote from textbook's website

Unfortunately, the publisher arranged for the generation of a different set of practice and homework problems in the global edition. The person doing this didn't do a very good job, and so these problems and their solutions have many errors. We have not created an errata for this edition.

# Course Components (Subject to Change)

## ■ Lectures

- Higher level concepts
- May be utilized to teach you some applied concepts, important tools and skills for labs, deeper clarification of certain concepts, exam coverage, **demonstrations**, etc.
- **May** run random online Quizzes on canvas

## ■ Lab Assignments (~5-6)

- The heart of the course
- About 1-2 weeks each
- Provide in-depth understanding of an aspect of systems
- Programming and measurement

## ■ Homework Assignments (~2)

- Mainly arithmetic calculations, etc.. (no coding)

## ■ Exam(s)

- Current Options: 1) Midterm + Final OR, a 2) Comprehensive Final
- Test your understanding of concepts & mathematical principles, etc...
- Will **most likely** be held online through **canvas**

# Getting Help

## ■ Instructor's office hours:

- Answering questions about theoretical concepts we did cover in lecture
- Time/Location: Check Canvas

## ■ TA's office hours:

- Technical help with labs, homework assignments, etc.
- Most likely will run tutoring hours (Details will be added to Canvas)
- Time/Location: Check Canvas

## ■ CANVAS

- PLEASE CHECK CANVAS CONSTANTLY FOR UPDATES



# Timeliness

## ■ Late Penalty

- You **may** submit your assignments up to 24 hours **late** without getting penalized **as long as this does not happen more than 2 times** throughout the semester.
- If you submit late within the 24-hour period **more than two times** throughout the semester, **OR** if you submit after the 24-hour period, you will risk receiving a late penalty of up to **15% per day**
- No late submission is allowed **3 days after due date**
  - Graders can easily miss your submission

## ■ Catastrophic events

- Major illness, death in family, ...
- Formulate a plan (with your academic advisor) to get back on track

## ■ Advice

- Once you start running late, **it's really really hard to catch up**

# During Lecture

## ■ Laptop use is encouraged

- No GAMING please :) !!!

## ■ Please no electronic communications

- No email, instant messaging, cell phone calls, etc

## ■ Be Present

- Attendance will not be taken but is very **strongly encouraged!!!**

## ■ Please NO recordings of ANY KIND

# Grading (Tentative)

## ■ Labs/Homeworks (70%)

- Most likely weighted according to effort

## ■ Exams (30%)

# Programs and Data

## ■ Topics

- Bits operations, arithmetic, assembly language programs
- Representation of C control and data structures
- Includes aspects of architecture and compilers

## ■ Possible Lab Assignments

- Lab0: Test/refresh your C programming abilities
- Lab1: Manipulating bits
- Lab2: The basics of Assembly/Disassembly
- Lab3: The basics of code injection attacks

# The Memory Hierarchy

## ■ Topics

- Memory technology, memory hierarchy, caches, disks, locality
- Includes aspects of architecture and OS

## ■ Possible Lab Assignments

- Lab4: Building a cache simulator and optimizing for locality.
  - Learn how to exploit locality in your programs.

# Virtual Memory

## ■ Topics

- Virtual memory, address translation, dynamic storage allocation
- Includes aspects of architecture and OS

## ■ Possible Lab Assignments

- L5: Writing your own malloc package
  - Get a real feel for systems-level programming

# Exceptional Control Flow

## ■ Topics

- Hardware exceptions, processes, process control, Unix signals, nonlocal jumps BN
- Includes aspects of compilers, OS, and architecture

## ■ Possible Lab Assignments

- Lab6: Writing your own Unix shell.
  - A first introduction to concurrency

*Welcome and Enjoy!*