

Bits, Bytes, and Integers – Part 2

CS2011: Introduction to Computer Systems
Lecture 4

Bits, Bytes, and Integers

■ Representing information as bits

■ Bit-level manipulations

■ Integers

- Representation: unsigned and signed
- Conversion, casting
- Expanding, truncating

previous lecture

-
- **Addition, negation, multiplication, shifting**

this lecture

■ Representations in memory, pointers, strings

■ Summary

Unsigned Addition

Operands: w bits

u

+ v

True Sum: $w+1$ bits

$u + v$

Discard Carry: w bits

$\text{UAdd}_w(u, v)$

Standard Addition Function

- Ignores carry output

Implements Modular Arithmetic

$$s = \text{UAdd}_w(u, v) = u + v \bmod 2^w$$

unsigned char

```

      1110 1001
    + 1101 0101
    -----
  
```

```

      E9
    + D5
    -----
  
```

```

      233
    + 213
    -----
  
```

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Unsigned Addition

Operands: w bits

u

+ v

True Sum: $w+1$ bits

$u + v$

Discard Carry: w bits

$\text{UAdd}_w(u, v)$

Standard Addition Function

- Ignores carry output

Implements Modular Arithmetic

$$s = \text{UAdd}_w(u, v) = u + v \bmod 2^w$$

unsigned char

```

      1110 1001
    + 1101 0101
    -----
  1 1011 1110
    -----
      1011 1110
  
```

```

      E9
    + D5
    -----
    1BE
    -----
      BE
  
```

= 446 - 256

```


      233
    + 213
    -----
    446
    -----
    190
  
```

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Visualizing (Mathematical) Integer Addition

Integer Addition

- 4-bit integers u, v
- Compute true sum $\text{Add}_4(u, v)$
- Values increase **linearly** with u and v
- Forms planar surface





Visualizing Unsigned Addition

■ Wraps Around

- If true sum $\geq 2^w$
- At most once
- E.g., $w = 4$, after wrap around \rightarrow max sum = 14

True Sum




Two's Complement Addition

Operands: w bits



True Sum: $w+1$ bits



Discard Carry: w bits



■ TAdd and UAdd have Identical Bit-Level Behavior

- Signed vs. unsigned addition in C:

```
int s, t, u, v;
s = (int) ((unsigned) u + (unsigned) v);
t = u + v;
```

- Will give $s == t$

1110 1001	E9	-23
+ 1101 0101	+ D5	+ -43
<hr/>	<hr/>	<hr/>
1 1011 1110	1BE	-66
<hr/>	<hr/>	<hr/>
1011 1110	BE	-66

TAdd Overflow

■ Functionality

- True sum requires $w+1$ bits
- Drop off MSB
- Treat remaining bits as 2's comp. integer

0 111...1


0 100...0

0 000...0

1 011...1

1 000...0

True Sum



Visualizing 2's Complement Addition


■ Values

- 4-bit two's comp.
- Range from -8 to +7

■ Wraps Around

- If $\text{sum} \geq 2^{w-1}$
 - **Becomes negative**
 - At most once
- If $\text{sum} < -2^{w-1}$
 - **Becomes positive**
 - At most once

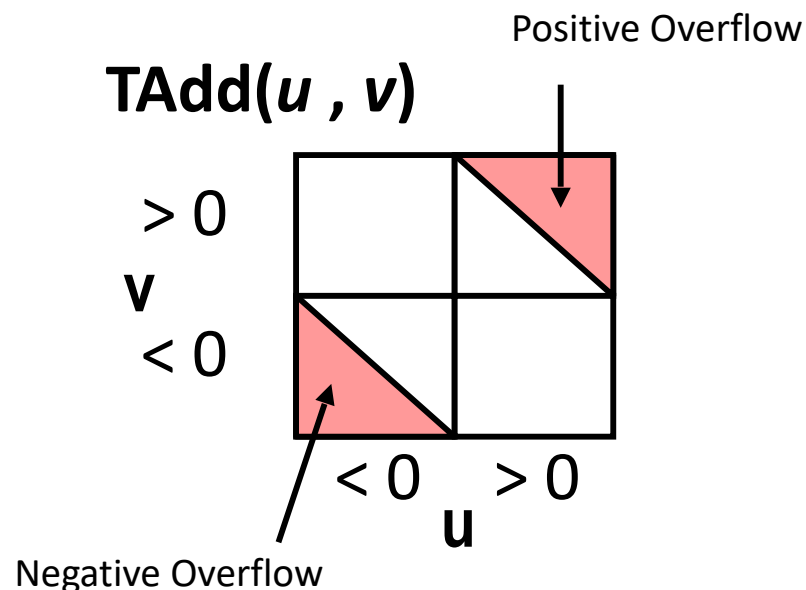
NegOver



Characterizing TAdd

■ Functionality

- True sum requires $w+1$ bits
- Drop off MSB
- Treat remaining bits as 2's comp. integer



$$TAdd_w(u, v) = \begin{cases} u + v + 2^w & u + v < TMin_w \text{ (NegOver)} \\ u + v & TMin_w \leq u + v \leq TMax_w \\ u + v - 2^w & TMax_w < u + v \text{ (PosOver)} \end{cases}$$

e.g., $-5 + -5$
e.g., $3 + -8$
e.g., $4 + 4$

Multiplication

■ Goal: Computing Product of w -bit numbers x, y

- Either signed or unsigned

■ But, exact results can be bigger than w bits


- **Unsigned**: up to $2w$ bits
 - Result range: $0 \leq x * y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$
 - E.g., 4-bits: $0 \leq x * y \leq (2^4 - 1)^2 = 2^8 - 2^5 + 1 = 256 - 32 + 1 = 225$
- **Two's complement** min (negative): Up to $2w-1$ bits
 - Result range: $x * y \geq (-2^{w-1}) * (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1}$
 - E.g., 4-bits: $x * y \geq (-2^3) * (2^3 - 1) = -8 * 7 = -56$
- **Two's complement** max (positive): Up to $2w$ bits, but only for $(TMin_w)^2$
 - Result range: $x * y \leq (-2^{w-1})^2 = 2^{2w-2}$
 - E.g., 4-bits: $x * y \leq (-2^3)^2 = (-8)^2 = 64$

■ So, maintaining exact results...

- would need to keep expanding word size with each product computed
- is done in software, if needed
 - e.g., by “arbitrary precision” arithmetic packages

Unsigned Multiplication in C

Operands: w bits



■ Standard Multiplication Function

- Ignores high order w bits

■ Implements Modular Arithmetic

$$\text{UMult}_w(u, v) = u \cdot v \bmod 2^w$$

$$\begin{array}{r}
 1110 \ 1001 \\
 * \quad 1101 \ 0101 \\
 \hline
 1100 \ 0001 \quad 1101 \ 1101 \\
 \quad 1101 \ 1101 \\
 \hline
 \end{array}$$

$$\begin{array}{r}
 \text{E9} \\
 * \quad \text{D5} \\
 \hline
 \text{C1DD} \\
 \quad \text{DD} \\
 \hline
 \end{array}$$


$$\begin{array}{r}
 233 \\
 * \quad 213 \\
 \hline
 49629 \\
 \quad 221 \\
 \hline
 \end{array}$$

Signed Multiplication in C

Operands: w bits

True Product: $2 \cdot w$ bits

Discard w bits: w bits



Standard Multiplication Function

- Ignores high order w bits
- Some of which are different for signed vs. unsigned multiplication
- Lower bits are the same

	1110	1001
*	1101	0101
<hr/>		
0000	0011	1101 1101
<hr/>		
		1101 1101

	E9	-23
*	D5	-43
<hr/>		
	03DD	989
<hr/>		
	DD	-35

- $u \ll k$ gives $u * 2^k$
- Both signed and unsigned


Diagram illustrating the construction of a $2k$ -bit vector u from a k -bit vector. The top row shows a k -bit vector with three dots in the middle. The bottom row shows a $2k$ -bit vector with a 1 in the middle, surrounded by zeros and dots. An asterisk is next to the $2k$ label.

$$u \cdot 2^k$$

$$\begin{array}{l} \text{UMult}_w(u, 2^k) \\ \text{TMult}_w(u, 2^k) \end{array} \quad \begin{array}{|c|c|c|c|c|c|c|c|} \hline \bullet & \bullet & \bullet & & 0 & \bullet & \bullet & \bullet \\ \hline \end{array}$$

- $u \ll 3 \quad == \quad u * 8$
- $(u \ll 5) - (u \ll 3) \quad == \quad u * 24$
- E.g., $2 * 24 = 2 * (32 - 8) = (2 * 32) - (2 * 8)$
 $(2 \ll 5) - (2 \ll 3)$
- Most machines shift and add faster than multiply
 - Compiler generates this code automatically

Unsigned Power-of-2 Divide with Shift



example with no binary point:


$$32 / 8 = 4 \dots 32 \gg 3 = 4$$

	Division	Computed	Hex	Binary
x	15213	15213	3B 6D	00111011 01101101
x >> 1	7606.5	7606	1D B6	00011101 10110110
x >> 4	950.8125	950	03 B6	00000011 10110110
x >> 8	59.4257813	59	00 3B	00000000 00111011

Signed Power-of-2 Divide with Shift

■ Quotient of Signed by Power of 2

- $x \gg k$ gives $\lfloor x / 2^k \rfloor$
- Uses **arithmetic** shift
- Rounds wrong direction when $x < 0$



	Division	Computed	Hex	Binary
y	-15213	-15213	C4 93	11000100 10010011
$y \gg 1$	-7606.5	-7607	E2 49	1 1100010 01001001
$y \gg 4$	-950.8125	-951	FC 49	1111 1100 01001001
$y \gg 8$	-59.4257813	-60	FF C4	11111111 11000100


Correct Power-of-2 Divide

■ Quotient of Negative Number by Power of 2


- Want $\lceil x / 2^k \rceil$ (Round Toward 0)
- Compute as $\lfloor (x + 2^k - 1) / 2^k \rfloor$
 - In C: $(x + (1 \ll k) - 1) \gg k$
 - Biases dividend toward 0

Case 1: No rounding

Dividend:




$+2^k - 1$



Divisor:

$/ 2^k$



$\lceil x / 2^k \rceil$

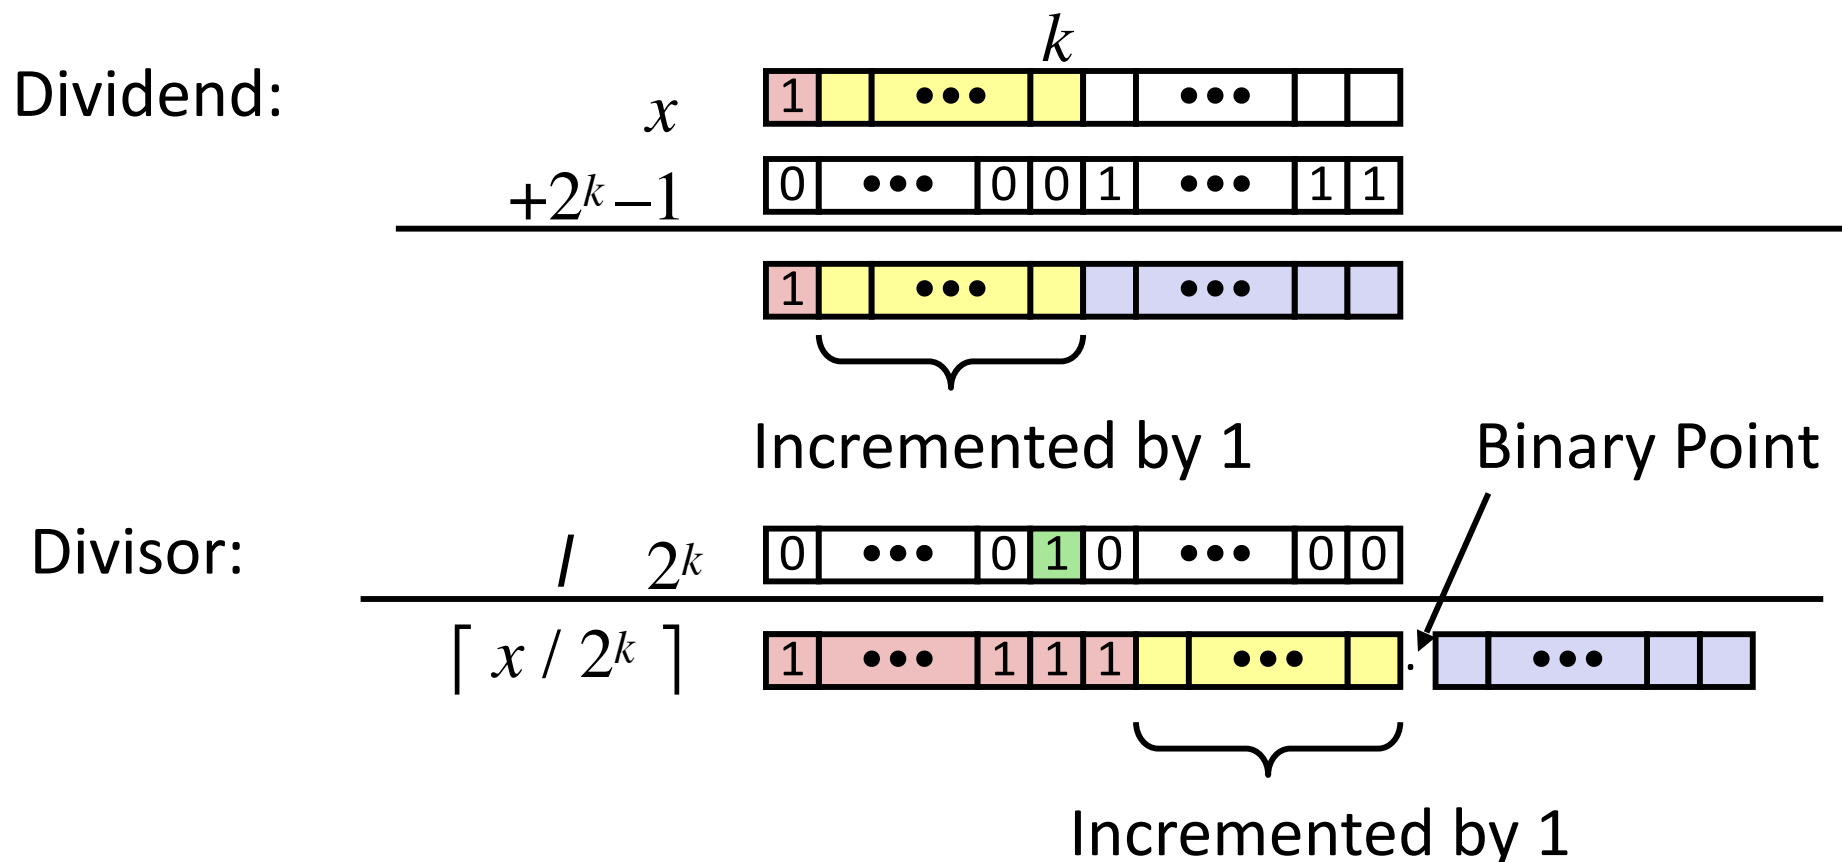


Binary Point

Biasing has no effect

Correct Power-of-2 Divide (Cont.)

Case 2: Rounding



Biasing adds 1 to final result

Correct Power-of-2 Divide (Example)

-17 / 16

Dividend:

$$\begin{array}{r}
 x \quad 101111 \quad -17 \\
 +2^k - 1 \quad 1111 \quad \text{Bias } (1 \ll 4) - 1 \\
 \hline
 111110
 \end{array}$$

Incremented by 1

Arithmetic Shift right by 4

$111110 \gg 4$

Binary Point

$$\begin{array}{r}
 \lceil x / 2^k \rceil = \lceil -17 / 2^4 \rceil = -1 \quad 11.1111 \\
 \hline
 \end{array}$$

Incremented by 1

Biasing adds 1 to final result

Negation: Complement & Increment

■ Negate through complement and increase

$\sim x + 1 == -x$ (works for both signed and unsigned)

■ Example

■ Observation: $\sim x + x == 1111\dots111 == -1$

$$\begin{array}{r}
 x \quad 10011101 \\
 + \quad \sim x \quad 01100010 \\
 \hline
 -1 \quad 11111111
 \end{array}$$

x = 15213

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
~x	-15214	C4 92	11000100 10010010
~x+1	-15213	C4 93	11000100 10010011
y	-15213	C4 93	11000100 10010011

Complement & Increment Examples (Special Cases)

$x = 0$

	Decimal	Hex	Binary
0	0	00 00	00000000 00000000
~ 0	-1	FF FF	11111111 11111111
$\sim 0 + 1$	0	00 00	00000000 00000000

Same for unsigned

$x = \text{TMin}$ (Signed Two's Complement)

	Decimal	Hex	Binary
x	-32768	80 00	10000000 00000000
$\sim x$	32767	7F FF	01111111 11111111
$\sim x + 1$	-32768	80 00	10000000 00000000

Bits, Bytes, and Integers

■ Representing information as bits

■ Bit-level manipulations

■ Integers

- Representation: unsigned and signed
- Conversion, casting
- Expanding, truncating

previous lecture

-
- Addition, negation, multiplication, shifting

this lecture

- **Summary**

■ Representations in memory, pointers, strings

Arithmetic: Basic Rules

■ Addition:

- Unsigned/signed: Normal addition followed by truncate, same operation on bit level
- Unsigned: addition mod 2^w
 - Mathematical addition + possible subtraction of 2^w
- Signed: modified addition mod 2^w (result in proper range)
 - Mathematical addition + possible addition or subtraction of 2^w

■ Multiplication:

- Unsigned/signed: Normal multiplication followed by truncate, same operation on bit level
- Unsigned: multiplication mod 2^w
- Signed: modified multiplication mod 2^w (result in proper range)

Why Should I Use Unsigned?

■ *Don't* use without understanding implications

- Easy to make mistakes

```
unsigned i;  
for (i = cnt-2; i >= 0; i--)  
    a[i] += a[i+1];
```

- Can be very subtle

```
#define DELTA sizeof(int)  
int i;  
for (i = CNT; i-DELTA >= 0; i-= DELTA)  
    . . .
```


Counting Down with Unsigned

■ Proper way to use unsigned as loop index

```
unsigned i;  
for (i = cnt-2; i < cnt; i--)  
    a[i] += a[i+1];
```

■ See Robert Seacord, *Secure Coding in C and C++*

- C Standard guarantees that unsigned addition will behave like modular arithmetic
 - $0 - 1 \rightarrow UMax$

■ Even better

```
size_t i;  
for (i = cnt-2; i < cnt; i--)  
    a[i] += a[i+1];
```

- Data type `size_t` defined as unsigned value with length = word size
- Code will work even if `cnt == UMax`

Why Should I Use Unsigned? (cont.)

■ **Do Use When Performing Modular Arithmetic**

- Multiprecision arithmetic

■ **Do Use When Using Bits to Represent Sets**

- Logical right shift, no sign extension

■ **Do Use In System Programming**

- Bit masks, device commands,...

—> unsigned numbers can be very **useful** and C language supports it
(**Java** does not support unsigned numbers)

Integer C Puzzles

Initialization

```
int x = foo();
int y = bar();
unsigned ux = x;
unsigned uy = y;
```

<code>x < 0</code>	\Rightarrow	<code>((x*2) < 0)</code>	✗
<code>ux >= 0</code>			✓
<code>x & 7 == 7</code>	\Rightarrow	<code>(x<<30) < 0</code>	✓
<code>ux > -1</code>			✗
<code>x > y</code>	\Rightarrow	<code>-x < -y</code>	✗
<code>x * x >= 0</code>			✗
<code>x > 0 && y > 0</code>	\Rightarrow	<code>x + y > 0</code>	✗
<code>x >= 0</code>	\Rightarrow	<code>-x <= 0</code>	✓
<code>x <= 0</code>	\Rightarrow	<code>-x >= 0</code>	✗
<code>(x -x)>>31 == -1</code>			✗
<code>ux >> 3 == ux/8</code>			✓
<code>x >> 3 == x/8</code>			✗
<code>x & (x-1) != 0</code>			✗

Bits, Bytes, and Integers

■ Representing information as bits

■ Bit-level manipulations

■ Integers

- Representation: unsigned and signed
- Conversion, casting
- Expanding, truncating

previous lecture


-
- Addition, negation, multiplication, shifting

this lecture

- Summary

■ Representations in memory, pointers, strings

Byte-Oriented Memory Organization



■ Programs refer to data by address

- Conceptually, envision it as a very large array of bytes
 - In reality, it's not, but can think of it that way
- An address is like an index into that array
 - and, a pointer variable stores an address

■ Note: system provides private address spaces to each “process”

- Think of a process as a program being executed
- So, a program can clobber its own data, but not that of others

Machine Words

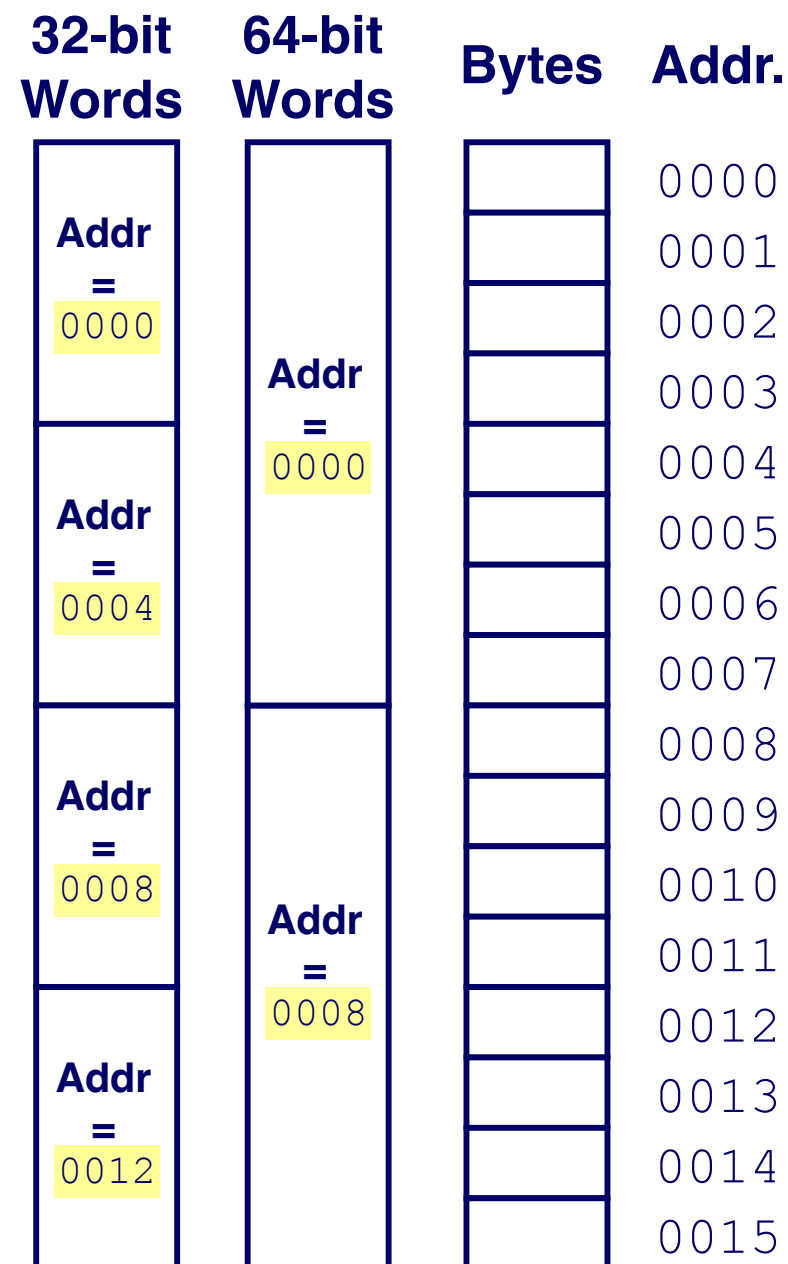
■ Any given computer has a “Word Size”

- Nominal size of integer-valued data
 - and of addresses
- Until recently, most machines used 32 bits (4 bytes) as word size
 - Limits addresses to 4GB (2^{32} bytes)
- Increasingly, machines have 64-bit word size
 - Potentially, could have 18 EB (exabytes) of addressable memory
 - That's 18.4×10^{18}
- Machines still support multiple data formats
 - Fractions or multiples of word size
 - Always integral number of bytes

Word-Oriented Memory Organization

■ Addresses Specify Byte Locations

- Address of first byte in word
- Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)



Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit	x86-64
<code>char</code>	1	1	1
<code>short</code>	2	2	2
<code>int</code>	4	4	4
<code>long</code>	4	8	8
<code>float</code>	4	4	4
<code>double</code>	8	8	8
<code>pointer</code>	4	8	8

Byte Ordering

■ So, how are the bytes within a multi-byte word ordered in memory?

■ Conventions


- Big Endian: Sun (Oracle SPARC), PPC Mac, *Internet*
 - Least significant byte has highest address
- Little Endian: *x86*, ARM processors running Android, iOS, and Linux
 - Least significant byte has lowest address

Byte Ordering Example


■ Example

- Variable x has 4-byte value of 0x01234567
- Address given by &x is 0x100

Big Endian



Little Endian




Representing Integers

Decimal: 15213


Binary: 0011 1011 0110 1101

Hex: 3 B 6 D


`int A = 15213;`



`long int C = 15213;`



`int B = -15213;`



Two's complement representation

Examining Data Representations

■ Code to Print Byte Representation of Data

- Casting pointer to unsigned char * allows treatment as a byte array

```
typedef unsigned char *pointer;

void show_bytes(pointer start, size_t len){
    size_t i;
    for (i = 0; i < len; i++)
        printf("%p\t0x%.2x\n", start+i, start[i]);
    printf("\n");
}
```

Printf directives:

%p: Print pointer

%x: Print Hexadecimal

show_bytes Execution Example

```
int a = 15213;  
printf("int a = 15213;\n");  
show_bytes((pointer) &a, sizeof(int));
```

Result (Linux x86-64):

```
int a = 15213;  
0x7fffb7f71dbc    6d  
0x7fffb7f71dbd    3b  
0x7fffb7f71dbe    00  
0x7fffb7f71dbf    00
```

Representing Pointers

```
int B = -15213;  
int *P = &B;
```

Sun	IA32	x86-64
EF	AC	3C
FF	28	1B
FB	F5	FE
2C	FF	82
		FD
		7F
		00
		00

Different compilers & machines assign different locations to objects

Even get different results each time run program

Representing Strings


■ Strings in C

- Represented by array of characters
- Each character encoded in ASCII format
 - Standard 7-bit encoding of character set
 - Character “0” has code 0x30
 - Digit i has code $0x30+i$
 - *man ascii* for code table
- String should be null-terminated
 - Final character = 0

■ Compatibility

- Byte ordering not an issue

```
char S[6] = "18213";
```



Reading Byte-Reversed Listings

■ Disassembly

- Text representation of binary machine code
- Generated by program that reads the machine code

■ Example Fragment

Address	Instruction Code	Assembly Rendition
8048365:	5b	pop %ebx
8048366:	81 c3 ab 12 00 00	add \$0x12ab,%ebx
804836c:	83 bb 28 00 00 00 00	cmpl \$0x0,0x28(%ebx)

■ Deciphering Numbers

- Value:
- Pad to 32 bits:
- Split into bytes:
- Reverse:

0x12ab

0x000012ab

00 00 12 ab

ab 12 00 00

Summary

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Representations in memory, pointers, strings
- Summary