

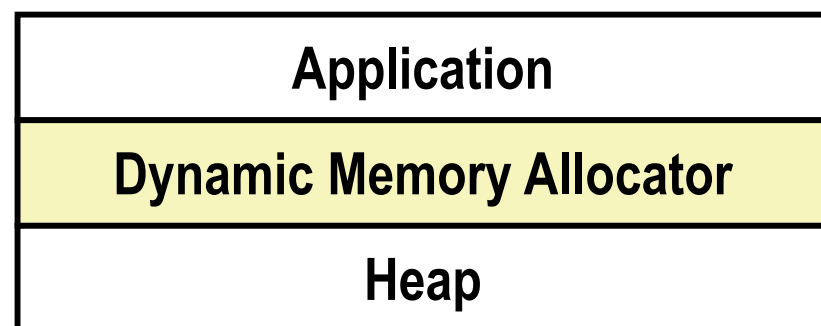
Dynamic Memory Allocation: Basic Concepts

CS2011: Introduction to Computer Systems
Lecture 13 (9.9)

Dynamic Memory Allocation: Basic Concepts

- Basic concepts
- Implicit free lists

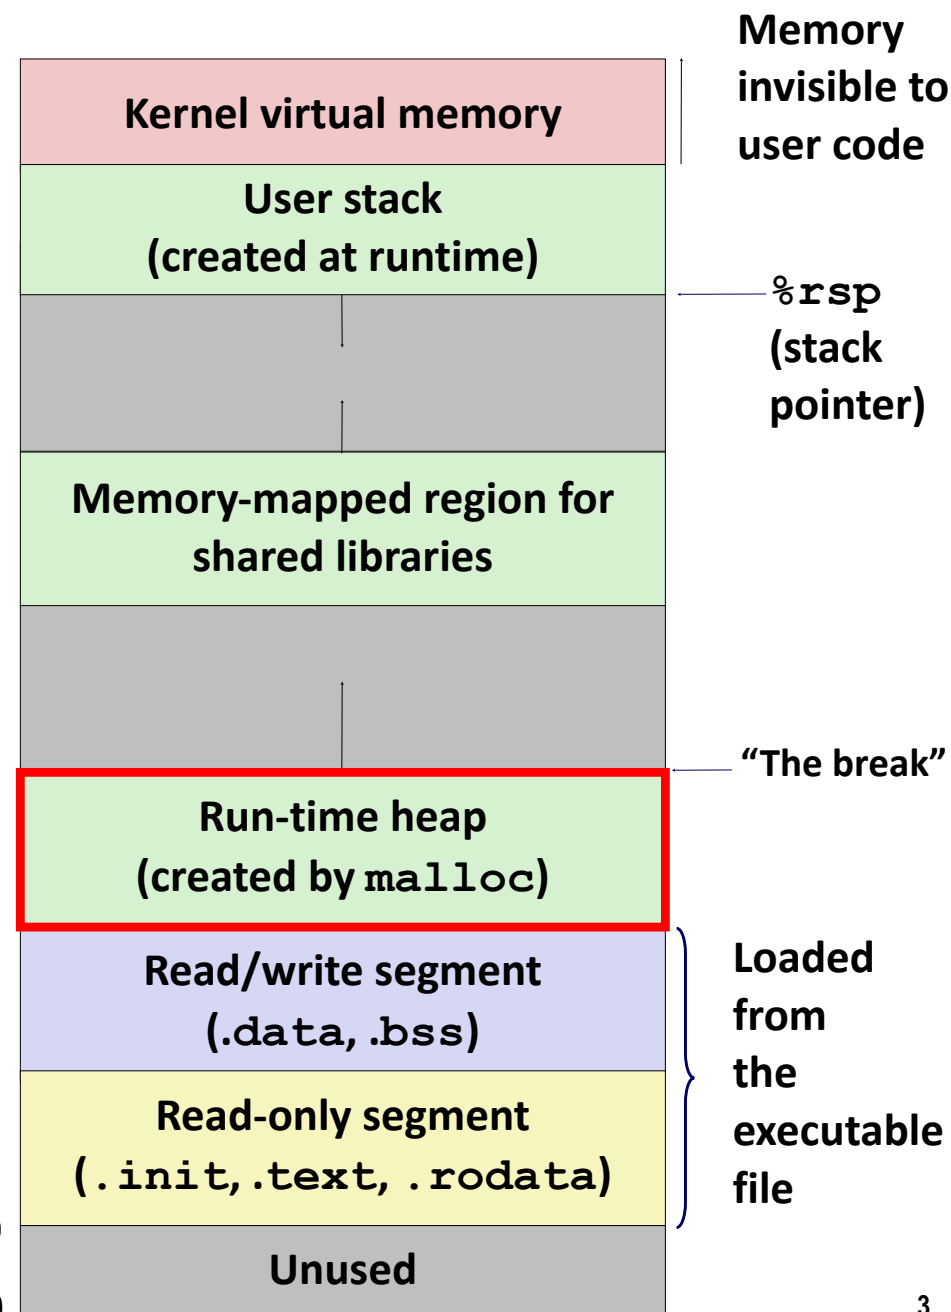
Dynamic Memory Allocation



- Programmers use *dynamic memory allocators* (such as `malloc`) to acquire virtual memory (VM) at runtime

- For data structures whose size is only known at runtime

- Dynamic memory allocators manage an area of process VM known as the *heap*



Dynamic Memory Allocation

- Allocator maintains heap as collection of variable sized *blocks*, which are either *allocated* or *free*
- Types of allocators
 - *1) Explicit allocator*: application allocates and frees space
 - e.g., `malloc` and `free` in `C`
 - *2) Implicit allocator*: application allocates, but does not free space
 - e.g., `new` and garbage collection in `Java`
- Will discuss simple *explicit* memory allocation in this lecture

The malloc Package

```
#include <stdlib.h>
```

```
void *malloc(size_t size)
```

- **Successful:**

- Returns a pointer to a memory block of at least **size** bytes *aligned* to a **16-byte boundary** (on x86-64)
- If **size == 0**, returns NULL

- **Unsuccessful:** returns NULL (0) and sets **errno**

```
void free(void *p)
```

- Returns the block pointed at by **p** to *pool* of available memory
- **p** must come from a previous call to **malloc**, **calloc**, or **realloc**

Other functions

- **calloc**: Version of **malloc** that **initializes allocated block to zero**
- **realloc**: **Changes the size** of a previously allocated block
- **sbrk**: Used **internally** by allocators to grow or shrink the heap

malloc Example

```
#include <stdio.h>
#include <stdlib.h>


void foo(long n) {
    long i, *p;

    /* Allocate a block of n longs */
    p = (long *) malloc(n * sizeof(long));
    if (p == NULL) {
        perror("malloc");
        exit(0);
    }

    /* Initialize allocated block */
    for (i=0; i<n; i++)
        p[i] = i;
    /* Do something with p */
    . . .
    /* Return allocated block to the heap */
    free(p);
}
```

Heap Visualization Convention / Assumption

- 1 square = 1 “word” = 8 bytes
- Allocations are **double-word (2-word) aligned**




Allocation Example (Conceptual)


```
p1 = malloc(32)
```




```
p2 = malloc(40)
```



```
p3 = malloc(48)
```



```
free(p2)
```



```
p4 = malloc(16)
```



Constraints

■ Applications

- Can issue **arbitrary sequence** of **malloc** and **free** requests (no special ordering)
- **free** request must be to a **malloc'd** block

■ Explicit Allocators

- Can't control number or size of allocated blocks
- Must **respond immediately** to **malloc** requests
 - *i.e.*, can't reorder or buffer requests
- Must **allocate blocks from free memory**
 - *i.e.*, can only place allocated blocks in free memory
- Must **align** blocks so they satisfy all alignment requirements
 - **16-byte (x86-64) alignment** on **64-bit systems**
- Can manipulate and modify only free memory blocks
- **Can't move** the allocated blocks once they are **malloc'd**
 - *i.e.*, compaction is not allowed. (Program wouldn't know how to free pointer if moved)

Performance Goal 1: Maximize Throughput

■ Given some sequence of `malloc` and `free` requests:

- $R_0, R_1, \dots, R_k, \dots, R_{n-1}$

■ Goals: **maximize throughput** and *peak memory utilization*

- These goals are often *conflicting*

■ Throughput:

- *Number of completed requests per unit time*

- Example:

- 500 `malloc` calls and 500 `free` calls in 1 seconds
- Throughput is 1,000 operations/second

Performance Goal 2: Maximize Memory Utilization

■ Given some sequence of `malloc` and `free` requests:

- $R_0, R_1, \dots, R_k, \dots, R_{n-1}$

■ After k requests we have:

■ **Def: Aggregate payload P_k**

- `malloc(p)` results in a block with a **payload** of p bytes
- The **aggregate payload** P_k is the **sum of currently allocated payloads**
- The **peak aggregate payload** $\max_{i \leq k} P_i$ is the **maximum** aggregate payload **at any point in the sequence up to request**

■ **Def: Current heap size H_k**

- Assume heap only grows when allocator uses `sbrk`, never shrinks

■ **Def: Peak memory utilization after $k+1$ requests**

- $U_k = (\max_{i \leq k} P_i) / H_k$

■ **Def: Overhead, O_k**

- **Fraction** of heap space **NOT** used for program data
- $O_k = (H_k / \max_{i \leq k} P_i) - 1.0$
- Goal: minimize overhead

Benchmark Example

■ Benchmark

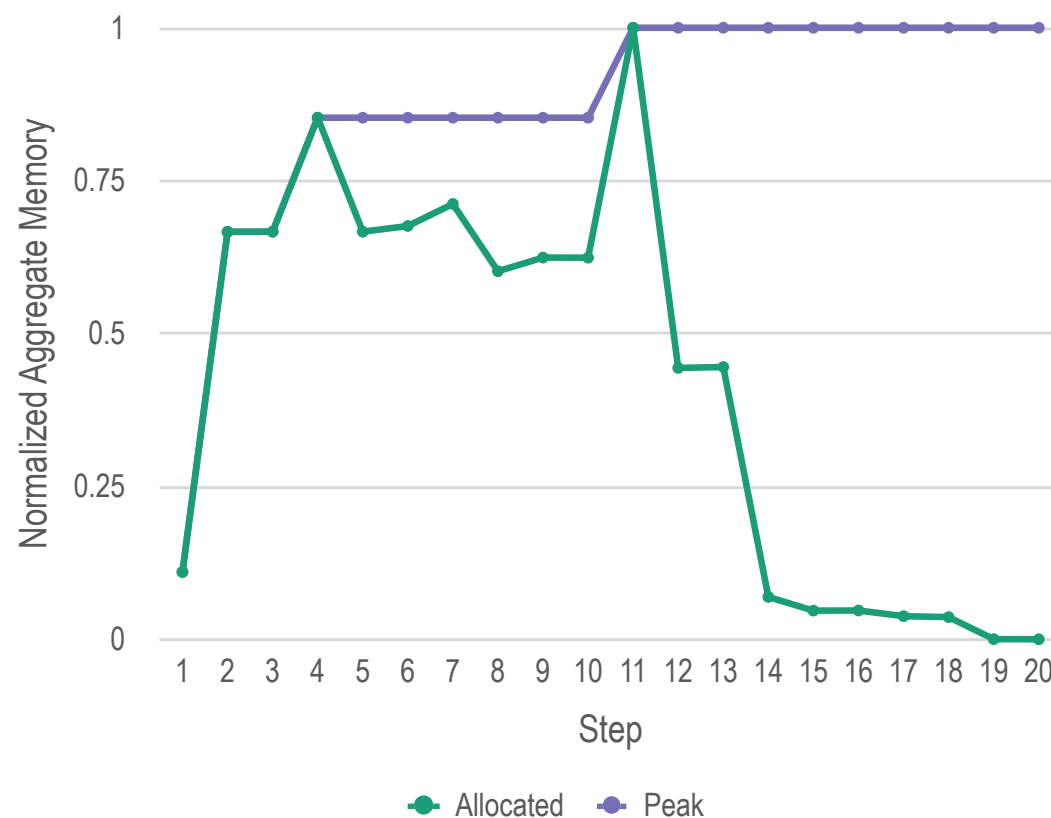
syn-array-short

- Example Trace
- Allocate & free 10 blocks
- **a** = **allocate**
- **f** = **free**
- Bias toward allocate at beginning & free at end
- Blocks number 1–10
- **Allocated**: Sum of all allocated amounts (i.e. *aggregate payload*)
- **Peak**: Max so far of Allocated (*peak aggregate payload*)

Step	Command	Delta	Allocated	Peak
1	a 0 9904	9904	9904	9904
2	a 1 50084	50084	59988	59988
3	a 2 20	20	60008	60008
4	a 3 16784	16784	76792	76792
5	f 3	-16784	60008	76792
6	a 4 840	840	60848	76792
7	a 5 3244	3244	64092	76792
8	f 0	-9904	54188	76792
9	a 6 2012	2012	56200	76792
10	f 2	-20	56180	76792
11	a 7 33856	33856	90036	90036
12	f 1	-50084	39952	90036
13	a 8 136	136	40088	90036
14	f 7	-33856	6232	90036
15	f 6	-2012	4220	90036
16	a 9 20	20	4240	90036
17	f 4	-840	3400	90036
18	f 8	-136	3264	90036
19	f 5	-3244	20	90036
20	f 9	-20	0	90036


Benchmark Visualization

Step	Command	Delta	Allocated	Peak
1	a 0 9904	9904	9904	9904
2	a 1 50084	50084	59988	59988
3	a 2 20	20	60008	60008
4	a 3 16784	16784	76792	76792
5	f 3	-16784	60008	76792
6	a 4 840	840	60848	76792
7	a 5 3244	3244	64092	76792
8	f 0	-9904	54188	76792
9	a 6 2012	2012	56200	76792
10	f 2	-20	56180	76792
11	a 7 33856	33856	90036	90036
12	f 1	-50084	39952	90036
13	a 8 136	136	40088	90036
14	f 7	-33856	6232	90036
15	f 6	-2012	4220	90036
16	a 9 20	20	4240	90036
17	f 4	-840	3400	90036
18	f 8	-136	3264	90036
19	f 5	-3244	20	90036
20	f 9	-20	0	90036



- Plot P_k (**allocated**) and $\max_{i \leq k} P_k$ (**peak**) as a function of k (step)
- Y-axis normalized — fraction of maximum

Typical Benchmark Behavior



- Longer sequence of mallocs & frees (40,000 blocks)
 - *Starts with all mallocs, and shifts toward all frees*
- Allocator must manage space efficiently the whole time (i.e. maximize Utilization)
- Production allocators can shrink the heap


Fragmentation

■ Poor memory utilization caused by *fragmentation*

- *Internal* fragmentation
- *External* fragmentation


Internal Fragmentation

- For a given block, *internal fragmentation* occurs if payload is smaller than block size



- Caused by
 - Overhead of *maintaining* heap *data structures*
 - Allocator implementation forces a minimum size of allocated blocks
 - *Padding* for *alignment* purposes
 - Explicit policy decisions
(e.g., to *return a big block to satisfy a small request*)
- Depends only on the pattern of *previous* requests
 - Thus, easy to measure
(“size of all allocated blocks” - “size of all payloads”)

Internal Fragmentation Effect



■ **Purple line: additional heap size due to allocator's data + padding for alignment**

- For this benchmark, 1.5% overhead
- Cannot achieve in practice
- Especially since cannot move allocated blocks


External Fragmentation

- Occurs when **there is enough aggregate heap memory, but no single free block is large enough**


```
p1 = malloc(32)
```



```
p2 = malloc(40)
```



```
p3 = malloc(48)
```



```
free(p2)
```




```
p4 = malloc(64)
```

Yikes! (what would happen now?)

- Depends on the pattern of future requests
 - Thus, difficult to measure (*e.g., if next request fits then no fragmentation*)

External Fragmentation Effect



■ **Green** line: additional heap size due to external fragmentation

■ **Best Fit: One allocation strategy**

- (To be discussed later)
- Total overhead = 8.3% on this benchmark

Implementation Issues

■ How do we know **how much memory to free** given just a pointer?

■ Free Block Organization

- How do we **keep track of the free blocks**?

■ Placement

- How do we **pick a block to use** for allocation -- many might fit?

■ Splitting

- What do we do with the **extra space when allocating** a structure that is smaller than the free block it is placed in?


■ Coalescing

- How do we **reuse a block** that has been freed?

Knowing How Much to Free


■ Standard method

- Keep the **length (in bytes)** of a block in the word *preceding* the block.
 - This word is often called the **header field** or **header**
 - **Length/Size** includes the header and any possible padding
 - Remember: Only payload needs to be aligned NOT the header
- Requires an extra word for every allocated block




Keeping Track of Free Blocks

Method 1: *Implicit list* using length—links all blocks



Need to **tag**
each block as
allocated/free

Method 2: *Explicit list* among the free blocks using pointers



Need space
for **pointers**

Method 3: *Segregated free list*

- Different combinations of free lists (e.g., implicit, explicit, etc) for different size classes (e.g., 4 bytes, 8 bytes, etc.)

Method 4: *Blocks sorted by size*

- Can use a balanced tree (e.g., Red-Black tree) with pointers within each free block, and the length used as a key


Dynamic Memory Allocation: Basic Concepts

- Basic concepts
- Implicit free lists

Method 1: Implicit Free List

- Called **implicit** because free blocks are linked implicitly by the size fields in the header
- For each block we need both size and allocation status
 - Could store this information in two words: wasteful!
- Standard trick
 - When blocks are **aligned**, some low-order address bits are always 0
 - If we impose that blocks are always **2-word aligned**, the lowest-order 4 bits of address are always 0
size is always a multiple of 16 (Remember we are assuming a word is 8 bytes)
 - Instead of storing an always-0 lowest-order bit, use it as an allocated/free flag
 - **When reading** the **Size** word, **must mask out this bit**
 (E.g., $0x0000000000000033 \& \sim 0xF = 0x0000000000000033 \& 0xFFFFFFFFFFFFFFF0 = 0x0000000000000032$)

*Format of
allocated and
free blocks*




a = 1: Allocated block

a = 0: Free block

Size: total block size including header, payload, and any padding

Payload: application data
(allocated blocks only)

Detailed Implicit Free List Example



Allocated blocks: shaded


Free blocks: unshaded

Headers: labeled with “size in words/allocated bit”

Headers are at non-aligned positions

→ Payloads are aligned

Implicit Free List Example Sequence



Double-word
aligned

Example Sequence:	Header Value:
Malloc (1)	0x11
Malloc (9)	0x21
Malloc (24)	0x21
Malloc (26)	0x31
Malloc (5)	0x11

Implicit List: Data Structures



■ Block declaration

```
typedef uint64_t word_t;
```

```
typedef struct block
{
    word_t header;
    unsigned char payload[0];           // Zero length array
} block_t;
```

■ Getting payload from block pointer // block_t *block

```
return (void *) (block->payload);
```

■ Getting header from payload // bp points to a payload

```
return (block_t *) ((unsigned char *) bp
    - offsetof(block_t, payload));
```

C function `offsetof(struct, member)` returns offset of member within struct

Implicit List: Header access

Size	a
------	---

■ Getting allocated bit from header (get_alloc)

```
return header & 0x1;
```

■ Getting size from header (get_size)

```
return header & ~0xFL; // L indicates Long
```

■ Initializing header

```
//block_t *block
```


```
block->header = size | alloc;
```

Implicit List: Traversing list



Find next block

```
static block_t *find_next(block_t *block)
{
    return (block_t *) ((unsigned char *) block
        + get_size(block));
}
```




Implicit List: Finding a Free Block (Placement)

■ *First fit:*

- **Search** list from beginning, choose *first* free block that fits:
- Finding space for **asize** bytes (including header):

```
static block_t *find_fit(size_t asize)
{
    block_t *block;
    for (block = heap_start; block != heap_end; block = find_next(block)) {
        if (!(get_alloc(block)) && (asize <= get_size(block)))
            return block;
    }
    return NULL; // No fit found
}
```



Implicit List: Finding a Free Block

■ *First fit:*

- Search list from beginning, choose **first** free block that fits:
- Can take **linear time** in total number of blocks (allocated and free)
- **Disadvantage:** In practice it can cause “**splinters/clutter**” of **small free blocks at beginning of list (due to splitting)**
 - **Increases the search time for larger blocks**


■ *Next fit:*

- Like first fit, but search list **starting where previous search finished or left off**
- Should often be **faster** than first fit: **avoids re-scanning unhelpful blocks**
 - **e.g., if size of next block is same as previously allocated block**
- **Disadvantage:** Some research suggests that **fragmentation is worse**

■ *Best fit:*

- Search the list, choose the **best** free block: fits, with fewest bytes left over
 - Examine **every free block** and choose the one with **smallest size** that fits
 - If you find a perfect match, you can stop the search there
- **Keeps fragments small**—usually **improves memory utilization**
- **Disadvantage:** Will typically run **slower** than first fit
- Still a greedy algorithm. No guarantee of optimality (Exhaustive search)

Comparing Strategies




■ Total Overheads (for this benchmark)

- Perfect Fit: 1.6%
- Best Fit: 8.3%
- First Fit: 11.9%
- Next Fit: 21.6%


Implicit List: Splitting Free Block on Allocation

■ Allocating in a free block: *splitting*


- Since allocated space might be smaller than free space, we **might** want to split the block
 - If placement policy produces good fits, some additional internal fragmentation is acceptable and no need for splitting



`split_block(p, 32)`



Implicit List: Splitting Free Block



// Warning: This code is incomplete


```
static void split_block(block_t *block, size_t asize){
    size_t block_size = get_size(block);

    // Split only if remainder of block is enough to fit
    // the smallest possible data
    if ((block_size - asize) >= min_block_size) {
        write_header(block, asize, true);
        block_t *block_next = find_next(block);
        write_header(block_next, block_size - asize, false);
    }
}
```

Implicit List: **Freeing** a Block

■ **Simplest** implementation:

- Need only **clear the “allocated” flag (bit 1 becomes 0)**
- But can lead to “**false fragmentation**” phenomena




`malloc(5*SIZ)` **Yikes!**

There is enough contiguous free space, but the allocator won't be able to find it

Implicit List: Coalescing Free Blocks

■ Join/merge (*coalesce*) with next/previous blocks, if they are free

- Coalescing with **next block**




- When to merge?

- Allocator may go for **immediate coalescing** (discussed here)
- Allocator may go for deferred coalescing (wait until later time)
 - To reduce the amount of possible repeated splitting and coalescing

Implicit List: Coalescing

■ Join (*coalesce*) with next block, if it is free

- Coalescing with **next block** is **simple** and takes **constant time**
 - We simply check the next block if free or allocated. If free, its size will be added to the size of the current block




- But, how do we coalesce with *previous block*?
 - How do we know where it starts?
 - How can we determine whether it's allocated?
 - Solution: **Boundary Tags**


Implicit List: Bidirectional Coalescing

■ *Boundary tags* [Knuth73]

- **Replicate** size/allocated word (i.e. **header**) at “bottom” (end) of free blocks
- Allows us to **traverse** the “list” **backwards**, but **requires extra space**
- Important and general technique!



*Format of
allocated and
free blocks*

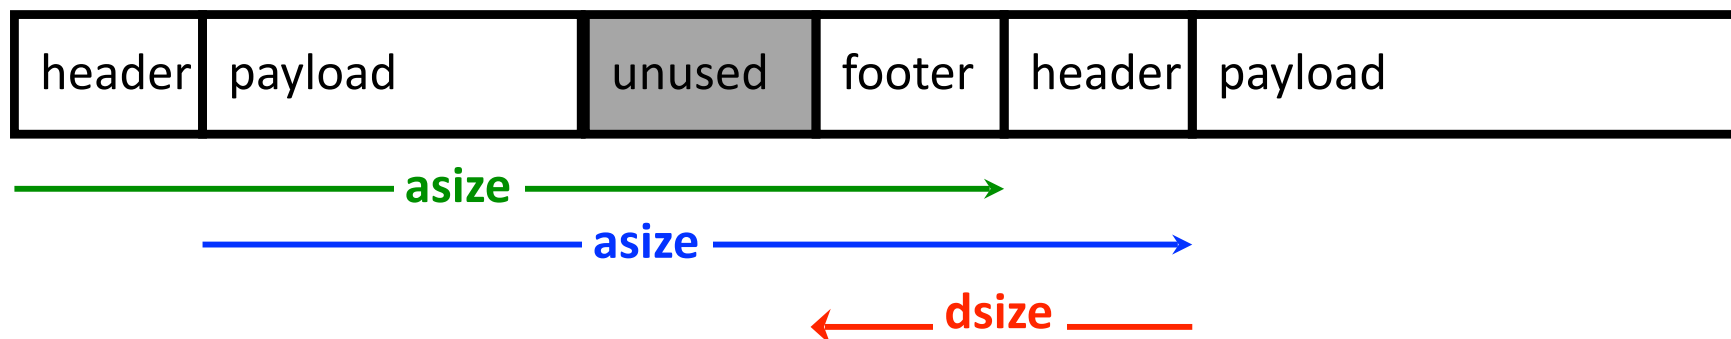


a = 1: Allocated block
a = 0: Free block

Size: Total block size

Payload: Application data
(allocated blocks only)

Implementation with Footers



■ Locating footer of **current block**

```
const size_t dsizes = 2*sizeof(word_t);

static word_t *header_to_footer(block_t *block)
{
    size_t asize = get_size(block);
    return (word_t *) (block->payload + asize - dsizes);
}
```

Implementation with Footers




←
1 word

■ Locating footer of **previous block**

```
static word_t *find_prev_footer(block_t *block)
{
    return &(block->header) - 1; // 1 word = 8 bytes
}
```

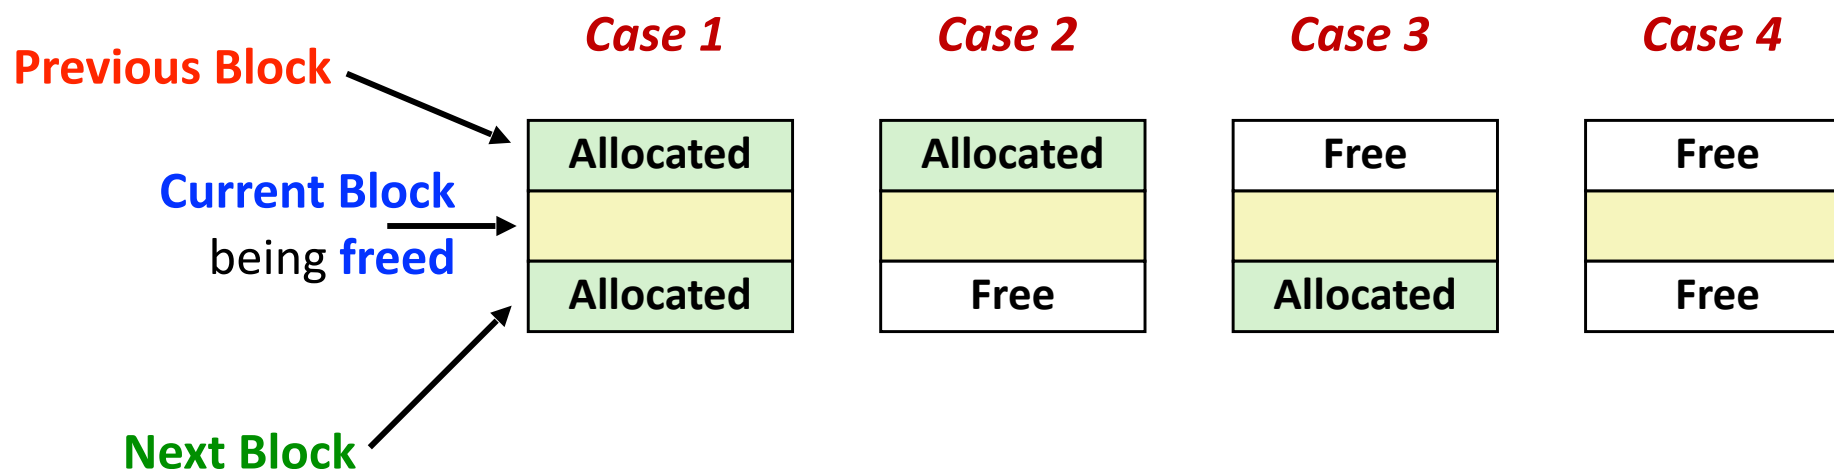

Splitting Free Block: Full Version




```
static void split_block(block_t *block, size_t asize) {
    size_t block_size = get_size(block);

    if ((block_size - asize) >= min_block_size) {
        write_header(block, asize, true);
        write_footer(block, asize, true);
        block_t *block_next = find_next(block);
        write_header(block_next, block_size - asize, false);
        write_footer(block_next, block_size - asize, false);
    }
}
```

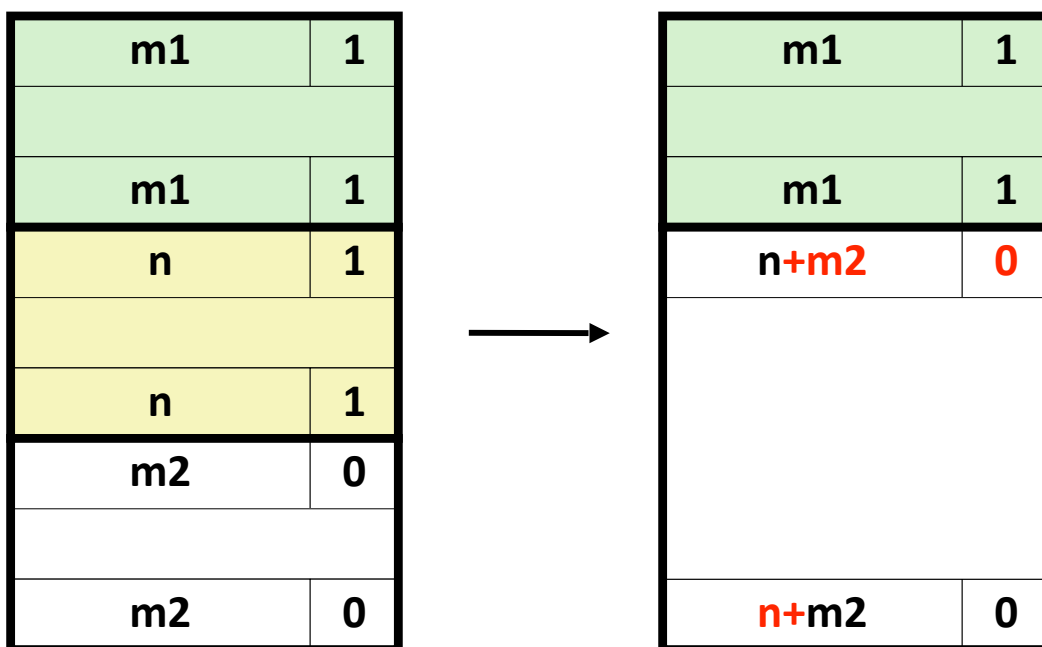
Constant Time Coalescing



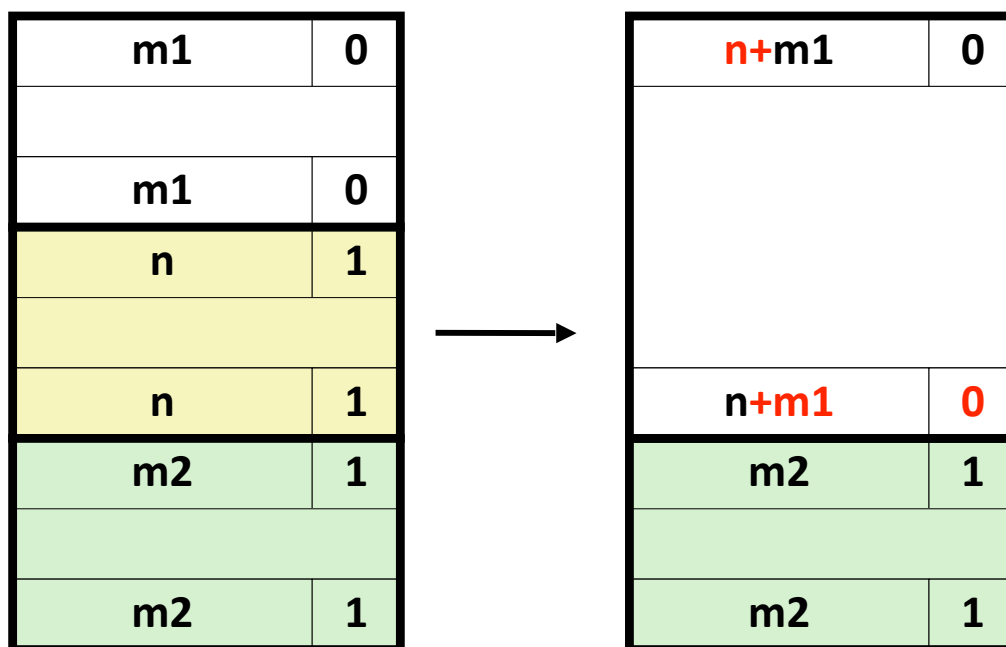
Constant Time Coalescing (Case 1)




Constant Time Coalescing (Case 2)



Constant Time Coalescing (Case 3)



Constant Time Coalescing (Case 4)




Disadvantages of Boundary Tags

- **Significant memory overhead** due to both header and footer **if** application uses **many small blocks** (e.g., graph with nodes each requiring one 1 word of payload)


- **Can it be optimized?**

- Which blocks need the footer tag?
- What does that mean?



No Boundary Tag for Allocated Blocks

- Boundary tag needed only for free blocks
- When sizes are multiples of 16, have **4 spare bits**




**Allocated
Block**

a = 1: Allocated block
a = 0: Free block
b = 1: 'Previous' block is allocated
b = 0: 'Previous' block is free

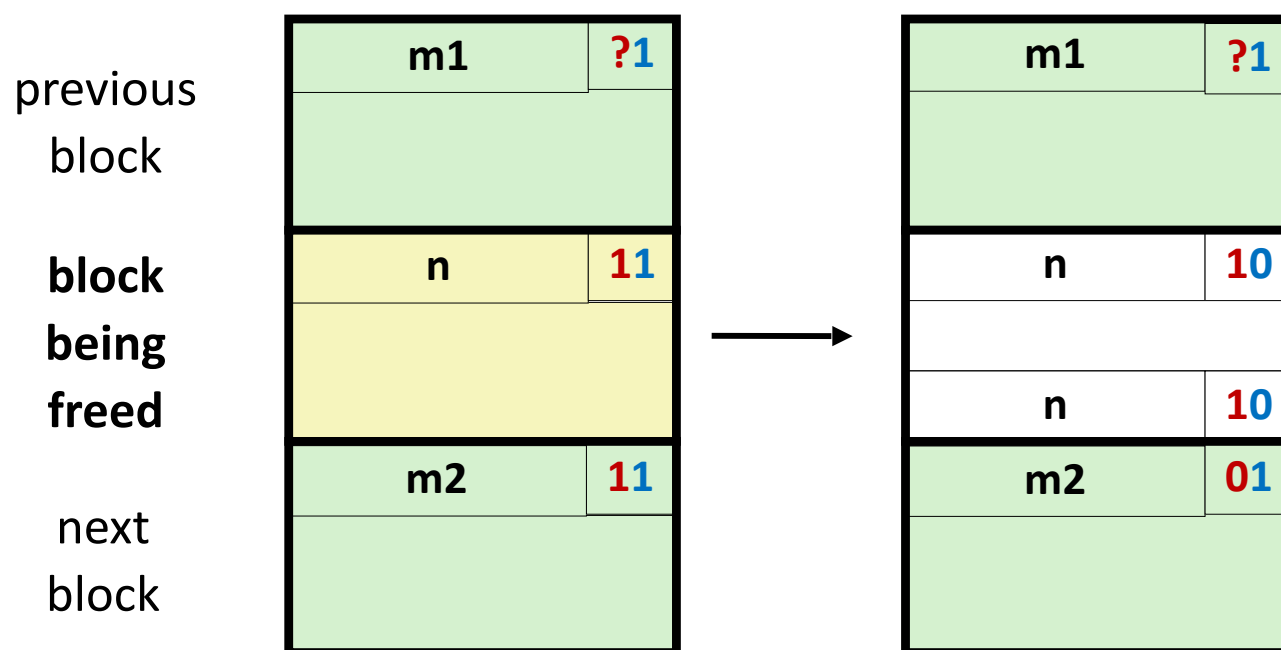
Size: block size

Payload: application data



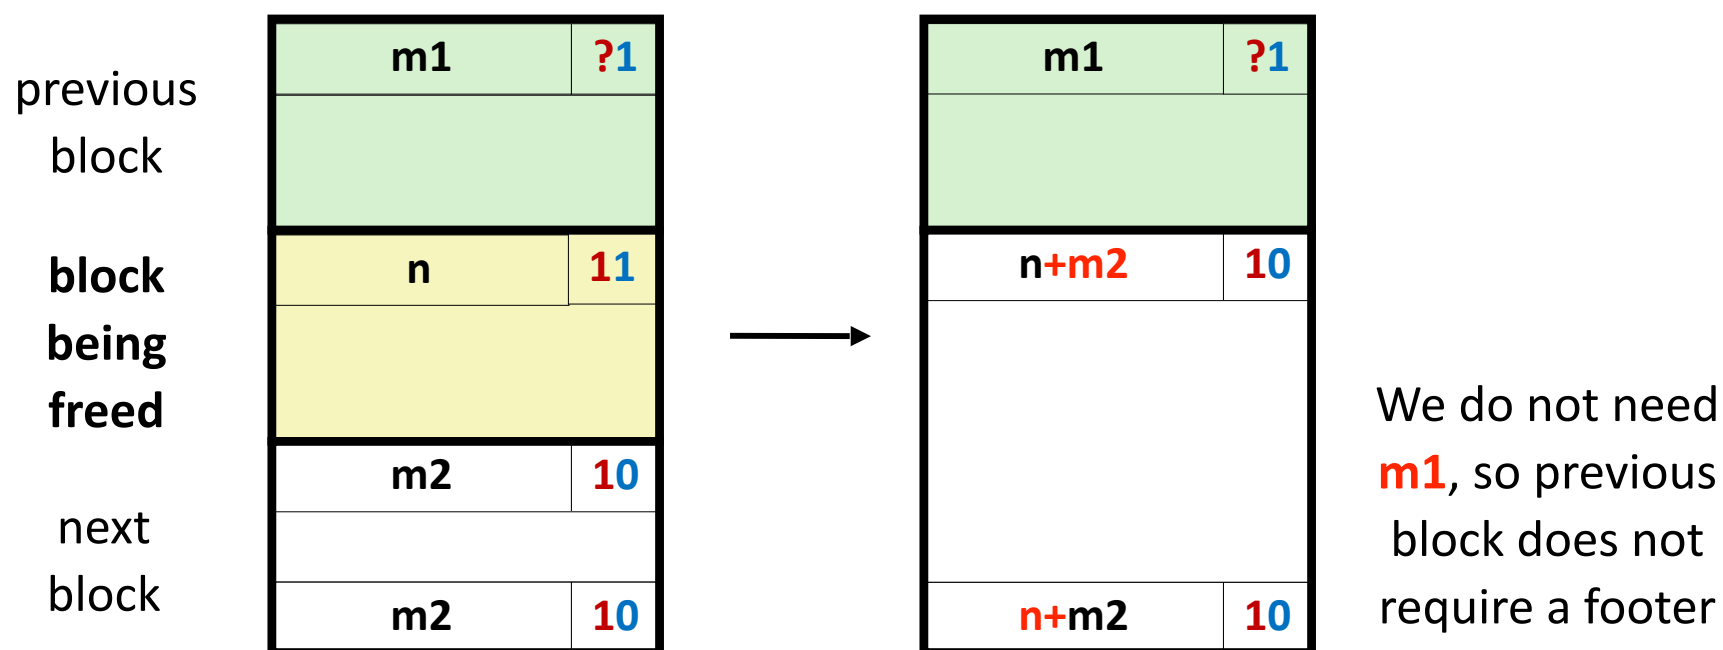
**Free
Block**

No Boundary Tag for Allocated Blocks (Case 1)



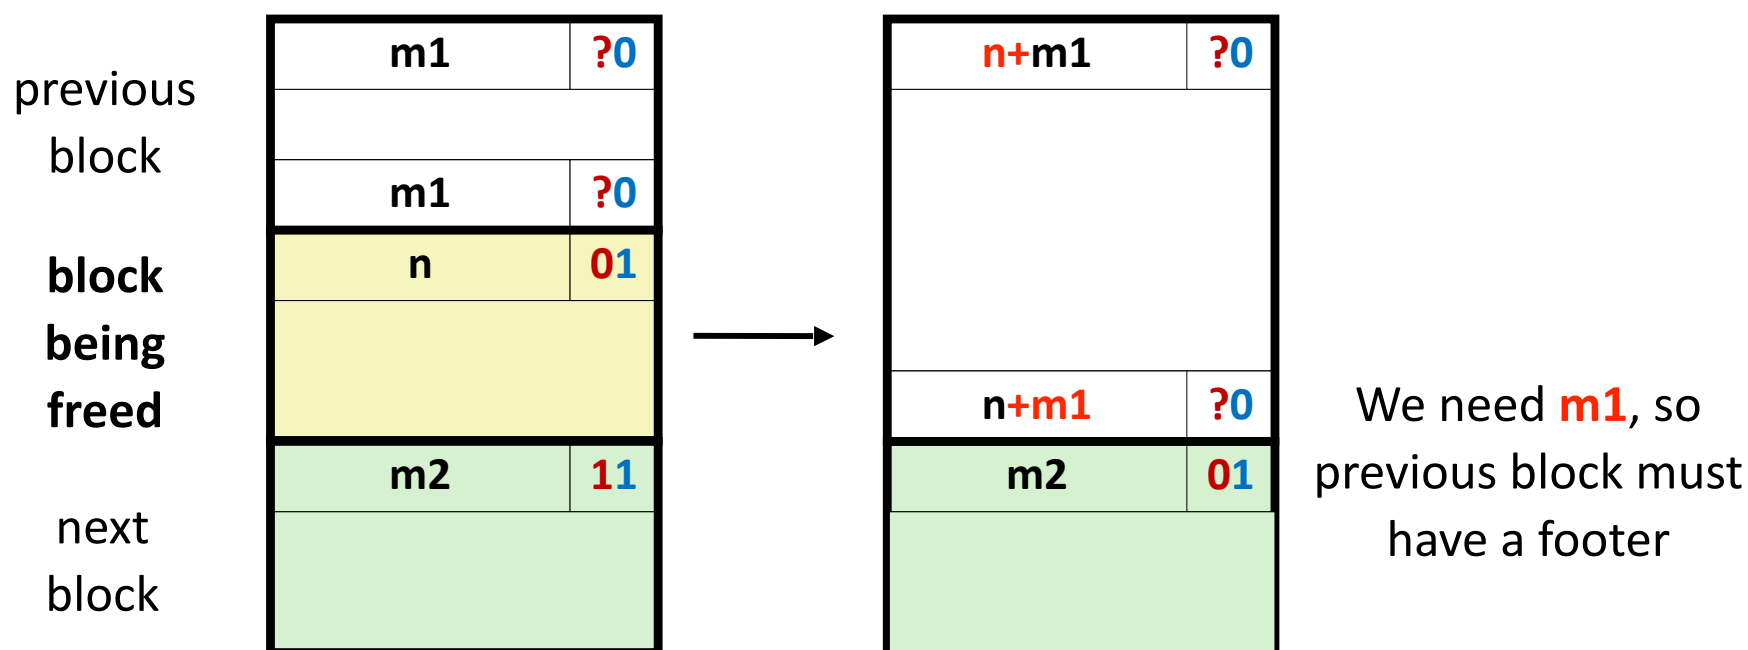
Header: Use 2 bits (address bits always zero due to alignment):
 (**previous block allocated**)<<1 | (**current block allocated**)

No Boundary Tag for Allocated Blocks (Case 2)



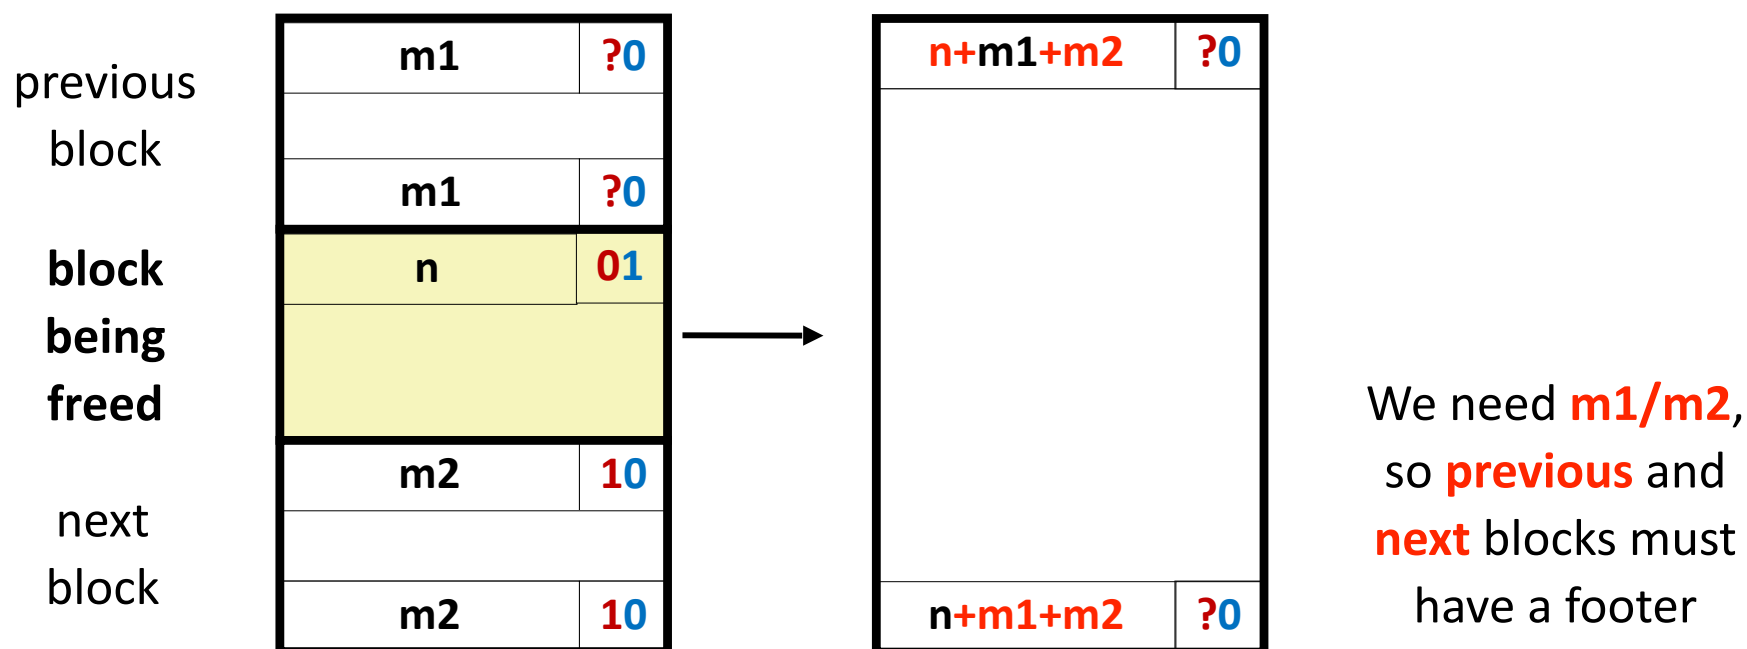
Header: Use 2 bits (address bits always zero due to alignment):
 (**previous block allocated**)<<1 | (**current block allocated**)

No Boundary Tag for Allocated Blocks (Case 3)




Header: Use 2 bits (address bits always zero due to alignment):
 (**previous block allocated**)<<1 | (**current block allocated**)

No Boundary Tag for Allocated Blocks (Case 4)



Header: Use 2 bits (address bits always zero due to alignment):
 (**previous block allocated**)<<1 | (**current block allocated**)

New Heap Structure After Adding Tags



■ Dummy footer before first header

- Marked as allocated
- Prevents accidental coalescing when freeing first block

■ Dummy header after last footer

- Prevents accidental coalescing when freeing final block

Top-Level Malloc Code

```
const size_t dsize = 2*sizeof(word_t);

void *mm_malloc(size_t size)
{
    // asize includes header and footer
    size_t asize = round_up(size + dsize, dsize);

    block_t *block = find_fit(asize);

    if (block == NULL)
        return NULL;

    // first utilize entire free area
    size_t block_size = get_size(block);
    write_header(block, block_size, true);
    write_footer(block, block_size, true);

    split_block(block, asize);

    // always return pointer to payload not block
    return header_to_payload(block);
}
```

$$\begin{aligned} \text{round_up}(n, m) \\ &= \\ m * ((n+m-1) / m) \end{aligned}$$

Top-Level Free Code

```
void mm_free(void *bp)
{
    block_t *block = payload_to_header(bp); // bp is pointer to payload
    size_t size = get_size(block);

    write_header(block, size, false);
    write_footer(block, size, false);

    coalesce_block(block);
}
```

Getting Additional Heap Memory

■ No fit found even after coalescing free blocks?

- 1) Ask the kernel for additional heap memory
 - Allocator calls **sbrk** function
- 2) Additional memory given by kernel is **transformed** into **a single large free block** and **attached** to the end of the heap
- 3) **Place requested block** in the new free block
 - Perform splitting if necessary

Summary of Key Allocator Policies

■ Placement policy:

- First-fit, next-fit, best-fit, etc.
- Trades off lower throughput for less fragmentation
- *Interesting observation:* **segregated free lists** approximate a best fit placement policy without having to search entire free list

■ Splitting policy:

- When do we go ahead and split free blocks?
- How much internal fragmentation are we willing to tolerate?

■ Coalescing policy:

- *Immediate coalescing:* coalesce each time **free** is called
- *Deferred coalescing:* try to improve performance of **free** by deferring coalescing until needed.

Implicit Lists: Summary

- **Implementation: very simple**
- **Allocate cost:**
 - linear time worst case
- **Free cost:**
 - constant time worst case
 - even with coalescing
- **Memory Overhead**
 - will depend on placement policy
 - First-fit, next-fit or best-fit
- **Not used in practice** for `malloc/free` because of linear-time allocation
 - used in many **special purpose applications**
- **However, the concepts of splitting and boundary tag coalescing are general to *all* allocators**